

JEARON WONG / AGENTIC LIFECYCLE GOVERNANCE INDUSTRY SERIES / 03
INSURABILITY & RISK TRANSFER

Agentic AI Insurability & Risk Transfer White Paper 2026

A Lifecycle Evidence Guide for Underwriting, Claims, and Enterprise Risk Transfer

CANDIDATE ID

AIIRWP-2026-v0.1-R6-CANDIDATE

STATUS

Internal candidate artifact

ARTIFACT ROLE

HTML/PDF internal visual candidates

PREPARED

2026-05-21

Internal candidate artifact only. Not public, not final, not sealed, not insurer-accepted, not coverage-ready, not underwriting-ready, not legal advice, not insurance advice, not a coverage opinion, not an underwriting standard, and not claims approval guidance.

Publication Contents

1. [A Lifecycle Evidence Guide for Underwriting, Claims, and Enterprise Risk Transfer](#)
2. [Candidate Boundary](#)
3. [Governing Thesis](#)
4. [Mandatory Boundary Sentences](#)
5. [Source Posture](#)
6. [Contents](#)
7. [Chapter 00](#)
8. [Chapter 01](#)
9. [Chapter 02](#)
10. [Chapter 03](#)
11. [Chapter 04](#)
12. [Chapter 05](#)
13. [Chapter 06](#)
14. [Chapter 07](#)
15. [Chapter 08](#)
16. [Chapter 09](#)
17. [Chapter 10](#)
18. [Chapter 11](#)
19. [Chapter 12](#)
20. [Chapter 13](#)
21. [Chapter 14](#)
22. [Chapter 15](#)
23. [Chapter 16](#)
24. [Appendix A](#)
25. [Appendix B](#)
26. [Appendix C](#)
27. [Appendix D](#)
28. [Appendix E](#)
29. [Appendix F](#)
30. [Appendix G](#)
31. [Source Register and Citation Notes](#)
32. [Final Source Use Rules](#)
33. [Final Source ID List](#)
34. [Required Pre-Publication Recheck](#)
35. [Finalization Decision](#)
36. [Source Marker Style](#)
37. [Marker Validation](#)
38. [R2A Citation Hygiene Checks](#)
39. [R6/R6B Citation Tasks](#)

Agentic AI Insurability & Risk Transfer White Paper 2026

A Lifecycle Evidence Guide for Underwriting, Claims, and Enterprise Risk Transfer

AIIRWP-2026-v0.1-R6-CANDIDATE

jearon.wong.com

Copyright © 2026 Jearon Wong. All rights reserved.

Candidate ID: AIIRWP-2026-v0.1-R6-CANDIDATE

Status: Internal candidate artifact source

Series: Agentic Lifecycle Governance Industry Series

Series Position: 03 / Insurability & Risk Transfer

Author: Jearon Wong

Prepared: 2026-05-21

Source Basis: R4-reviewed R3 internal draft plus WP3-R0/R1/R2/R2A/R4 source and boundary QA.

> Internal source package only. Not published. Not final. Not sealed. Not insurer-accepted. Not coverage-ready. Not underwriting-ready. Not an underwriting standard. Not legal or insurance advice. Not claims approval guidance.

Candidate Boundary

This candidate source preserves the R2A source accuracy gate and the R4 editorial/source/boundary QA result. It is assembled for later internal artifact generation planning and author review. It does not create a public route, public artifact, public manifest, public checksum, public metadata, sitemap entry, entity graph entry, Evidence Registry entry, research index entry, homepage entry, or public DOCX.

Governing Thesis

Agentic AI cannot become broadly insurable until insurers can distinguish the insured legal subject from the agentic risk object, map human and organizational responsibility to agentic actions, and reconstruct the lifecycle evidence that connects authority, action, loss, remediation, and coverage boundaries.

Mandatory Boundary Sentences

- AI agents are not the insured legal subject.
- The company may be insured. The AI agent usually is not.
- Logs do not make AI agents insurable. Reconstructable lifecycle evidence does.
- Tool permission is not coverage authority.
- Framework trace is not claim evidence by itself.
- Workflow completion is not an insured outcome.
- Agentic AI risk is not only model risk. It is lifecycle risk.

Source Posture

External source usage is limited to verified, scoped, boundary safe claims from the R1/R2/R2A source base. Public market examples are treated as market signals or product examples, not proof of consensus. Technical framework docs are technical capability sources only, not insurance conclusions. AIO and AIRM are Jearon Wong synthesis / analytical object models unless directly source-supported.

Contents

1. Chapter 00: The Plain-English Problem: Why Agentic AI Breaks Today's Insurance Logic

2. Chapter 01: The Insured Subject Problem: Who Is Covered When an Agent Acts?

3. Chapter 02: The Insurable Object Problem: What Exactly Is Being Covered?

4. Chapter 03: The Responsibility Mapping Problem: How Human Liability and Agentic AI Risk Connect

5. Chapter 04: What AI Insurance Covers Today, and Why Agentic AI Still Falls Through the Gap

6. Chapter 05: Why Agentic AI Is Not Yet Broadly Insurable

7. Chapter 06: From Compliance and Auditability to Insurability

8. Chapter 07: Why Logs, Traces, and Vendor Assurances Are Not Claim Evidence

9. Chapter 08: The Agentic Insurability Object Model

10. Chapter 09: Coverage Boundaries, Authority, Delegation, and Exclusion Triggers

11. Chapter 10: Loss Event Reconstruction and Causality Tracing

12. Chapter 11: Third-Party, Vendor, Model, and Tool Dependency Risk

13. Chapter 12: Evidence Requirements for Underwriting

14. Chapter 13: Evidence Requirements for Claims Review

15. Chapter 14: Aggregation and Accumulation Risk

16. Chapter 15: Agentic Insurability Readiness Model

17. Chapter 16: Conclusion: From Agentic Deployment to Insurable Risk Transfer

18. Appendix A: Agentic Insurability Object Checklist

19. Appendix B: Underwriting Evidence Request Template

20. Appendix C: Claims Reconstruction Evidence Package

21. Appendix D: AIO-to-MRO Mapping

22. Appendix E: AIO-to-Audit-Evidence-Chain Mapping

23. Appendix F: AIRM Readiness Matrix

24. Appendix G: Boundary and Non-Claim Language

Chapter 00

The Plain-English Problem: Why Agentic AI Breaks Today’s Insurance Logic

Insurance still starts with plain questions: who is covered, what risk is covered, what happened, who was responsible, and what evidence supports the claim. [SRC: INS-04] [SRC: INS-05] [SRC: INS-06] [SRC: CLAIM-01] [SRC: CLAIM-02]

Public market examples show AI-specific cover at the edge, but they remain narrow, conditional, and product-specific rather than broad lifecycle risk transfer. [SRC: MKT-01] [SRC: MKT-02] [SRC: MKT-03] [SRC: MKT-05] [SRC: MKT-08]

This paper defines the bridge from policy subject to claim review as:

Legal Subject -> Human Responsibility Role -> Agent / MAS Role -> Agentic Work Unit -> Loss Event -> Claim Evidence Chain -> Coverage / Exclusion Decision. [SYNTHESIS: Jearon Wong] [INT: INT-01] [INT: INT-05] [INT: INT-06]

Bridge Figure



Insurance Basics in Plain English

Question	Plain-English answer	Why it matters for agentic AI	Source
Who is covered?	A person, firm, officer, vendor, or organization.	The AI system is not automatically the insured party.	[SRC: INS-04] [SRC: INS-06] [SRC: INS-08]
What is covered?	A policy-defined risk, limit, or exposure.	A model name or workflow label is not enough.	[SRC: INS-05] [SRC: INS-07]
What happened?	A loss event with facts and timing.	Agentic work needs event reconstruction.	[SRC: CLAIM-01] [SRC: CLAIM-02]
Who was responsible?	A mapped human or organizational role.	Responsibility cannot stay hidden inside automation.	[SRC: INS-01] [INT: INT-01] [INT: INT-05]
What evidence exists?	Records that support review.	Logs are inputs, not the full claim file.	[SRC: CLAIM-01] [SRC: CLAIM-03] [INT: INT-05]

Discussion

Agentic AI Insurability & Risk Transfer White Paper 2026 | Internal Candidate Artifact

The insurance problem is not that agentic AI is mysterious. It is that the usual insurance questions break apart when action, authority, and responsibility are split across a person, a company, a tool, a model, and a workflow.

[SRC: INS-01] [SRC: CLAIM-01] [INT: INT-01]

AIIRWP-2026-v0.1-R6-CANDIDATE

jearonwong.com

Copyright © 2026 Jearon Wong. All rights reserved.

AI governance sources already expect documentation, controls, and reviewability. That does not make those records insurance evidence by themselves, but it does show why the evidence question comes first.

[SRC: INS-01] [SRC: AI-01]

[SRC: AI-08] [INT: INT-05]

WP1 / WP2 Bridge

- WP1 contributes MROs, ALCS, accepted outcome, authority boundary, and remediation closure. [INT: INT-01]
[INT: INT-02]
- WP2 contributes the Audit Evidence Chain and the logs-versus-evidence distinction. [INT: INT-04] [INT: INT-05]
- This white paper translates those into the claim-review and risk-transfer question. [SYNTHESIS: Jearon Wong]

Boundary

- Not a coverage opinion.
- Not an underwriting standard.
- Not claims approval guidance.
- Not a legal liability determination.

Chapter 01

AIIRWP-2026-v0.1-R6-CANDIDATE

jearonwong.com

Copyright © 2026 Jearon Wong. All rights reserved.

The Insured Subject Problem: Who Is Covered When an Agent Acts?

Insurance analysis starts with an insured party, policy subject, or insured role. [SRC: INS-04] [SRC: INS-06] [SRC: INS-08]

An AI agent can act, recommend, draft, route, or invoke tools, but that does not make the agent the insured legal subject. The company may be insured. The AI agent usually is not. [SRC: INS-04] [SRC: INS-06] [INT: INT-01] [SYNTHESIS: Jearon Wong]

Argument

The subject question comes before the model question. If the policy attaches to a person, firm, officer, or organization, then the analysis begins there and works outward to responsibility and evidence. [SRC: INS-04] [SRC: INS-08]

Additional insured and D&O/E&O sources help show that insured status is a policy relationship, not a model attribute. [SRC: INS-04] [SRC: INS-07] [SRC: INS-08]

Subject Map

Insurance question	Traditional answer	Agentic AI problem	Required mapping	Source
Who is insured?	Named insured, additional insured, officer, professional, vendor, organization.	The agent may act without being the insured subject.	Legal subject first, agent second.	[SRC: INS-04] [SRC: INS-06] [SRC: INS-08]
Who is accountable?	A legal or organizational party named or implied by the contract.	HITL does not equal responsibility structure.	Human role must be explicit.	[SRC: INS-01] [SRC: MKT-07] [INT: INT-01]
What is not enough?	A tool label or model name.	The AI agent name is not the policy subject.	Policy language must anchor the subject.	[SRC: INS-05] [SRC: INS-06]

Discussion

The safe wording is simple: an AI agent is generally a tool or actor, not the insured legal subject. That does not settle liability, and it does not decide coverage. It only preserves the insurance starting point. [SRC: INS-04] [SRC: INS-06] [SRC: INS-08]

This chapter should use examples where a human, a firm, or an officer is the covered subject while the agent is the acting layer. That makes the policy question visible before the automation question arrives. [INT: INT-01] [INT: INT-04] [INT: INT-05]

WP1 / WP2 Bridge

- WP1 MROs support the subject and responsibility split. [INT: INT-01]

- WP2 adds the audit-evidence lens for who accepted, reviewed, or remediated the output. [INT: INT-04]
- Agentic AI Insurability & Risk Transfer White Paper 2026 | Internal Candidate Artifact
- [INT: INT-05]

- Chapter 1 should not drift into legal advice or jurisdiction-specific doctrine. [SYNTHESIS: Jearon Wong]

Boundary

WP2-2026-v0.1-R6-CANDIDATE

jearonwong.com

Copyright © 2026 Jearon Wong. All rights reserved.

- Not a legal opinion.
 - Not a coverage decision.
 - Not a liability determination.
-

Chapter 02

The Insurable Object Problem: What Exactly Is Being Covered?

The insured subject is not the same as the risk object. A policy can cover a person or organization while the exposure is a narrower activity, event, or operation. [SRC: INS-05] [SRC: INS-07] [SRC: MKT-01] [SRC: MKT-02]

Public AI and cyber materials describe performance, error, misuse, cyber events, or AI-linked service harm. They do not make a model name or workflow label into the insurable object by itself. [SRC: MKT-01] [SRC: MKT-02] [SRC: MKT-03] [SRC: INS-09] [SRC: INS-10]

Claim

This white paper defines the bounded agentic work unit, operation, exposure, or loss-triggering activity as the reviewable risk object. That is author synthesis, not a market standard. [SYNTHESIS: Jearon Wong] [INT: INT-01] [INT: INT-06]

Layer Table

Layer	What it is	Insurance function	Agentic AI gap	Source
Subject	Covered person or organization	Identifies who can be indemnified or defended	AI agent is usually not the subject	[SRC: INS-04] [SRC: INS-06] [SRC: INS-08]
Work unit	Bounded agentic task or operation	Makes exposure reviewable	No external source names this as a policy object	[SYNTHESIS: Jearon Wong] [INT: INT-06]
Workflow label	System or product name	Operational shorthand only	Label alone does not bound loss	[SRC: TECH-01] [SRC: TECH-04]
Vendor/model name	Provider or model identity	Useful dependency context	Not the same as the insurable object	[SRC: MKT-01] [SRC: TECH-01] [SRC: TECH-04]
Exposure boundary	Scope, authority, time, data, tools, outcome	Enables review of what was at stake	Must be reconstructed, not assumed	[SRC: INS-05] [SRC: CLAIM-01] [INT: INT-05]

Discussion

Chapter 2 should keep the distinction sharp: the object of review is not "the agent" in the abstract. It is the bounded thing the agent did, under the authority it had, with the tools it used, inside the scope the enterprise accepted. [SRC: INS-01] [SRC: CLAIM-01] [INT: INT-01] [INT: INT-05]

That bounded object is the analytic bridge between technical logs and insurance review. It is a synthesis object that makes the exposure legible without pretending to be a policy form. [SYNTHESIS: Jearon Wong] [INT: INT-06]

WP1 / WP2 Bridge

- WP1 MRO-02, MRO-04, MRO-05, and MRO-08 help bound the exposure.

- WP2 auditability work helps separate workflow labels from evidence objects. [INT: INT-04] [INT: INT-05]
- Agentic AI Insurability & Risk Transfer White Paper 2026 | Internal Candidate Artifact

Boundary

- Not an underwriting standard.
 - Not a policy form.
 - Not a claim decision.
- AIIRWP-2026-v0.1-R6-CANDIDATE jearonwong.com Copyright © 2026 Jearon Wong. All rights reserved.

Chapter 03

The Responsibility Mapping Problem: How Human Liability and Agentic AI Risk Connect

Human-in-the-loop is not the same thing as responsibility mapping. A button click can confirm an action, but it does not by itself establish authority, liability structure, or responsibility transfer. [SRC: INS-01] [SRC: CLAIM-03] [SRC: MKT-07] [INT: INT-05]

Argument

The insurance bridge needs to show how a human role, an agent role, a work unit, evidence, and a loss event connect back to a legal subject. [SYNTHESIS: Jearon Wong] [INT: INT-01] [INT: INT-04] [INT: INT-05]

Governance and D&O / professional-liability sources support the idea that responsibility is organizational and role-based. They do not say the AI agent becomes the liable subject. [SRC: INS-07] [SRC: INS-08] [SRC: MKT-07]

Responsibility Matrix

Human role	Agent role	Work unit	Evidence needed	Loss question	Legal subject
Approver	Tool-using agent	Bounded task	Approval record, authority scope, timestamp	What was approved?	Company or officer
Operator	Delegating employee	Service workflow	Delegation, handoff, output record	Who ran the workflow?	Firm or team
Reviewer	Multi-agent controller	Escalated exception	Review notes, exception route, acceptance/rejection	Who accepted the result?	Organization or officer
Remediator	Recovery owner	Post-loss closure	Fix record, recheck, closure state	Was it contained?	Legal subject with duty

Discussion

This chapter should say plainly that responsibility mapping is an evidence problem before it is a legal problem. The paper is not deciding liability; it is showing what evidence a later liability or coverage review would need. [SRC: CLAIM-01] [SRC: CLAIM-02] [SRC: CLAIM-03]

The precise bridge is the internal relationship between human role, agent role, and work unit. This white paper uses that bridge to separate oversight from responsibility and responsibility from legal conclusion. [INT: INT-01] [INT: INT-04] [INT: INT-05]

WP1 / WP2 Bridge

- WP1 human-role and MAS mapping becomes the responsibility path.
- WP2 audit evidence becomes the reviewable chain for acceptance, exception, and remediation.
- HITL is only one input; it is not the whole responsibility model. [INT: INT-01] [INT: INT-04] [INT: INT-05]

Boundary

Agentic AI Insurability & Risk Transfer White Paper 2026 | Internal Candidate Artifact

- Not a legal liability determination.
- Not a coverage opinion.
- Not a certification or standard.

AIRWP-2026-V0.1-R6-CANDIDATE

jearonwong.com

Copyright © 2026 Jearon Wong. All rights reserved.

Chapter 04

What AI Insurance Covers Today, and Why Agentic AI Still Falls Through the Gap

Public market examples show AI-specific products, broker framing, cyber-linked AI cover, and D&O/E&O discussion around AI risk. They do not show a broad, standardized lifecycle risk-transfer market. [SRC: MKT-01] [SRC: MKT-02]

[SRC: MKT-03] [SRC: MKT-05] [SRC: MKT-06] [SRC: MKT-07] [SRC: MKT-08]

Current AI Insurance Focus

Focus	What it may cover	Why it is not enough for agentic AI	Source
AI-specific product examples	Model performance, error, or bespoke AI risk	Product existence is not market consensus	[SRC: MKT-01] [SRC: MKT-02]
AI-linked cyber	Data breach, misuse, LLMjacking, response costs	Cyber controls are not claim evidence by themselves	[SRC: MKT-03] [SRC: MKT-04] [SRC: INS-09] [SRC: INS-10]
E&O / professional liability	Professional mistakes or omissions	The AI agent is not the insured professional subject	[SRC: INS-07] [SRC: MKT-06]
D&O / corporate exposure	Governance, oversight, disclosure, board risk	Does not create an agentic lifecycle object	[SRC: INS-08] [SRC: MKT-07]
Silent AI exposure	AI may appear inside existing lines	Silent exposure is a market-risk signal, not a coverage conclusion	[SRC: MKT-05] [SRC: MKT-08] [SRC: CYB-02]

Discussion

This chapter should read like a market map, not a verdict. The public source base points to fragmented and conditional coverage at the edges. [SRC: MKT-01] [SRC: MKT-02] [SRC: MKT-05] [SRC: MKT-08]

The safe language is "certain product materials describe" and "some sources discuss." The unsafe language is "the industry believes" or "AI is broadly covered." [SYNTHESIS: Jearon Wong] [INT: INT-01] [INT: INT-03]

Boundary

- Use market examples as examples.
- Do not convert product pages into consensus claims.
- Do not convert broker views into policy interpretation.

Chapter 05

Why Agentic AI Is Not Yet Broadly Insurable

Agentic AI is not yet broadly insurable because insurers still need to bound the object, attribute the act, reconstruct the event, and understand aggregation before they can review the loss with confidence. [SRC: INS-01] [SRC: CLAIM-01] [SRC: CYB-02] [SRC: CYB-03] [SRC: CYB-04]

Gap Table

Insurance question	Current AI stack often provides	Missing lifecycle object	Source
What is the object?	Model name, workflow name, vendor name	Bounded agentic work unit	[SRC: MKT-01] [SRC: TECH-01] [SYNTHESIS: Jearon Wong]
Who is responsible?	Logs, approvals, handoffs	Human-agent responsibility map	[SRC: CLAIM-03] [INT: INT-01] [INT: INT-05]
What happened?	Events and traces	Loss event record with context	[SRC: CLAIM-01] [SRC: CLAIM-02] [INT: INT-05]
What was remediated?	Alerts or tickets	Remediation and recovery record	[SRC: CLAIM-02] [SRC: CLAIM-03] [INT: INT-05]
What depends on what?	Technical dependency lists	Dependency and aggregation view	[SRC: CYB-02] [SRC: CYB-03] [SRC: CYB-04] [INT: INT-03]

Discussion

The core claim should be narrow: current sources show market experimentation and governance pressure, but not a standardized lifecycle object layer for agentic AI. [SRC: MKT-01] [SRC: MKT-05] [SRC: MKT-08] [SRC: CLAIM-01] [SRC: TECH-01]

That means the gap is not just the LLM, the framework, or the vendor. It is the missing layer that turns action into reviewable risk. [SYNTHESIS: Jearon Wong] [INT: INT-06]

WP1 / WP2 Bridge

- WP1 shows why lifecycle failure modes need object boundaries.
- WP2 shows why auditability needs evidence chains, not just logs. [INT: INT-01] [INT: INT-04] [INT: INT-05]
- This white paper converts those into insurability terms.

Boundary

- Do not say "AI agents are uninsurable."
- Do not say "AI agents are insurable."
- Do not imply the market has settled the question.

Chapter 06

From Compliance and Auditability to Insurability

WP1 and WP2 do not become insurance standards. They become source truth for translating compliance and auditability into claim-relevant language. [INT: INT-01] [INT: INT-04] [INT: INT-05] [INT: INT-06] [INT: INT-07]

Translation Table

WP1 / WP2 object	Insurance translation in this paper	Why it matters	Source
MRO	AIO review object family	Gives the risk object vocabulary	[INT: INT-01] [INT: INT-06]
ALCS	Claim reconstructability lens	Shows how visible the lifecycle is	[INT: INT-01] [INT: INT-02]
Audit Evidence Chain	Claim Evidence Chain	Connects event, responsibility, and recovery	[INT: INT-04] [INT: INT-05]
AARM	AIRM readiness vocabulary	Describes evidence readiness without certification	[INT: INT-07]
Enterprise failure scenarios	Insurance failure modes	Shows what can break in loss review	[INT: INT-01]

Discussion

The useful sentence is simple: compliance and auditability are necessary ingredients, but they do not by themselves establish insurability. [SRC: INS-01] [SRC: CLAIM-01] [SYNTHESIS: Jearon Wong]

This chapter should help readers see the lineage. WP1 gives the lifecycle objects. WP2 gives the evidence-chain lens. this white paper gives the insurance translation of both. [INT: INT-01] [INT: INT-04] [INT: INT-05] [INT: INT-06] [INT: INT-07]

Boundary

- Not an assurance opinion.
- Not a coverage opinion.
- Not a certification claim.

Chapter 07

Why Logs, Traces, and Vendor Assurances Are Not Claim Evidence

Technical traces are useful inputs, not claim evidence by themselves. They can show activity, timing, and flow, but they do not automatically prove authority, responsibility, remediation, or coverage boundary. [SRC: CLAIM-01]


[SRC: CLAIM-02]

[SRC: CLAIM-03]

[SRC: TECH-01]

[SRC: TECH-04]

Sequence Figure



Discussion

Logs and traces can help reconstruct what happened. They cannot, by themselves, answer the insurance questions that matter most: was the action authorized, who owned it, what failed, what was remediated, and what evidence belongs in the claim file. [SRC: CLAIM-01] [SRC: CLAIM-02] [SRC: CLAIM-03] [INT: INT-05]

Vendor assurances can support risk discussion. They are not coverage authority. Framework traces can support technical reconstruction. They are not claim evidence by themselves. [SYNTHESIS: Jearon Wong] [INT: INT-05]

[INT: INT-06]

Source-Layer Distinction

Signal	What it shows	What it does not show	Source
Logs	Events, timestamps, service activity	Authority or accepted outcome	[SRC: CLAIM-01] [SRC: TECH-01] [INT: INT-05]
Traces	Workflow path and service flow	Responsibility transfer or legal causation	[SRC: CLAIM-01] [SRC: TECH-04] [INT: INT-05]
Vendor assurances	Market position or product intent	Coverage boundary or claim approval	[SRC: MKT-03] [SRC: MKT-05] [SRC: MKT-08]
Incident guidance	Response and recovery vocabulary	Insurance decision	[SRC: CLAIM-01] [SRC: CLAIM-02] [SRC: CLAIM-03]

Boundary

Agentic AI Insurability & Risk Transfer White Paper 2026 | Internal Candidate Artifact

- Do not say traces prove the claim.
- Do not say logs are useless.
- Do not say a vendor statement creates insurance evidence.

AIRWP-2026-V0.1-R6-CANDIDATE

jearonwong.com

Copyright © 2026 Jearon Wong. All rights reserved.

Chapter 08

The Agentic Insurability Object Model

AIO v2 is a Jearon Wong synthesis: an analytical insurability object model, not an external standard, policy form, or insurer requirement. [SYNTHESIS: Jearon Wong] [INT: INT-06]

AIO Catalog

AIO	Plain-English label	Insurance use	Boundary note	Source
AIO-01	Legal insured subject	Names the covered party	Not legal advice	[SRC: INS-04] [SRC: INS-06] [SRC: INS-08]
AIO-02	Insurable agentic work unit	Binds exposure to a bounded action	Synthesis object	[SYNTHESIS: Jearon Wong] [INT: INT-06]
AIO-03	Human-agent responsibility map	Links action back to human role	Not liability finding	[SRC: INS-01] [SRC: INS-08] [INT: INT-01]
AIO-04	Coverage boundary	Frames limits, exclusions, scope	Not coverage opinion	[SRC: INS-05] [SRC: INS-07] [SRC: MKT-05]
AIO-05	Authority and delegation boundary	Separates permission from business authority	Not coverage authority	[SRC: TECH-01] [SRC: TECH-02] [INT: INT-05]
AIO-06	Loss event record	Anchors event, time, consequence	Not proof of loss alone	[SRC: CLAIM-01] [SRC: CLAIM-02]
AIO-07	Causality reconstruction trace	Links sequence and contributors	Not legal causation proof	[SRC: CLAIM-01] [SRC: TECH-04]
AIO-08	Control failure record	Shows failed or bypassed control	Not negligence finding	[SRC: INS-01] [SRC: AI-01] [SRC: CLAIM-01]
AIO-09	Claim evidence chain	Packages reviewable evidence	Not claim approval	[SRC: CLAIM-01] [SRC: CLAIM-02] [INT: INT-04]
AIO-10	Remediation and recovery record	Shows containment and closure	Not legal closure	[SRC: CLAIM-02] [SRC: CLAIM-03]
AIO-11	Vendor / model / tool dependency map	Shows dependency concentration	No vendor ranking	[SRC: CYB-02] [SRC: TECH-01] [SRC: TECH-03] [SRC: TECH-04] [SRC: TECH-05]
AIO-12	Exclusion trigger / boundary breach map	Helps review boundary facts	Not exclusion determination	[SRC: INS-05] [SRC: MKT-03] [SRC: MKT-08]
AIO-13	Aggregation and accumulation risk view	Shows correlated exposure	Not actuarial proof	[SRC: CYB-01] [SRC: CYB-02] [SRC: CYB-03] [SRC: CYB-04]
AIO-14	Dispute-ready claim package	Organizes evidence for review	Not guaranteed payment	[SRC: CLAIM-01] [SRC: CLAIM-02] [INT: INT-05]

Discussion

The AIO family is the paper's central object layer. It turns an agentic system into something underwriting, claims, and dispute review can discuss without pretending that a model name is enough. [SYNTHESIS: Jearon Wong]
[INT: INT-06]

Boundary

Agentic AI Insurability & Risk Transfer White Paper 2026 | Internal Candidate Artifact

- AIO is not a standard.
- AIO is not an insurer product requirement.
- AIO is not a legal conclusion.

AIRWP-2026-00.1-R8-CANDIDATE

jearonwong.com

Copyright © 2026 Jearon Wong. All rights reserved.

Chapter 09

Coverage Boundaries, Authority, Delegation, and Exclusion Triggers

Coverage boundaries are review questions, not coverage opinions. Authority and delegation boundaries help show whether the action was within the permitted scope, but they do not by themselves decide the claim. [SRC: INS-05]

[SRC: INS-07] [SRC: MKT-03] [SRC: MKT-05] [SRC: MKT-08] [INT: INT-05]

Boundary Table

Action	Tool permission	Business authority	Confirmation required	Coverage risk
Drafts a message	Allowed by tool	Maybe allowed	Low	May be low unless content causes harm
Submits a payment	Allowed by tool	Needs explicit scope	High	Higher if outside delegated authority
Changes a record	Allowed by tool	Needs role-based approval	High	Boundary review may be needed
Escalates an exception	Allowed by tool	Needs owner review	Medium	Coverage depends on facts and terms
Calls a vendor tool	Allowed by tool	Depends on business scope	Medium	Dependency and exclusion review may matter

Discussion

This chapter should keep three things separate: technical permission, business authority, and insurance boundary. Tool permission is not coverage authority. [SRC: TECH-01] [SRC: TECH-02] [SYNTHESIS: Jearon Wong]

Exclusion triggers and boundary breaches are evidence contexts. They are not denial rules in the abstract. The paper should keep the language careful and policy-neutral. [SRC: INS-05] [SRC: MKT-03] [SRC: MKT-05] [SRC: MKT-08]

WP1 / WP2 Bridge

- WP1 authority and confirmation boundary becomes the review of permission versus delegated scope.
- WP2 exception traceability becomes the claim boundary inquiry.
- AIO-04, AIO-05, and AIO-12 carry the chapter's analytical load. [INT: INT-01] [INT: INT-04] [INT: INT-05] [INT: INT-06]

Boundary

- Not a coverage opinion.
- Not a denial rule.
- Not an underwriting standard.

Chapter 10

Loss Event Reconstruction and Causality Tracing

The claim question is not just what happened. It is how the event can be reconstructed, what failed, what was remediated, and what evidence chain supports review. [SRC: CLAIM-01] [SRC: CLAIM-02] [SRC: CLAIM-03]

Reconstruction Table

Reconstruction element	Evidence needed	Source type	Boundary risk
Loss event record	Time, event, effect, owner	Incident / disclosure guidance	Not a full claim by itself
Causality trace	Sequence, dependencies, contributors	Incident response + technical trace	Not legal causation proof
Control failure record	Failed or bypassed control	Governance / incident guidance	Not negligence finding
Remediation record	Fix, recheck, closure	Response and recovery guidance	Not settlement or legal closure
Dispute-ready package	Combined review file	Internal synthesis	Not guaranteed payment

Discussion

Logs and traces can support reconstruction, but reconstruction must reach beyond the raw record. It needs role, authority, and remediation context so the claim file can be reviewed rather than guessed. [SRC: CLAIM-01]
[SRC: CLAIM-02] [SRC: CLAIM-03] [INT: INT-05]

This chapter should distinguish sequence from causation. Sequence can be documented. Legal causation cannot be assumed from telemetry. [SRC: TECH-04] [INT: INT-06]

WP1 / WP2 Bridge

- WP1 accepted outcome and remediation closure map to the post-loss story.
- WP2 Audit Evidence Chain becomes this paper's Claim Evidence Chain.
- AIO-06 to AIO-10 do most of the work here. [INT: INT-01] [INT: INT-04] [INT: INT-05] [INT: INT-06]

Boundary

- Not legal causation.
- Not claim approval guidance.
- Not a settlement promise.

Chapter 11

Third-Party, Vendor, Model, and Tool Dependency Risk

Agentic AI risk is rarely single-node risk. It spreads across models, tools, vendors, subprocessors, runtimes, templates, and repeated workflows. [SRC: CYB-02] [SRC: CYB-03] [SRC: CYB-04] [SRC: MKT-08] [SRC: TECH-01]

[SRC: TECH-02] [SRC: TECH-03] [SRC: TECH-04] [SRC: TECH-05]

Dependency Table

Dependency layer	Risk	Evidence needed	Aggregation concern	Source
Model provider	Model change, outage, error	Dependency map, version record	Shared model exposure	[SRC: MKT-01] [SRC: TECH-01]
Tool provider	Tool failure or misuse	Tool inventory, permission scope	Shared tool exposure	[SRC: TECH-02] [SRC: TECH-05]
Vendor chain	Subprocessor or upstream service issue	Contract and handoff map	Concentrated dependency	[SRC: CYB-02] [SRC: CYB-03]
Runtime / framework	Checkpoint, handoff, or persistence issues	Trace and state record	Shared orchestration exposure	[SRC: TECH-04]
Cross-project reuse	Same template across many work units	Work-unit registry	Correlated loss across projects	[INT: INT-03] [INT: INT-06]

Discussion

This chapter should frame dependency risk as evidence for concentration, not as a vendor scorecard. The reader should understand the chain without seeing a ranking. [SYNTHESIS: Jearon Wong] [INT: INT-03]

Cyber accumulation sources are useful only as analogy and risk framing. They do not become direct actuarial proof for agentic AI. [SRC: CYB-02] [SRC: CYB-03] [SRC: CYB-04]

Boundary

- No vendor ranking.
- No procurement recommendation.
- No claim that one provider is safer by default.

Chapter 12

AIIRWP-2026-v0.1-R6-CANDIDATE

jearonwong.com

Copyright © 2026 Jearon Wong. All rights reserved.

Evidence Requirements for Underwriting

Underwriting discussion needs visible evidence categories before it can assess an agentic system. That is not the same thing as a formal underwriting standard. [SRC: INS-01] [SRC: INS-09] [SRC: INS-10] [SRC: AI-01] [SRC: CLAIM-01]

Evidence Table

Evidence category	Why it matters	AIO	Boundary note
Subject evidence	Shows who the insured party is	AIO-01	Not legal advice
Work-unit evidence	Shows what the agentic system actually did	AIO-02	Not a policy form
Authority evidence	Shows what was permitted	AIO-05	Not coverage authority
Dependency evidence	Shows what systems and vendors were involved	AIO-11	No vendor ranking
Control evidence	Shows what control points existed	AIO-08	Not sufficient by itself
Aggregation evidence	Shows correlated exposure	AIO-13	Not actuarial proof
Remediation evidence	Shows how failure was handled	AIO-10	Not a guarantee of insurability

Discussion

The chapter should say that pre-bind evidence helps underwriting discussion because it makes scope visible. It should not present a checklist as a universal insurer requirement. [SYNTHESIS: Jearon Wong] [INT: INT-01] [INT: INT-06]

Boundary

- Not an underwriting standard.
- Not an insurer checklist.
- Not a pricing model.

Chapter 13

Evidence Requirements for Claims Review

Claims review needs a reconstructable package, not just a log dump. The review has to connect event, authority, responsibility, loss, and remediation. [SRC: CLAIM-01] [SRC: CLAIM-02] [SRC: CLAIM-03] [SRC: INS-05]

Claims Table

Claims evidence category	Why it matters	AIO	Boundary note
Loss event record	Anchors the incident	AIO-06	Not proof of payment
Causality trace	Shows sequence and contributors	AIO-07	Not legal causation proof
Responsibility map	Shows who owned the action	AIO-03	Not liability finding
Boundary map	Shows scope and exclusions context	AIO-04, AIO-12	Not coverage opinion
Remediation record	Shows containment and closure	AIO-10	Not legal closure
Dispute-ready package	Makes review efficient	AIO-14	Not claim approval guidance

Discussion

The claims story should be written in the language of reconstruction and review. It should not promise payment or imply that a package guarantees approval. [SYNTHESIS: Jearon Wong] [INT: INT-04] [INT: INT-05] [INT: INT-06]

Boundary

- Not claim approval guidance.
- Not a coverage opinion.
- Not a settlement promise.

Chapter 14

AIIRWP-2026-v0.1-R6-CANDIDATE

jearonwong.com

Copyright © 2026 Jearon Wong. All rights reserved.

Aggregation and Accumulation Risk

The concern is correlated loss, not isolated error. Shared models, shared tools, shared templates, and shared vendors can create a portfolio shape that behaves more like accumulation than one-off failure. [SRC: CYB-01] [SRC: CYB-02]

[SRC: CYB-03] [SRC: CYB-04] [SRC: MKT-08]

Aggregation Table

Aggregation driver	Example	Evidence needed	Reinsurance concern
Shared model	Many teams use the same model version	Model/version map	Concentrated loss path
Shared tool	Same orchestration tool across workflows	Tool inventory	Multi-workflow correlation
Shared template	Same prompt or agent recipe reused widely	Work-unit registry	Repeated exposure pattern
Shared vendor	Same provider in many deployments	Dependency map	Upstream concentration
Shared sector	Same use case across one industry	Portfolio grouping	Sector-level shock

Discussion

Cyber accumulation sources are the best public analogy in the current source base. They help explain the concern, but they do not prove a direct actuarial model for agentic AI. [SRC: CYB-02] [SRC: CYB-03] [SRC: CYB-04]

[SYNTHESIS: Jearon Wong]

Boundary

- Not an actuarial model.
- Not a quantified accumulation study.
- Not a vendor ranking.

Chapter 15

Agentic Insurability Readiness Model

AIRM is readiness vocabulary, not certification, insurer acceptance, an actuarial score, a coverage guarantee, claims approval guidance, or a procurement benchmark. It is a Jearon Wong synthesis. [SYNTHESIS: Jearon Wong] [INT: INT-07]

AIRM Matrix

Level	Label	Evidence visibility	Insurance reading	Boundary
L0	Uninsurable Black Box	Very low	Hard to review	Analytical label only
L1	Logged but Not Attributable	Low	Events visible, ownership weak	Logs are not enough
L2	Bounded but Weakly Reconstructable	Medium	Some scope visible, chain incomplete	Not underwriting-ready
L3	Evidence-Linked and Claim-Reviewable	Good	Reviewable after loss	Not claim-approved
L4	Underwriting-Ready Lifecycle System	High	Pre-loss evidence is strong	Not insurer acceptance
L5	Dispute-Ready Risk Transfer Architecture	Very high	Review and dispute readiness are robust	Not certification

Discussion

The readiness ladder lets the paper talk about maturity without pretending there is a certified industry scale. It can say what is more visible, more reviewable, and more disputable, without claiming a guarantee. [SRC: CLAIM-01]

[SRC: CYB-02] [INT: INT-07]

WP1 / WP2 Bridge

- WP1 ALCS supports the lifecycle-conformance side.
- WP2 AARM becomes the source shape for the readiness translation.
- AIRM should remain clearly separate from any certification logic. [INT: INT-02] [INT: INT-04] [INT: INT-07]

Boundary

- Not a certification.
- Not a benchmark.
- Not insurer acceptance.

Chapter 16

AIIRWP-2026-v0.1-R6-CANDIDATE

jearonwong.com

Copyright © 2026 Jearon Wong. All rights reserved.

Conclusion: From Agentic Deployment to Insurable Risk Transfer

Agentic AI becomes more insurable not because it becomes more intelligent, but because its work becomes bounded, evidenced, attributable, reconstructable, remediated, and disputable.

[SRC: INS-01] [SRC: CLAIM-01] [SRC: CLAIM-02] [SRC: CLAIM-03] [SRC: MKT-08] [INT: INT-01] [INT: INT-04] [INT: INT-05] [INT: INT-06]

Closing Argument

- The insured subject must be clear.
- The insurable object must be bounded.
- The responsibility path must be visible.
- The loss event must be reconstructable.
- The evidence package must be reviewable.
- The boundary must be legible.

WP4 will later translate compliance, auditability, and insurability into enterprise implementation synthesis. This chapter should stop there and not promise final or guaranteed transfer.

[SYNTHESIS: Jearon Wong]

Boundary

- Not a public release claim.
- Not a final or sealed claim.
- Not insurer acceptance.
- Not a coverage promise.

Appendix A

Agentic Insurability Object Checklist

This appendix is an internal candidate template. It is illustrative only and not operational, normative, or insurer-approved.

Checklist Skeleton

Item	Question	Evidence pointer	AIO
Subject	Who is the legal subject?	Policy / entity record	AIO-01
Work unit	What exact work was done?	Work-unit registry	AIO-02
Authority	Was it within scope?	Authority record	AIO-05
Loss	What happened?	Incident record	AIO-06
Remediation	What was fixed?	Closure record	AIO-10
Dependency	What depended on what?	Dependency map	AIO-11
Boundary	What may be excluded or disputed?	Boundary note	AIO-04 / AIO-12

Boundary Notes

- Do not imply legal sufficiency.
- Do not imply claims approval.
- Do not imply underwriting standard.

Appendix B

Underwriting Evidence Request Template

This appendix is a draft request structure, not a required insurer form.

Template Skeleton

Evidence category	Prompt	AIO	Boundary
Subject	Who is the insured party?	AIO-01	Not legal advice
Scope	What work units are in scope?	AIO-02	Not a policy form
Authority	What can the agent do?	AIO-05	Not coverage authority
Dependencies	Which vendors and tools are involved?	AIO-11	No vendor ranking
Controls	What controls and review points exist?	AIO-08	Not sufficient by itself
Aggregation	Where can correlated loss arise?	AIO-13	Not actuarial proof

Appendix C

Claims Reconstruction Evidence Package

This appendix is an internal package template for review readiness, not a claims approval kit.

Package Skeleton

1. Loss event record.
2. Causality trace.
3. Responsibility map.
4. Boundary / exclusion context.
5. Remediation record.
6. Dispute-ready summary.

Table

Package element	What it answers	AIO
Event record	What happened?	AIO-06
Trace	How did it happen?	AIO-07
Responsibility map	Who owned it?	AIO-03
Boundary map	What context mattered?	AIO-04 / AIO-12
Remediation	How was it fixed?	AIO-10
Review package	How can it be disputed or reviewed?	AIO-14

Boundary Notes

- Not a guarantee of payment.
- Not legal causation proof.
- Not legal closure.

Appendix D

AIO-to-MRO Mapping

This appendix maps this paper's object language back to WP1 source truth.

AIO	WP1 / GAIC relation	Why it matters	Boundary
AIO-01	MRO subject / role logic	Anchors the insured party	Internal translation only
AIO-02	MRO work-unit logic	Binds exposure to bounded action	Not a standard
AIO-03	MRO responsibility logic	Connects action back to role	Not liability finding
AIO-04	MRO boundary logic	Shows what is in or out of scope	Not coverage opinion
AIO-05	MRO authority logic	Distinguishes permission from power	Not coverage authority
AIO-06 to AIO-10	MRO event / remediation / accepted outcome logic	Supports reconstruction	Not claim approval
AIO-11 to AIO-14	MRO dependency / aggregation / dispute logic	Supports review and dispute readiness	Not vendor ranking

Appendix E

AIO-to-Audit-Evidence-Chain Mapping

This appendix maps this paper's object language back to WP2 evidence logic.

AIO	WP2 relation	Why it matters	Boundary
AIO-06	Audit event record	Shows what happened	Not proof of claim
AIO-07	Evidence reconstruction trace	Shows sequence	Not legal causation
AIO-08	Control failure evidence	Shows failed control point	Not negligence finding
AIO-09	Audit evidence chain	Packages evidence for review	Not claim approval
AIO-10	Remediation closure	Shows response and recovery	Not legal closure
AIO-14	Dispute-ready package	Supports review and challenge	Not guaranteed payment

Appendix F

AIRM Readiness Matrix

This appendix is the long-form version of the readiness ladder.

Visibility			
L0	None	AIIRWP-2026-v0.1-R6-CANDIDATE	jearonwong.com
L1	Logs only		Copyright © 2026 Jearon Wong. All rights reserved.
L2	Partial boundary		
L3	Reviewable chain		
L4	Strong pre-loss		
L5	Dispute-ready		

Attribution			
L0	None		
L1	Weak		
L2	Partial		
L3	Good		
L4	Strong		
L5	Very strong		

Reconstruction			
L0	None		
L1	Weak		
L2	Partial		
L3	Good		
L4	Strong		
L5	Very strong		

Underwriting use			
L0	None		
L1	Minimal		
L2	Cautious		
L3	Possible		
L4	Better		
L5	Best available		

L0	None
L1	Minimal
L2	Cautious
L3	Reviewable
L4	Strong
L5	Dispute-ready

Boundary risk

L0	Extreme
L1	High
L2	Medium
L3	Lower
L4	Lower
L5	Lowest

Boundary Notes

- Not a certification.
- Not a benchmark.
- Not insurer acceptance.

Appendix G

AIIRWP-2026-v0.1-R6-CANDIDATE

jearonwong.com

Copyright © 2026 Jearon Wong. All rights reserved.

Boundary and Non-Claim Language

This appendix is the safety rail for the internal candidate.

Required Safe Phrases

- market signals suggest
- some sources indicate
- in certain policy contexts
- may be treated as
- does not by itself establish
- should be understood as an analytical object, not a standard
- author synthesis
- not a coverage opinion
- not an underwriting standard
- not claims approval guidance
- technical trace, not claim evidence by itself
- tool permission, not coverage authority
- necessary layer, not sufficient insurability layer

Forbidden Phrases

- The insurance industry believes
 - AI agents are insurable
 - AI agents are uninsurable
 - Coverage applies
 - Claims will be paid
 - Insurers will accept
 - This is underwriting-ready
 - This proves legal liability
 - MPLP makes agentic AI insurable
 - Validation Lab certifies insurability
-

Source Register and Citation Notes

This section preserves internal citation review context for R6 artifact generation. It is not a public footnote system and does not create a public source registry.

AIIRWP R5 Source Register Finalization Plan

Status: Internal source-register finalization plan.

Boundary: This plan does not create final public footnotes, public source notes, public route metadata, or publication claims.

Final Source Use Rules

- Insurance basics sources support subject, policy, limits, D&O, E&O, and cyber terminology only; they do not support coverage opinions.
- AI insurance market sources support market signals, product examples, broker/reinsurer framing, and fragmentation language only; they do not prove industry consensus.
- Cyber accumulation sources support analogy and risk framing; they do not provide direct actuarial proof for agentic AI.
- Technical framework sources support technical capabilities only; they do not establish insurance evidence, claim evidence, legal authority, or underwriting sufficiency.
- Internal WP1/WP2/WP3 source truth supports framework translation and author synthesis, not external validation.
- AIO and AIRM remain Jearon Wong synthesis / analytical object models.

Final Source ID List

Source tier	Tier 1
Source role	Primary evidence
Chapter usage	0, 4, 5, 12
Access caveat	Public PDF accessible.
Citation risk	State adoption varies; model bulletin is not universal law.

INS-02

Source tier	Tier 1
Source role	Supporting context
Chapter usage	0, 4
Access caveat	Public page accessible.
Citation risk	Press release; use bulletin for primary detail.

INS-03

Source tier	Tier 1
Source role	Boundary reference
Chapter usage	4, 5, 12
Access caveat	Public PDF accessible.
Citation risk	Adoption list can change; verify again before publication.

INS-04

Source tier	Tier 1
Source role	Primary evidence for insurance terminology
Chapter usage	1, 9
Access caveat	Public page accessible.
Citation risk	General industry definition, not policy-specific advice.

INS-05

Source tier	Tier 1
Source role	Primary evidence for insurance terminology
Chapter usage	4, 9, 13
Access caveat	Public page accessible.
Citation risk	General definition; not policy interpretation.

Source tier	Tier 1
Source role	Legal/industry terminology
Chapter usage	1
Access caveat	Public page accessible.
Citation risk	U.S.-oriented legal encyclopedia; use only for plain-English basics.

INS-07

Source tier	Tier 1
Source role	Primary terminology
Chapter usage	3, 4, 9, 13
Access caveat	Public page accessible.
Citation risk	Legal education source, not coverage advice.

INS-08

Source tier	Tier 1
Source role	Primary terminology
Chapter usage	1, 3, 4, 9
Access caveat	Public page accessible.
Citation risk	General education source; not policy advice.

INS-09

Source tier	Tier 1
Source role	Primary context
Chapter usage	4, 5, 12, 13
Access caveat	Public page accessible.
Citation risk	Small-business oriented; not all enterprise forms.

INS-10

Source tier	Tier 1
Source role	Broker context
Chapter usage	4, 5, 12
Access caveat	Public page accessible.
Citation risk	Broker explanatory material; not market consensus.

Source tier	Tier 1
Source role	AI governance reference
Chapter usage	5, 6, 12, 15
Access caveat	Public PDF accessible.
Citation risk	Not insurance-specific.

AIRWIP-2026-v0.1-R6-CANDIDATE

jearonwong.com

Copyright © 2026 Jearon Wong. All rights reserved.

AI-02

Source tier	Tier 1
Source role	Regulation context
Chapter usage	4, 5, 6, 9, 12
Access caveat	Public text accessible.
Citation risk	Do not convert into insurance advice or coverage conclusions.

MKT-01

Source tier	Tier 1
Source role	Primary market evidence
Chapter usage	4, 5, 12
Access caveat	Public page accessible.
Citation risk	Marketing/product page; do not infer broad market consensus.

MKT-02

Source tier	Tier 1
Source role	Primary insurer/reinsurer report
Chapter usage	4, 5, 9, 12
Access caveat	Public PDF accessible.
Citation risk	Reinsurer-authored; not industry standard.

MKT-03

Source tier	Tier 1
Source role	Primary market evidence
Chapter usage	4, 5, 9
Access caveat	Public page may block automated access; verify manually before final citation.
Citation risk	Insurer product announcement; not consensus or policy wording.

MKT-04

Source tier	Tier 1
Source role	Risk engineering context
Chapter usage	4, 5, 7, 11, 12
Access caveat	Public PDF accessible through QBE link.
Citation risk	Risk engineering guidance, not claims evidence by itself.

GenAI Insurability & Risk Transfer White Paper 2026 | Internal Candidate Artifact

AIRWIP-2026-v0.1-R6-CANDIDATE

jearonwong.com

Copyright © 2026 Jearon Wong. All rights reserved.

MKT-05

Source tier	Tier 2
Source role	Broker thought leadership
Chapter usage	4, 5, 9, 14
Access caveat	Public page may rate-limit automated checks.
Citation risk	Thought-leadership format; verify final page text before publication.

MKT-06

Source tier	Tier 1
Source role	Broker report
Chapter usage	4, 5, 9, 12
Access caveat	Public PDF accessible.
Citation risk	Broker perspective; not policy interpretation.

MKT-07

Source tier	Tier 1
Source role	Broker context
Chapter usage	3, 4, 9
Access caveat	Public page accessible.
Citation risk	Podcast/insight format; use as context, not legal conclusion.

MKT-08

Source tier	Tier 1
Source role	Industry association report
Chapter usage	4, 5, 14, 16
Access caveat	Public page accessible.
Citation risk	Report is GenAI-focused, not all agentic AI.

Source tier	Tier 1
Source role	Insurer market/risk report
Chapter usage	4, 5, 14
Access caveat	Public PDF accessible.
Citation risk	Survey-based risk ranking; do not use as coverage proof.

CYB-02

Source tier	Tier 1
Source role	Primary aggregation/reinsurance context
Chapter usage	5, 11, 14, 15
Access caveat	Public page/PDF accessible.
Citation risk	Cyber-specific; AI mapping is author synthesis.

CYB-03

Source tier	Tier 1
Source role	Marketplace/reinsurance context
Chapter usage	4, 5, 11, 14
Access caveat	Public page accessible.
Citation risk	Press release; use report for detailed claims if cited later.

CYB-04

Source tier	Tier 1
Source role	Reinsurer market report
Chapter usage	4, 5, 14
Access caveat	Public page accessible if final URL remains stable.
Citation risk	Not AI-specific in every section.

CYB-05

Source tier	Tier 1
Source role	Cyber insurer/threat report
Chapter usage	4, 5, 7, 11
Access caveat	Public page accessible; report details may require form.
Citation risk	Insurer/threat-intelligence view; not coverage consensus.

CLAIM-01

Agentic Cyberinsurability & Risk Transfer White Paper 2026 | Internal Candidate Artifact

Source tier	Tier 1
Source role	Incident-response evidence context
Chapter usage	7, 10, 13
Access caveat	Public page/PDF accessible.
Citation risk	Cybersecurity source, not insurance claims standard.

AIRWIP-2026-v0.1-R6-CANDIDATE

jearonwong.com

Copyright © 2026 Jearon Wong. All rights reserved.

CLAIM-02

Source tier	Tier 1
Source role	Incident playbook reference
Chapter usage	7, 10, 13
Access caveat	Public page/PDF accessible.
Citation risk	Federal playbook, not private-sector insurance rule.

CLAIM-03

Source tier	Tier 1
Source role	Disclosure/governance context
Chapter usage	3, 10, 13
Access caveat	Public page/PDF accessible.
Citation risk	Securities disclosure context, not claim approval.

TECH-01

Source tier	Tier 1
Source role	Technical context
Chapter usage	2, 3, 7, 11, 12
Access caveat	Public docs accessible.
Citation risk	Do not infer insurance evidence or coverage authority.

TECH-02

Source tier	Tier 1
Source role	Technical context
Chapter usage	2, 3, 7, 11
Access caveat	Public docs accessible.
Citation risk	Not insurance or governance standard.

TECH-03

Source tier	Tier 1
Source role	Technical context
Chapter usage	2, 3, 11
Access caveat	Public docs accessible.
Citation risk	Not responsibility or insurance evidence by itself.

Aspirin & Insurability & Risk Transfer White Paper 2026 | Internal Candidate Artifact

AIRWIP-2026-v0.1-R6-CANDIDATE

jearonwong.com

Copyright © 2026 Jearon Wong. All rights reserved.

TECH-04

Source tier	Tier 1
Source role	Technical context
Chapter usage	7, 10, 11, 12
Access caveat	Public docs accessible.
Citation risk	Product docs; do not convert to insurance evidence claim.

TECH-05

Source tier	Tier 1
Source role	Technical context
Chapter usage	2, 3, 7, 11
Access caveat	Public docs accessible.
Citation risk	Product docs; not insurance authority.

INT-01

Source tier	Tier 1 internal
Source role	Internal framework source
Chapter usage	0-16
Access caveat	Repo file accessible.
Citation risk	Internal author framework; label as internal source truth.

INT-02

Source tier	Tier 1 internal
Source role	Internal source QA
Chapter usage	6, 8, 12, 15
Access caveat	Repo file accessible.
Citation risk	Internal QA; not external validation.

Source tier	Tier 1 internal
Source role	Internal evidence QA
Chapter usage	6, 7, 8, 11, 12, 13
Access caveat	Repo file accessible.
Citation risk	Internal QA; not third-party acceptance.

INT-04

Source tier	Tier 1 internal
Source role	Internal WP2 source truth
Chapter usage	6, 7, 13, 15
Access caveat	Repo file accessible.
Citation risk	Internal preparatory source register.

INT-05

Source tier	Tier 1 internal
Source role	Internal WP2 source truth
Chapter usage	7, 10, 13
Access caveat	Repo file accessible.
Citation risk	Internal author synthesis with source references.

INT-06

Source tier	Tier 1 internal
Source role	WP3 architecture source
Chapter usage	8, 12, 13, 15
Access caveat	Repo file accessible.
Citation risk	Author synthesis, not external standard.

INT-07

Source tier	Tier 1 internal
Source role	WP3 architecture source
Chapter usage	15
Access caveat	Repo file accessible.
Citation risk	Author synthesis, not certification or insurer acceptance.

Required Pre-Publication Recheck

Before any public artifact generation or public staging, R6/R7 must recheck:

Copyright © 2026 Jearon Wong. All rights reserved.

- all live URLs and redirects;
- current A2A documentation URL, replacing the old `google-a2a.github.io` source path;
- split LangGraph durable execution and persistence URLs;
- any blocked, rate-limited, paywalled, or partial-access sources;
- insurer, broker, reinsurer, product, and startup materials for market-signal framing;
- source-note wording around legal, insurance, underwriting, claims, coverage, and aggregation risk.

Finalization Decision

PASS for R5 preparation. Final public footnote style and public source-note rendering remain R6/R6B tasks after artifact generation is authorized.

AIIRWP R5 Citation Normalization Report

Status: Internal citation/source-note normalization report.

Boundary: This report prepares source normalization only. It does not create final public footnotes, public source notes, or public artifact citations.

Source Marker Style

Candidate source uses the R3/R4 marker model: jearonwong.com Copyright © 2026 Jearon Wong. All rights reserved.

- external claims: [SRC: ID]
- internal framework claims: [INT: ID]
- author synthesis: [SYNTHESIS: Jearon Wong]

Marker Validation

Source marker	Candidate count	Inventory status
Agentic AI Insurability & Risk Transfer White Paper 2026 Internal Candidate Artifact		
AI-01	3	valid
AI-08	1	missing
CLAIM-01	27	valid
AIRWP-2026-v0.1-R6-CANDIDATEjearonwong.comCopyright © 2026 Jearon Wong. All rights reserved.		
CLAIM-02	16	valid
CLAIM-03	13	valid
CYB-01	2	valid
CYB-02	11	valid
CYB-03	8	valid
CYB-04	7	valid
INS-01	12	valid
INS-04	10	valid
INS-05	10	valid
INS-06	9	valid
INS-07	7	valid
INS-08	11	valid
INS-09	3	valid
INS-10	3	valid
INT-01	26	valid
INT-02	3	valid
INT-03	4	valid
INT-04	17	valid
INT-05	33	valid
INT-06	19	valid
INT-07	6	valid
MKT-01	10	valid
MKT-02	6	valid
MKT-03	8	valid
MKT-04	1	valid
MKT-05	9	valid
MKT-06	2	valid
MKT-07	5	valid
MKT-08	12	valid
TECH-01	11	valid
TECH-02	4	valid

Source marker	Candidate count	Inventory status
Agentic AI Insurability & Risk Transfer White Paper 2026 Internal Candidate Artifact		
TECH-o3	2	valid
TECH-o4	9	valid
TECH-o5	3	valid
AIRWP-2026-v0.1-R6-CANDIDATE		
jearonwong.com		
Copyright © 2026 Jearon Wong. All rights reserved.		

Missing marker result: BLOCKER - missing markers: AI-o8

R2A Citation Hygiene Checks

Check	Result	Notes
AIIRWP-2026-v0.1-R6-CANDIDATE	jearonwong.com	Copyright © 2026 Jearon Wong. All rights reserved.
Old A2A URL marker remains	PASS	Candidate source uses [SRC: TECH-03] ; final source notes must use current A2A docs, not old google-a2a.github.io .
LangGraph references split correctly	PASS for R5 plan	Candidate source uses [SRC: TECH-04] ; final source notes must keep durable execution and persistence URLs split.
Old Coalition source used for critical claims	PASS	Candidate body does not use [SRC: CYB-05] .
QBE/WTW/Allianz/OpenAI caveated	PASS	Access-caveated sources are framed as signals or technical context, not sole support.
Insurer/broker/product sources as market signals	PASS	Market framing remains fragmented and conditional.
Technical framework docs used only for capabilities	PASS	Candidate preserves necessary-but-insufficient framing.
AIO/AIRM synthesis preserved	PASS	Candidate uses synthesis markers and boundary language.
Fake citations or unsupported quotes	PASS	No fabricated quotations introduced in R5.

R6/R6B Citation Tasks

- Decide whether source IDs remain visible, become endnotes, or appear in a source-note appendix.
- Recheck all URLs before artifact generation.
- Add access caveat notes for blocked/rate-limited/partial sources.
- Preserve the R2A source accuracy gate when formatting final citations.