

JEARON WONG / AGENTIC LIFECYCLE GOVERNANCE INDUSTRY  
SERIES

# Agentic AI Auditability & Assurance White Paper 2026

A Lifecycle Evidence Guide for Audit,  
Assurance, and Enterprise AI Governance

Jearon Wong - Protocol Architect for the Agent Era

| Authority | Responsibility | Agent Work | Evidence | Closure |
|-----------|----------------|------------|----------|---------|
|-----------|----------------|------------|----------|---------|

|                                                                                                     |                                                                     |                                                                          |
|-----------------------------------------------------------------------------------------------------|---------------------------------------------------------------------|--------------------------------------------------------------------------|
| <b>DOCUMENT ID</b><br>AIAAWP-2026-<br>v0.1-R4-<br>CANDIDATE                                         | <b>VERSION</b><br>v0.1-publication-<br>candidate                    | <b>DATE</b><br>May 18, 2026                                              |
| <b>STATUS</b><br>Public research<br>candidate;<br>deployment<br>verification<br>pending; not final, | <b>SERIES</b><br>Agentic Lifecycle<br>Governance<br>Industry Series | <b>VISUAL SOURCE OF TRUTH</b><br>HTML/PDF<br>primary; DOCX<br>derivative |

not sealed, not  
live

Public research candidate. Deployment verification pending. Not final, not sealed, not live, not legal advice, not an audit standard, not certification, not an assurance opinion, not legal compliance proof, not regulator approval, not procurement guidance, not a vendor ranking, and not an endorsement by any firm or professional body.

**PUBLICATION BOUNDARY**

---

## Publication Candidate Status

This artifact is a public research publication candidate staged for deployment verification. It is not a final publication, sealed artifact, production-live verification record, audit standard, certification, assurance opinion, legal compliance proof, regulator approval, procurement recommendation, vendor ranking, or endorsement.

It applies the GAIC main white paper publication standard to AIAAWP while preserving AIAAWP's public-facing title, subtitle, and document ID.

# Publication Contents

HTML anchors are active in the public research candidate. PDF page-number pinning passed internal final candidate QA; deployment/live verification remains pending.

## Front Matter

1. Relationship to the Agentic Lifecycle Governance Industry Series
2. Executive Thesis
3. Boundary Note
4. Artifact Integrity Note
5. Register

## Chapters 0-8

1. 0. Executive Summary
2. 1. Scope, Audience, and Non-Audit-Standard Boundary
3. 2. Why AI Agent Auditability Is Different from Model Governance
4. 3. The Audit Object Shift: From Model Output to Lifecycle Work
5. 4. Why Logs Are Not Audit Evidence Chains
6. 5. The Agentic Audit Object Model
7. 6. MRO-to-Audit-Evidence Mapping
8. 7. Evidence Request List for Agentic Systems
9. 8. Lifecycle Walkthrough for AI Agent / MAS Work

## Chapters 9-16

1. 9. Evidence Partitioning Across Agents, Tools, Roles, Vendors, and Projects
2. 10. Privacy, Selective Disclosure, and Audit Evidence Minimization
3. 11. Exception, Dispute, and Remediation Closure
4. 12. Third-Party Validation Without Certification Claims
5. 13. Agentic Auditability Readiness Model
6. 14. Enterprise Implementation: CIO / CTO / CCO Readiness
7. 15. How Audit and Assurance Firms Can Use This Framework
8. 16. Conclusion

## Appendices

1. Appendix A — Agentic Audit Evidence Request List
2. Appendix B — Agentic System Walkthrough Template
3. Appendix C — MRO-to-Audit-Evidence Mapping
4. Appendix D — Auditability Readiness Matrix
5. Appendix E — Exception / Remediation Closure Checklist
6. Appendix F — Boundary Language
7. Appendix Citation Map

## Sources

1. Source Note

2. Source Note and Citation Register
3. Source Note and Citation Register
4. Source Use Rules
5. R4 Source Register Result
6. R4 Citation Style
7. Package Citation Map
8. Chapter Citation Map
9. R4 Citation Map Result

---

## Relationship to the Agentic Lifecycle Governance Industry Series

The Agentic Lifecycle Governance Industry Series is fixed as a 3+3 structure: three core industry white papers and three practitioner guides. The Global AI Compliance White Paper 2026 defines the compliance object layer for agentic lifecycle governance. This paper defines the audit evidence object layer. Guide 1 will later translate this paper into audit-ready technical implementation planning. Guide 2 will later translate the Global AI Compliance White Paper 2026 and this paper into compliance operating-model planning. The later insurability white paper will address insurability. This public research candidate stages the auditability white paper for deployment verification; it does not replace those later assets.

---

## Executive Thesis

AI agent auditability cannot be built on logs alone. An agentic system becomes auditable only when its lifecycle work can be reconstructed through responsibility-linked evidence chains. **Agentic AI Auditability** is the ability to reconstruct, test, and evidence agentic lifecycle work across authority, responsibility, tools, outcomes, exceptions, and remediation. This definition is author synthesis in this paper grounded in audit evidence language, AI governance guidance, provenance/logging concepts, privacy/evidence-retention guidance, and GAIC's MRO object layer.

---

## Source Note

Big Four sources are used as market context only. Professional audit and assurance sources support terminology, evidence concepts, internal-control language, and

assurance boundaries. AI governance/public guidance sources support governance context. Evidence/provenance/observability/logging sources support the logs-versus-evidence distinction. Privacy sources support evidence minimization, selective disclosure, and retention tension. GAIC source truth supports MROs, Validation Lab boundaries, RCCS-M/ALCS context, and companion-paper boundaries. Inline source IDs in the candidate source map to the package source register and package citation map.

---

## Boundary Note

This candidate source is not an audit standard, certification, legal compliance proof, assurance opinion, regulator-approved method, procurement recommendation, vendor ranking, Big Four endorsement, audit body endorsement, MPLP requirement, or MPLP industry-standard claim. AARM is a proposed readiness model, not a score, benchmark, certification, assurance result, or vendor assessment.

---

# 0. Executive Summary

**Chapter purpose:** State the auditability problem, the object shift, the evidence-chain answer, the AARM bridge, and the non-standard boundary. **Reader question:** Why does agentic AI need a new auditability frame?

---

## Key Claims

- AI agent auditability cannot be built on logs alone.
- Agentic systems become auditable only when lifecycle work can be reconstructed through responsibility-linked evidence chains.
- Auditability shifts the review object from model output alone to lifecycle work across authority, responsibility, tools, outcomes, exceptions, privacy treatment, and remediation.
- GAIC MROs can be translated into audit evidence objects and evidence requests.
- AARM is a proposed readiness model for describing auditability readiness; it does not issue assurance, certification, or legal compliance proof.

---

## Source Grounding Note

Market context is grounded in BF-01, BF-02, and BF-05. Audit evidence language is grounded in AUD-01 and AUD-02. AI governance context is grounded in AI-01 and AI-08. Observability context is grounded in EVID-02. MROs, Validation Lab boundaries, and companion-paper boundaries are GAIC-derived source truth.

---

## Author Synthesis Note

Agentic AI Auditability, Agentic Audit Object, Audit Evidence Chain, lifecycle-responsibility-linked agent work, and AARM are author synthesis in this paper. They are not presented as externally issued audit standards or certification schemes.

---

## Main Text

AI agent auditability cannot be built on logs alone. Logs matter. Traces matter. Workflow histories, tool-call records, telemetry, and incident records all help reconstruct technical activity. But agentic auditability asks a different question: not merely what happened inside a system, but whether the work can be reviewed as lifecycle responsibility.

The shift is subtle and material. Traditional AI governance often begins with the model, the dataset, the output, the control environment, and the policy framework. Those layers remain necessary.

An agentic system, however, does not only produce outputs. It delegates work, calls tools, transfers tasks across agents, accepts or disputes outcomes, escalates exceptions, stores evidence, crosses vendor or processor boundaries, and may require remediation. A log can record that a tool call occurred. By itself, it usually does not establish who authorized the action, which human role owned the work, whether the agent acted inside a delegated authority boundary, whether the outcome was accepted, what exception occurred, or whether remediation closed.

This paper defines **Agentic AI Auditability** as the ability to reconstruct, test, and evidence agentic lifecycle work across authority, responsibility, tools, outcomes, exceptions, and remediation. This is an authored definition. It is grounded in professional audit evidence language, AI governance guidance, provenance concepts, observability/logging documentation, privacy guidance, and GAIC's existing lifecycle responsibility object layer.

It is not a legal definition, audit standard, assurance engagement requirement, certification criterion, or regulator-approved compliance test.

The central contribution of this paper is the Audit Evidence Object layer. The Global AI Compliance White Paper 2026 defined Missing Regulatory Objects for agentic lifecycle governance. This paper translates those objects into auditability terms: evidence requests, walkthrough structures, responsibility maps, evidence partitioning, privacy-aware disclosure, third-party review boundaries, and readiness levels.

The Agentic Audit Object is the unit of review. The Audit Evidence Chain is the responsibility-linked structure that connects lifecycle work to authority, role, tool action, evidence pointer, accepted outcome, exception state, and closure state.

This matters because audit and assurance readers increasingly face AI systems whose behavior cannot be understood through model cards, output samples, policy attestations, or raw telemetry alone. Big Four and professional sources already discuss AI assurance, trusted AI, responsible AI, audit transformation, governance, risk, controls, and agentic AI adoption. Those discussions create the market context for this paper.

This paper does not claim that any firm endorses this framework. Instead, it offers an object-model layer adjacent to those discussions: a disciplined way to ask what must exist before agentic work can be reconstructed.

The paper's recurring distinction is simple:

- Logs record activity.
- Observability explains system behavior.
- Audit evidence supports responsibility review.
- Responsibility-linked evidence chains make agentic lifecycle work reconstructable.

The distinction does not make logs unimportant. Logs and traces are often evidence ingredients. They may identify time, sequence, tool, service, user, or execution path. But audit evidence requires more than event capture. It requires relevance, reliability, sufficiency in context, relationship to review objective, and linkage to responsibility. For

agentic systems, that means linking activity to lifecycle work units, human roles, agent roles, delegated authority, accepted outcomes, exception handling, privacy treatment, and remediation closure.

This paper also introduces the **Agentic Auditability Readiness Model (AARM)**. AARM describes whether an agentic system is unobservable, log-visible, trace-linked, evidence-structured, auditability-ready, or assurance-ready for planning purposes. AARM is deliberately bounded. It is not a score, benchmark, certification, audit opinion, assurance conclusion, legal compliance proof, regulator approval, procurement recommendation, or vendor ranking. Its purpose is to help organizations and reviewers discuss readiness before they overstate assurance.

For enterprises, this paper prepares the path to two practitioner guides. Guide 1 will translate auditability into technical architecture for CIO, CTO, platform, security, and engineering teams. Guide 2 will translate auditability into compliance operating-model design for CCO, governance, legal, risk, internal audit, and policy-to-evidence teams. This paper defines what audit-ready means; the later guides will address how to build and govern toward it.

For audit and assurance firms, this framework can be used as a discussion and readiness structure. It can support evidence request framing, walkthrough scoping, object-model review, and internal method conversations. It does not replace professional methodology, professional judgment, engagement acceptance, independence rules, attestation criteria, assurance standards, or audit procedures. The paper is useful only if its boundary is kept visible.

Executive Orientation Table

| Executive Orientation Table |                                                                                                                                                  |                                                                                                            |                                                       |
|-----------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------|-------------------------------------------------------|
| Paper element               | What it means                                                                                                                                    | Source / synthesis status                                                                                  | Boundary                                              |
| Agentic AI Auditability     | Ability to reconstruct, test, and evidence agentic lifecycle work across authority, responsibility, tools, outcomes, exceptions, and remediation | author synthesis in this paper grounded by AUD-01, AUD-02, AI-01, AI-08, EVID-01, EVID-02, and GAIC-SOURCE | Not audit standard, legal proof, or assurance opinion |
| Agentic Audit Object        | Reviewable lifecycle work object for agentic systems                                                                                             | author synthesis in this paper derived from GAIC MROs and provenance/audit evidence concepts               | Not mandatory schema                                  |
|                             |                                                                                                                                                  |                                                                                                            |                                                       |

| Paper element           | What it means                                                                                                                         | Source / synthesis status                                                      | Boundary                                      |
|-------------------------|---------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------|-----------------------------------------------|
| Audit Evidence Chain    | Responsibility-linked chain connecting work unit, authority, role, tool, evidence, outcome, exception, privacy treatment, and closure | author synthesis in this paper                                                 | Not a certification criterion                 |
| MRO-to-evidence mapping | Translation of GAIC MROs into audit evidence objects and evidence requests                                                            | GAIC-derived plus audit/control language                                       | MROs are not legal mandates                   |
| AARM                    | Proposed readiness model for auditability                                                                                             | author synthesis in this paper grounded by R4B/R1 and audit/governance sources | Not score, certification, or assurance result |

Cross-Links

- See Appendix A for evidence request categories.
- See Appendix C for the full MRO-to-audit-evidence mapping.
- See Appendix D for the AARM readiness matrix.
- See Appendix F for boundary language.

Boundary Note

This chapter does not claim this paper is published, final, sealed, regulator-approved, audit-body-endorsed, Big-Four-endorsed, legally sufficient, certifying, or assurance-producing.

---

# 1. Scope, Audience, and Non-Audit-Standard Boundary

**Chapter purpose:** Define who this paper is for, what it does, what it does not do, how sources are used, and how GAIC/MPLP/Validation Lab relate to the draft. **Reader question:** Who should use this paper, and what authority must it not claim?

---

## Key Claims

- This paper is a lifecycle evidence guide and auditability readiness framework.
- This paper is not an audit standard, assurance opinion, certification, legal compliance proof, regulator approval, procurement recommendation, vendor ranking, or Big Four/audit body endorsement.
- Big Four sources are market context only.
- GAIC is the source for MROs, AARM planning, Validation Lab boundary, and companion-paper scope.
- MPLP may be discussed as one optional lifecycle protocol path, never as required or industry-standard.

---

## Source Grounding Note

Assurance and attestation boundaries use AUD-03 and BOUND-03. Certification and conformity-assessment boundaries use BOUND-01 and BOUND-02. Governance role context uses AUD-07. Big Four AI assurance context uses BF-02 and BF-03 only as market context.

---

## Author Synthesis Note

This paper's scope, Audit Evidence Object layer, Agentic Audit Object, Audit Evidence Chain, and AARM are author synthesis. They are draft framework constructs, not externally adopted professional requirements.

---

## Main Text

This paper is written for readers who need to understand whether agentic AI work can be reconstructed. Its primary audience includes audit and assurance professionals, internal audit teams, technology risk teams, AI governance committees, enterprise control owners, CIOs, CTOs, CCOs, CROs, legal/privacy stakeholders, and platform leaders responsible for agentic systems. Its secondary audience includes advisory teams, AI assurance practitioners, risk engineering teams, agent runtime builders, procurement and vendor-risk teams, and standards or policy readers exploring agentic AI governance.

The paper's scope is deliberately narrow. It does not try to restate all model governance, responsible AI, cybersecurity, privacy, or enterprise risk management. Those fields remain necessary. This paper focuses on the lifecycle evidence layer that becomes necessary when AI systems act through agents, tools, memory, delegation, multi-agent handoffs, and remediation workflows. The question is not simply whether a model behaved acceptably. The question is whether the lifecycle work can be reconstructed across authority, responsibility, tools, outcomes, exceptions, and closure.

This paper is best understood as a lifecycle evidence guide. It defines the proposed object layer needed for audit-ready agentic work, maps GAIC MROs into audit evidence objects, proposes evidence request categories, provides walkthrough and partitioning patterns, introduces AARM readiness levels, and preserves boundaries around privacy, third-party validation, and assurance. It is a draft architecture for auditability, not a professional standard.

The source hierarchy matters. Professional audit and assurance sources provide language for evidence, controls, assurance boundaries, internal audit, attestation, and professional limits. AI governance sources provide context for risk management, accountability, monitoring, documentation, and human oversight. Evidence, provenance, observability, and log-management sources help distinguish raw event data from responsibility-linked evidence. Privacy sources ground the tension between reviewability, minimization, selective disclosure, and retention. Big Four sources provide market context: they show that AI assurance, trusted AI, audit transformation, and agentic AI are active enterprise topics. They do not endorse this paper.

GAIC source truth plays a different role. The Global AI Compliance White Paper 2026 created the MRO object layer and the broader Agentic Lifecycle Governance framing. This paper uses those GAIC-derived objects as its internal foundation. When this paper discusses MRO-01 through MRO-16, RCCS-M/ALCS context, the Evidence-Based Validation Pattern, Validation Lab boundary language, or companion-paper sequencing, it is relying on GAIC source truth, not external audit standards.

MPLP should be handled with care. MPLP may be described as one protocol path whose lifecycle responsibility semantics can help express evidence objects. This paper does not require MPLP, does not rank MPLP above other approaches, does not claim MPLP is an industry standard, and does not claim MPLP proves compliance, auditability, assurance readiness, or enterprise readiness. Guide 1 may later discuss implementation paths, but this paper must remain implementation-neutral.

Validation Lab is also boundary-sensitive. GAIC describes Validation Lab as a non-certifying evidence adjudication example. This paper may use that boundary as an example of third-party evidence review without certification. It must not describe Validation Lab as a certification body, conformity assessment body, regulator, audit firm, assurance provider, legal authority, or proof of compliance.

The safest reading of this paper is this: it provides a structured way to ask what evidence must exist before agentic AI work can be meaningfully reviewed. It does not determine whether that evidence is sufficient for a particular audit, assurance engagement, legal analysis, regulatory process, insurance underwriting decision, procurement decision, or board conclusion. Those determinations belong to qualified

professionals operating under their own methods, scopes, independence requirements, criteria, and legal obligations.

## Scope Table

| Scope Table      |                                                                     |                                                               |                                 |
|------------------|---------------------------------------------------------------------|---------------------------------------------------------------|---------------------------------|
| Area             | This paper does                                                     | This paper does not do                                        | Source / boundary grounding     |
| Auditability     | Defines proposed lifecycle evidence objects and readiness questions | Issue audit opinions or audit procedures                      | AUD-01, AUD-02, AUD-03          |
| Assurance        | Frames readiness for possible assurance planning                    | Provide assurance conclusion or engagement standard           | AUD-03, BOUND-03                |
| Certification    | Distinguishes validation from certification                         | Certify systems or define certification criteria              | BOUND-01, BOUND-02              |
| Legal/privacy    | Frames evidence/privacy tension                                     | Provide legal advice or legal compliance proof                | PRIV-01 to PRIV-05, AI-09       |
| Big Four context | Uses public materials as market context                             | Claim endorsement, adoption, or need                          | BF-01 to BF-05                  |
| GAIC/MRO         | Uses MROs as GAIC-derived object layer                              | Claim MROs are law or external standards                      | GAIC-SOURCE                     |
| MPLP             | Treats MPLP as optional protocol path                               | Require MPLP or claim industry-standard status                | GAIC-SOURCE                     |
| Validation Lab   | Gives non-certifying evidence adjudication example                  | Claim certification, assurance opinion, or regulator approval | GAIC-SOURCE, BOUND-01, BOUND-02 |

---

## Cross-Links

- Appendix F contains reusable boundary language.
- Chapter 12 expands the validation/certification boundary.
- Chapter 15 addresses professional use by audit and assurance firms.

---

## Boundary Note

This chapter intentionally repeats non-claim language because this paper's usefulness depends on not overstating its authority.

---

---

## 2. Why AI Agent Auditability Is Different from Model Governance

**Chapter purpose:** Explain why model governance is necessary but insufficient for agentic auditability. **Reader question:** What changes when AI systems act through agents, tools, memory, delegation, and multi-agent workflows?

---

### Key Claims

- Model governance remains necessary for agentic systems, but it does not fully reconstruct lifecycle work.
- Agentic AI introduces delegated authority, tool action, role separation, handoffs, accepted outcomes, exception states, and remediation closure.
- Auditability must therefore move beyond model/output review into lifecycle responsibility review.

---

### Source Grounding Note

Big Four market context appears in BF-01, BF-04, and BF-05. AI governance context appears in AI-01, AI-02, AI-04, AI-05, and AI-08. Internal audit context appears in AUD-06. GAIC source truth provides MRO and RCCS-M/ALCS lifecycle-object context.

---

### Author Synthesis Note

The claim that "model governance is necessary but insufficient" for agentic auditability is author synthesis in this paper grounded in the shift from model/output governance to lifecycle responsibility objects.

---

### Main Text

Model governance is not obsolete. It remains part of responsible AI practice. Organizations still need to understand model purpose, data, evaluation, limitations, monitoring, robustness, bias, security, human oversight, and change management. Public AI governance frameworks and professional guidance continue to provide important language for risk management, accountability, documentation, monitoring, controls, and internal audit review.

The limitation is not that model governance is wrong. The limitation is that agentic systems create work that is not reducible to model behavior. An agent may plan a sequence, invoke a tool, pass work to another agent, store state, retrieve memory, request human confirmation, act under delegated authority, trigger an external

workflow, or remediate an exception. A model output may be only one event in a broader lifecycle.

For auditability, that broader lifecycle matters. If an agent drafts a message, the model output may be reviewed. If an agent sends the message through a CRM tool, updates a customer record, triggers a refund, routes a contract, or changes a production configuration, the relevant review object is no longer only the output. It is the work unit: who initiated it, what authority permitted it, what agent and tool executed it, what evidence was retained, who accepted the result, what exception occurred, and how any remediation closed.

This creates an object gap. Traditional governance may document the model, evaluate outputs, approve use cases, and monitor performance. Agentic auditability must also document authority boundaries, role mappings, tool-action evidence, responsibility transfer, evidence partitioning, privacy treatment, and closure. These are lifecycle objects, not merely model attributes.

The distinction becomes clearer when human oversight is considered. A human-in-the-loop control may be meaningful only if the loop is tied to an identifiable responsibility state. Who was the human? What role did they occupy? What were they asked to confirm? Was the confirmation advisory, blocking, or final? Was it tied to the delegated authority boundary? Did it create an accepted outcome state? Did it leave evidence that another reviewer can reconstruct? Without these objects, "human oversight" may be real operationally but weak evidentially.

Agentic systems also complicate responsibility because agent roles are not human roles. Human-like names, personas, and task labels may help interfaces, but they can blur governance. An "analyst agent" is not the same thing as an analyst. It is a bounded execution role with tool permissions, instructions, constraints, evidence obligations, and escalation paths. A human or organizational role still owns intent, authority, acceptance, and remediation.

The point of this paper is to place model governance beside lifecycle evidence architecture. Model governance asks whether the model and its use are governed. Agentic auditability asks whether the agentic work can be reconstructed. Both are needed. The second cannot be inferred from the first.

|                                                                                                                                                                                                    |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <div>Table 1: Traditional AI Audit vs Agentic AI Auditability</div> <div>Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.</div> |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

**Model documentation**

|                                           |                                                                  |
|-------------------------------------------|------------------------------------------------------------------|
| <b>AGENTIC<br/>AUDITABILITY<br/>FOCUS</b> | Lifecycle work reconstruction                                    |
| <b>MISSING LIFECYCLE<br/>OBJECT</b>       | Agentic Audit Object                                             |
| <b>EVIDENCE NEEDED</b>                    | Work unit ID, scope, evidence chain, accepted outcome            |
| <b>SOURCE /<br/>SYNTHESIS NOTE</b>        | Mixed: AI-01, AI-08, GAIC-SOURCE, author synthesis in this paper |
| <b>BOUNDARY</b>                           | Not a new audit standard                                         |

**Output review**

|                                           |                                                                     |
|-------------------------------------------|---------------------------------------------------------------------|
| <b>AGENTIC<br/>AUDITABILITY<br/>FOCUS</b> | Accepted outcome review                                             |
| <b>MISSING LIFECYCLE<br/>OBJECT</b>       | Accepted outcome state                                              |
| <b>EVIDENCE NEEDED</b>                    | Reviewer role, acceptance/rejection/dispute state, evidence pointer |
| <b>SOURCE /<br/>SYNTHESIS NOTE</b>        | GAIC-derived + AUD-01/AUD-06                                        |
| <b>BOUNDARY</b>                           | Acceptance is not legal proof                                       |

**Monitoring**

|                                           |                                                                   |
|-------------------------------------------|-------------------------------------------------------------------|
| <b>AGENTIC<br/>AUDITABILITY<br/>FOCUS</b> | Authority and exception monitoring                                |
| <b>MISSING LIFECYCLE<br/>OBJECT</b>       | Authority drift record                                            |
| <b>EVIDENCE NEEDED</b>                    | Baseline authority, observed action, escalation, closure evidence |
| <b>SOURCE /<br/>SYNTHESIS NOTE</b>        | AI-02, EVID-04, GAIC-SOURCE                                       |
| <b>BOUNDARY</b>                           | Drift is not automatically a legal breach                         |

**Access control**

|                                   |                                                   |
|-----------------------------------|---------------------------------------------------|
| <b>AGENTIC AUDITABILITY FOCUS</b> | Delegated authority review                        |
| <b>MISSING LIFECYCLE OBJECT</b>   | Delegated authority boundary                      |
| <b>EVIDENCE NEEDED</b>            | Scope, condition, expiry, tool action, escalation |
| <b>SOURCE / SYNTHESIS NOTE</b>    | AI-08, AUD-06, GAIC-SOURCE                        |
| <b>BOUNDARY</b>                   | IAM permission is not full business authority     |

**Control testing**

|                                   |                                                           |
|-----------------------------------|-----------------------------------------------------------|
| <b>AGENTIC AUDITABILITY FOCUS</b> | Lifecycle responsibility walkthrough                      |
| <b>MISSING LIFECYCLE OBJECT</b>   | Responsibility map and evidence chain                     |
| <b>EVIDENCE NEEDED</b>            | Human role, agent role, tool record, partitioned evidence |
| <b>SOURCE / SYNTHESIS NOTE</b>    | AUD-04, AUD-05, EVID-01                                   |
| <b>BOUNDARY</b>                   | Walkthrough is not formal audit procedure                 |

**Table note:** Mixed source-grounded, GAIC-derived, and author-synthesis table. It is not a certification table, audit procedure, procurement tool, or vendor comparison.

**Cross-Links**

- Chapter 3 defines the audit object shift.
- Chapter 5 defines the Agentic Audit Object Model.
- Appendix C maps MROs to audit evidence objects.

**Boundary Note**

This chapter does not attack model governance or claim model governance is useless. It argues that agentic auditability requires an additional lifecycle responsibility layer.

---

## 3. The Audit Object Shift: From Model Output to Lifecycle Work

**Chapter purpose:** Define the proposed review object for agentic auditability. **Reader question:** What exactly should audit stakeholders reconstruct?

---

### Key Claims

- Agentic auditability requires a shift from model/output review to lifecycle work review.
- The Agentic Audit Object is a proposed review object for lifecycle-responsibility-linked agent work.
- Lifecycle-responsibility-linked agent work connects intent, authority, role, agent action, tool action, evidence, outcome, exception, privacy treatment, and closure.

---

### Source Grounding Note

Audit evidence language is grounded in AUD-01 and AUD-02. Governance context is grounded in AI-01 and AI-08. Provenance concepts are grounded in EVID-01. MRO source truth is GAIC-derived.

---

### Author Synthesis Note

Agentic Audit Object and lifecycle-responsibility-linked agent work are author-synthesis constructs. They are not claimed as existing professional standards, legal categories, or regulator-defined objects.

---

### Main Text

An audit object is the thing being reviewed. In many AI governance conversations, the object is a model, dataset, output, use case, control, system, or policy. Those objects remain important. But agentic systems create another review object: lifecycle work.

Lifecycle work is not a single output. It is the chain of actions and decisions by which an agentic system receives intent, acts under authority, uses tools, transfers work, records evidence, produces a result, reaches acceptance or dispute, handles exceptions, and closes remediation. If that chain cannot be reconstructed, the system may be observable but not audit-ready.

This paper calls the proposed review unit the **Agentic Audit Object**. It is a structured representation of lifecycle-responsibility-linked agent work. It is not a mandatory data schema. It is a conceptual object model that helps audit, assurance, governance, and technology teams ask whether the relevant lifecycle work can be reviewed.

The Agentic Audit Object exists because agentic systems blur the boundary between process and output. A model-generated draft may be an output. A tool action that sends the draft, records the transaction, updates a system, triggers a workflow, or changes a downstream state is lifecycle work. The auditability question changes from "Was the output acceptable?" to "Can the lifecycle work that produced and accepted that outcome be reconstructed?"

The object shift also changes the meaning of responsibility. Responsibility cannot be inferred from a username, service account, model name, trace ID, or tool log alone. A responsibility-linked review object should identify the human role responsible for intent, authority, review, acceptance, exception handling, and remediation. It should separately identify the agent role and tool surface that executed the work. It should preserve the difference between technical execution and business responsibility.

Provenance concepts help explain why this structure matters. Provenance language distinguishes entities, activities, agents, and relationships. Audit evidence language asks whether evidence is relevant and reliable for the review objective. this paper combines those ideas with GAIC's MRO layer: lifecycle work must be addressable, linked, partitioned, and reviewable.

The Agentic Audit Object therefore contains more than a log. It contains the work unit, the authority boundary, the human role, the agent role, the tool-action record, the evidence pointer, the accepted outcome state, the exception state, the remediation closure state, and the privacy/selective disclosure profile. Those fields may live in different systems. The model is not saying every enterprise must store them in one database. It says that the relationships must be reconstructable.

This object shift is also where this paper becomes distinct from the Global AI Compliance White Paper 2026. The Global AI Compliance White Paper 2026 identifies the missing regulatory objects for agentic lifecycle governance. This paper asks how those objects become evidence. The same MRO can be read as a governance gap, an engineering object, or an audit evidence object. This paper focuses on the third reading.

### Agentic Audit Object Overview

| Object component    | Review question            | Evidence examples                                 | Source / synthesis status                                          | Boundary             |
|---------------------|----------------------------|---------------------------------------------------|--------------------------------------------------------------------|----------------------|
| Lifecycle work unit | What work is under review? | Work unit ID, scope, task intent, lifecycle stage | author synthesis in this paper grounded by EVID-01 and GAIC-SOURCE | Not mandatory schema |
|                     |                            |                                                   |                                                                    |                      |

| Object component                  | Review question                                             | Evidence examples                                                   | Source / synthesis status               | Boundary                       |
|-----------------------------------|-------------------------------------------------------------|---------------------------------------------------------------------|-----------------------------------------|--------------------------------|
| Authority boundary                | Was the action authorized within scope?                     | Delegation record, condition, expiry, escalation path               | GAIC-derived + AI-01/AI-08              | Not proof of legal delegation  |
| Human role responsibility         | Who owned intent, review, acceptance, or closure?           | Role map, review record, acceptance owner                           | GAIC-derived + AUD-06/AUD-07            | Not legal liability assignment |
| Agent role responsibility surface | What agent role executed or transformed the work?           | Agent role ID, constraints, capability boundary                     | GAIC-derived author synthesis + EVID-01 | Agent is not a legal person    |
| Tool-action evidence              | What external or consequential action occurred?             | Tool call record, affected system, reversibility, rollback evidence | GAIC-derived + EVID-02/EVID-03          | Not legal liability conclusion |
| Accepted outcome                  | Was the result accepted, disputed, rejected, or remediated? | Acceptance state, reviewer role, dispute reason                     | GAIC-derived + AUD-01/AUD-04            | Not compliance proof           |
| Exception / remediation closure   | How was deviation handled and closed?                       | Exception record, corrective action, recheck, closure owner         | GAIC-derived + EVID-04/AUD-04           | Not legal settlement           |

**Table note:** Mixed GAIC-derived and author-synthesis table grounded in audit evidence, provenance, and governance sources. It is not an audit standard or legal template.

## Cross-Links

- Chapter 5 expands the object model fields.
- Chapter 6 maps each MRO into audit evidence objects.

- Appendix B provides a walkthrough template.

---

## Boundary Note

The Agentic Audit Object is a proposed object model in this paper for reviewability. It does not create legal categories, professional requirements, certification criteria, or mandatory implementation schema.

---

---

## 4. Why Logs Are Not Audit Evidence Chains

**Chapter purpose:** Ground the core this paper claim that logs and traces are useful but insufficient for agentic auditability. **Reader question:** What can logs show, and what do they fail to prove?

---

### Key Claims

- Logs, traces, metrics, and workflow histories can support auditability.
- They do not automatically prove authority, responsibility, accepted outcome, exception handling, privacy treatment, or remediation closure.
- Audit Evidence Chains increase evidentiary value by linking activity to lifecycle responsibility metadata.

---

### Source Grounding Note

Audit evidence language is grounded in AUD-01 and AUD-02. Observability and logging context is grounded in EVID-02 and EVID-03. Provenance concepts are grounded in EVID-01. Incident/governance context is grounded in EVID-04. AI Act logging context may be referenced as legal context only through AI-06.

---

### Author Synthesis Note

"Logs are not audit evidence chains" is author synthesis in this paper. It is grounded in the distinction between observability/logging and professional audit evidence concepts.

---

### Main Text

Logs are necessary but not sufficient. That sentence is the safest starting point for agentic auditability. The point is not to diminish logs, traces, metrics, monitoring systems, workflow histories, or observability platforms. They are often the first evidence ingredients available. Without them, reconstruction may collapse into narrative memory.

The problem is that logs usually describe activity, not responsibility. A log may show that a service called a tool at a time. A trace may show that a workflow crossed several components. A metric may show latency, error rate, or throughput. An incident record may show that something failed. Those are useful signals. They do not automatically tell a reviewer whether the action was authorized, which human role owned it, which agent role was acting, what business scope applied, whether the tool action was reversible, whether the outcome was accepted, what privacy treatment applied, or whether remediation closed.

Audit evidence is tied to a review objective. In professional audit language, evidence is evaluated for relevance and reliability within context. For agentic systems, the review objective is not merely to explain system behavior. It is to reconstruct lifecycle responsibility. That means the evidence must connect the technical event to authority, role, tool action, outcome state, exception state, privacy treatment, and closure.

Observability explains how systems behave. Auditability asks whether lifecycle work can be reviewed. The same trace can be valuable for both, but it plays different roles. In observability, the trace helps engineers understand execution path and performance. In auditability, the trace is an input to a larger evidence chain. It must be joined with responsibility records, authority records, evidence pointers, acceptance records, and exception/remediation records.

Consider a tool call. A log may show the tool name, timestamp, endpoint, response code, and service identity. The auditability question is broader: Who authorized the tool action? Was the agent operating inside delegated scope? Which data or system was affected? Was the action reversible? Did a human review occur before or after the action? Was the result accepted, disputed, or remediated? Was sensitive data minimized or redacted in review evidence? Was there a closure state?

The same distinction applies to incident records. An incident log may show that an error occurred. It may not show whether the affected outcome was withdrawn, corrected, re-reviewed, accepted, escalated, or closed. It may not show whether the issue was due to authority drift, responsibility transfer, cross-project reuse, tool-action side effects, privacy leakage, or vendor/runtime substitution. Without linkage to lifecycle objects, the incident is visible but not fully reconstructable.

This paper calls the needed structure an **Audit Evidence Chain**. An Audit Evidence Chain is not just a sequence of logs. It is a responsibility-linked reconstruction path. It ties a lifecycle work unit to the relevant evidence objects: authority boundary, human role, agent role, tool action, evidence pointer, accepted outcome, exception, remediation, privacy treatment, and third-party review boundary where relevant.

This structure also protects against over-retention. A naive response to auditability gaps is to collect everything. That can create privacy and confidentiality risks. A better approach is to preserve evidence pointers, partition keys, disclosure profiles, redaction rules, retention logic, and integrity records so reviewers can reconstruct what they are authorized to inspect without unnecessary exposure.

## Table 2: Logs vs Audit Evidence Chains

Table 2: Logs vs Audit Evidence Chains

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

**Tool-call log**

|                                             |                                                                                 |
|---------------------------------------------|---------------------------------------------------------------------------------|
| <b>WHAT LOGS/TRACES CAN SHOW</b>            | Tool, endpoint, timestamp, service, response                                    |
| <b>WHAT THEY DO NOT AUTOMATICALLY PROVE</b> | Business authority, human owner, reversibility, accepted outcome                |
| <b>REQUIRED LINKAGE</b>                     | Delegated authority record, human role map, tool-action evidence, outcome state |
| <b>RELATED MRO</b>                          | MRO-02, MRO-05, MRO-04                                                          |
| <b>SOURCE / SYNTHESIS NOTE</b>              | EVID-02/EVID-03 + GAIC-SOURCE                                                   |

**Trace span**

|                                             |                                                      |
|---------------------------------------------|------------------------------------------------------|
| <b>WHAT LOGS/TRACES CAN SHOW</b>            | Execution path, service dependencies, timing         |
| <b>WHAT THEY DO NOT AUTOMATICALLY PROVE</b> | Lifecycle responsibility or review sufficiency       |
| <b>REQUIRED LINKAGE</b>                     | Work unit ID, agent role, evidence pointer, role map |
| <b>RELATED MRO</b>                          | MRO-01, MRO-03, MRO-08                               |
| <b>SOURCE / SYNTHESIS NOTE</b>              | EVID-01/EVID-02 + author synthesis in this paper     |

**Workflow state**

|                                             |                                                               |
|---------------------------------------------|---------------------------------------------------------------|
| <b>WHAT LOGS/TRACES CAN SHOW</b>            | Step completion, state transition                             |
| <b>WHAT THEY DO NOT AUTOMATICALLY PROVE</b> | Whether state is a governance acceptance state                |
| <b>REQUIRED LINKAGE</b>                     | Acceptance criteria, reviewer role, dispute/remediation state |
| <b>RELATED MRO</b>                          | MRO-04, MRO-16                                                |
| <b>SOURCE / SYNTHESIS NOTE</b>              | AUD-01/AUD-04 + GAIC-SOURCE                                   |

### Incident log

|                                             |                                                                 |
|---------------------------------------------|-----------------------------------------------------------------|
| <b>WHAT LOGS/TRACES CAN SHOW</b>            | Error event, detection, response action                         |
| <b>WHAT THEY DO NOT AUTOMATICALLY PROVE</b> | Remediation closure or legal fault                              |
| <b>REQUIRED LINKAGE</b>                     | Exception owner, corrective action, recheck, closure acceptance |
| <b>RELATED MRO</b>                          | MRO-07, MRO-16                                                  |
| <b>SOURCE / SYNTHESIS NOTE</b>              | EVID-04/AUD-04                                                  |

### Access log

|                                             |                                                           |
|---------------------------------------------|-----------------------------------------------------------|
| <b>WHAT LOGS/TRACES CAN SHOW</b>            | User/service access event                                 |
| <b>WHAT THEY DO NOT AUTOMATICALLY PROVE</b> | Delegated business authority or scope                     |
| <b>REQUIRED LINKAGE</b>                     | Authority boundary, confirmation rule, risk class, expiry |
| <b>RELATED MRO</b>                          | MRO-02, MRO-07                                            |
| <b>SOURCE / SYNTHESIS NOTE</b>              | AI-01/AI-08 + author synthesis in this paper              |

### Retention log

|                                             |                                                                           |
|---------------------------------------------|---------------------------------------------------------------------------|
| <b>WHAT LOGS/TRACES CAN SHOW</b>            | Data/evidence storage or deletion event                                   |
| <b>WHAT THEY DO NOT AUTOMATICALLY PROVE</b> | Whether retention/disclosure is legally sufficient                        |
| <b>REQUIRED LINKAGE</b>                     | Privacy treatment, minimization rule, retention basis, disclosure profile |
| <b>RELATED MRO</b>                          | MRO-10, MRO-12, MRO-13                                                    |
| <b>SOURCE / SYNTHESIS NOTE</b>              | PRIV-01 to PRIV-05                                                        |

**Table note:** Mixed source-grounded and author-synthesis table. Logs are treated as useful inputs, not as useless artifacts. The table is not a vendor critique or audit procedure.

---

## Cross-Links

- Appendix A lists evidence request categories.
- Chapter 9 addresses evidence partitioning.
- Chapter 10 addresses minimization and selective disclosure.

---

## Boundary Note

This chapter does not claim observability vendors fail, logs are useless, or telemetry can never support audit evidence. It claims only that raw logs do not automatically become responsibility-linked audit evidence chains.

---

---

## 5. The Agentic Audit Object Model

**Chapter purpose:** Define the proposed object model for audit-ready agentic lifecycle work. **Reader question:** What makes lifecycle work audit-ready?

---

### Key Claims

- Audit-ready agentic lifecycle work should be addressable, reconstructable, partitioned, privacy-aware, and reviewable.
- The Agentic Audit Object Model binds lifecycle work unit, authority boundary, human role responsibility, agent role responsibility, tool-action evidence, evidence pointer, accepted outcome, exception state, remediation closure, and privacy treatment.
- The model is a proposed review architecture, not a mandatory schema or standard.

---

### Source Grounding Note

The model uses audit evidence language from AUD-01, control vocabulary from AUD-04, AI governance context from AI-01, and provenance structure from EVID-01. MRO dependencies are GAIC-derived.

---

### Author Synthesis Note

The Agentic Audit Object Model is author synthesis in this paper. It is a conceptual object model for auditability, not a required implementation format.

---

### Main Text

The Agentic Audit Object Model translates lifecycle work into a reviewable structure. It does not require that every enterprise use the same database table, evidence pack format, or protocol. It asks whether the relationships needed for review can be reconstructed.

The first requirement is an addressable lifecycle work unit. A reviewer needs to know what work is under review: the task, workflow, agent run, delegated action, tool call, handoff, or remediation event. Without a stable work unit, evidence remains scattered across logs, tickets, approvals, messages, policies, and memories.

The second requirement is an authority boundary. Agentic systems often combine technical permission with business action. Technical permission answers whether a service can call a tool. Business authority answers whether that action was allowed

under scope, condition, risk class, and responsibility. Auditability requires evidence of the authority boundary, not merely evidence that the tool call succeeded.

The third requirement is role mapping. Human role responsibility and agent role responsibility surface must be separated. The human role owns intent, approval, review, acceptance, escalation, or closure. The agent role describes bounded execution capability, constraints, tool access, and evidence obligations. Confusing these roles weakens accountability and makes review harder.

The fourth requirement is tool-action evidence. A tool action is often where AI-generated output becomes operational consequence. The evidence object should identify the tool, action, target system, affected data or process, initiator, authority basis, reversibility, rollback path, and owner. This does not assign legal liability. It creates reviewability.

The fifth requirement is an evidence pointer. Auditability does not always require raw evidence to be duplicated into a central repository. It may require pointers to evidence, integrity hashes, redaction profiles, retention rules, and partition keys. The pointer must be strong enough that authorized reviewers can reconstruct the chain.

The sixth requirement is an accepted outcome state. An output is not the same as an accepted outcome. The model should distinguish produced, reviewed, accepted, rejected, disputed, remediated, and closed states. This is especially important when a downstream process treats agent work as completed before a responsible human or governance workflow has accepted it.

The seventh requirement is exception and remediation closure. Auditability is weak if exceptions disappear into operational tickets without lifecycle linkage. Exceptions should connect to authority, role, affected outcome, corrective action, recheck, closure owner, and reopen criteria.

The eighth requirement is privacy treatment. Evidence is not improved by unlimited collection. Reviewability must be balanced with minimization, selective disclosure, redaction, retention, and access control. The object model should identify which evidence can be disclosed, to whom, for what review purpose, and under what boundary.

Taken together, these fields create the Audit Evidence Chain. The chain does not replace professional judgment about sufficiency. It makes judgment possible by giving reviewers a reconstructable path.

## Agentic Audit Object Model Table

### Agentic Audit Object Model Table

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

**Lifecycle work unit ID**

|                                  |                                                    |
|----------------------------------|----------------------------------------------------|
| <b>PURPOSE</b>                   | Names the work under review                        |
| <b>EXAMPLE EVIDENCE</b>          | Workflow/run/task ID, scope, intent                |
| <b>RELATED MRO</b>               | MRO-08                                             |
| <b>SOURCE / SYNTHESIS STATUS</b> | author synthesis in this paper grounded in EVID-01 |
| <b>BOUNDARY</b>                  | Not mandatory schema                               |

**Authority boundary**

|                                  |                                                       |
|----------------------------------|-------------------------------------------------------|
| <b>PURPOSE</b>                   | Records delegated scope and limits                    |
| <b>EXAMPLE EVIDENCE</b>          | Delegation record, condition, expiry, escalation rule |
| <b>RELATED MRO</b>               | MRO-02, MRO-07                                        |
| <b>SOURCE / SYNTHESIS STATUS</b> | GAIC-derived + AI-01/AI-08                            |
| <b>BOUNDARY</b>                  | Not legal delegation proof                            |

**Human role responsibility**

|                                  |                                                             |
|----------------------------------|-------------------------------------------------------------|
| <b>PURPOSE</b>                   | Maps human ownership                                        |
| <b>EXAMPLE EVIDENCE</b>          | Intent owner, review owner, acceptance owner, closure owner |
| <b>RELATED MRO</b>               | MRO-01                                                      |
| <b>SOURCE / SYNTHESIS STATUS</b> | GAIC-derived + AUD-06/AUD-07                                |
| <b>BOUNDARY</b>                  | Not legal liability assignment                              |

**Agent role responsibility surface**

|                                  |                                                               |
|----------------------------------|---------------------------------------------------------------|
| <b>PURPOSE</b>                   | Defines bounded agent function                                |
| <b>EXAMPLE EVIDENCE</b>          | Agent role, constraints, capability boundary, evidence duties |
| <b>RELATED MRO</b>               | MRO-03, MRO-06                                                |
| <b>SOURCE / SYNTHESIS STATUS</b> | GAIC-derived author synthesis + EVID-01                       |
| <b>BOUNDARY</b>                  | Agent is not human role                                       |

**Tool-action evidence**

|                                  |                                                      |
|----------------------------------|------------------------------------------------------|
| <b>PURPOSE</b>                   | Records consequential action                         |
| <b>EXAMPLE EVIDENCE</b>          | Tool ID, target system, affected data, rollback path |
| <b>RELATED MRO</b>               | MRO-05                                               |
| <b>SOURCE / SYNTHESIS STATUS</b> | GAIC-derived + EVID-02/EVID-03                       |
| <b>BOUNDARY</b>                  | Not liability conclusion                             |

**Evidence pointer**

|                                  |                                                   |
|----------------------------------|---------------------------------------------------|
| <b>PURPOSE</b>                   | Links to review evidence                          |
| <b>EXAMPLE EVIDENCE</b>          | URI, hash, evidence pack reference, partition key |
| <b>RELATED MRO</b>               | MRO-08, MRO-12                                    |
| <b>SOURCE / SYNTHESIS STATUS</b> | Mixed                                             |
| <b>BOUNDARY</b>                  | Not blanket disclosure                            |

**Accepted outcome**

|                                  |                                              |
|----------------------------------|----------------------------------------------|
| <b>PURPOSE</b>                   | Records outcome governance state             |
| <b>EXAMPLE EVIDENCE</b>          | Accepted/rejected/disputed/remediated status |
| <b>RELATED MRO</b>               | MRO-04                                       |
| <b>SOURCE / SYNTHESIS STATUS</b> | GAIC-derived + AUD-01/AUD-04                 |
| <b>BOUNDARY</b>                  | Not compliance proof                         |

**Exception state**

|                                  |                                               |
|----------------------------------|-----------------------------------------------|
| <b>PURPOSE</b>                   | Records deviation or dispute                  |
| <b>EXAMPLE EVIDENCE</b>          | Exception ID, trigger, owner, escalation path |
| <b>RELATED MRO</b>               | MRO-07, MRO-16                                |
| <b>SOURCE / SYNTHESIS STATUS</b> | GAIC-derived + EVID-04                        |
| <b>BOUNDARY</b>                  | Not legal breach finding                      |

|                                  |                                                               |
|----------------------------------|---------------------------------------------------------------|
| <b>Remediation closure</b>       |                                                               |
| <b>PURPOSE</b>                   | Records corrective action and closure                         |
| <b>EXAMPLE EVIDENCE</b>          | Corrective action, recheck, closure reviewer, reopen criteria |
| <b>RELATED MRO</b>               | MRO-16                                                        |
| <b>SOURCE / SYNTHESIS STATUS</b> | GAIC-derived + AUD-04/AUD-05                                  |
| <b>BOUNDARY</b>                  | Not settlement or regulator closure                           |

|                                  |                                                 |
|----------------------------------|-------------------------------------------------|
| <b>Privacy treatment</b>         |                                                 |
| <b>PURPOSE</b>                   | Defines minimization and disclosure             |
| <b>EXAMPLE EVIDENCE</b>          | Redaction profile, retention rule, access scope |
| <b>RELATED MRO</b>               | MRO-10 to MRO-13                                |
| <b>SOURCE / SYNTHESIS STATUS</b> | PRIV sources + GAIC synthesis                   |
| <b>BOUNDARY</b>                  | Not legal advice                                |

**Table note:** This table is mixed source-grounded, GAIC-derived, and author-synthesis. It is a drafting model, not an audit standard, certification checklist, or legal template.

Cross-Links

- Chapter 6 maps the model to all MROs.
- Appendix A converts fields into evidence requests.
- Appendix B turns fields into a walkthrough template.

Boundary Note

The Agentic Audit Object Model should be read as a review architecture. It does not require MPLP, Cognitive OS, Validation Lab, or any particular vendor system.

---

## 6. MRO-to-Audit-Evidence Mapping

**Chapter purpose:** Translate GAIC MROs into audit evidence objects and evidence requests. **Reader question:** How does each MRO become evidence?

---

### Key Claims

- GAIC MROs can be mapped into reviewable audit evidence objects.
- The mapping turns lifecycle governance objects into evidence request categories.
- MROs remain GAIC-derived governance objects; they are not legal mandates, audit standards, or certification criteria.

---

### Source Grounding Note

MRO names and meanings use GAIC source truth. Audit evidence language uses AUD-01. AI audit/control context uses AUD-05 and AI-01.

---

### Author Synthesis Note

The MRO-to-audit-evidence mapping is author synthesis in this paper derived from GAIC objects and externally sourced audit/control terminology.

---

### Main Text

The Global AI Compliance White Paper 2026 introduced the sixteen Missing Regulatory Objects for agentic and multi-agent system governance. This paper reads those objects through an auditability lens. Each MRO asks a governance question. This paper adds the evidence question: what object would a reviewer need to inspect to reconstruct the lifecycle work?

This translation is not a claim that MROs are law. It is not a claim that the presence of an MRO proves compliance. It is not a claim that absence proves non-compliance. MROs are GAIC-derived lifecycle responsibility objects. Their value for this paper is that they give auditability a vocabulary for the work that logs alone do not capture.

The mapping also prevents a common failure in agentic governance: treating "evidence" as a flat export. Agentic systems do not need a pile of logs; they need evidence structured around review questions. Who owned the action? What authority applied? Which agent role acted? What tool created consequence? Was the outcome accepted? What exception occurred? What data was exposed? What changed after vendor, model, or runtime substitution? How did remediation close?

The sixteen MROs can be grouped into six auditability clusters:

- 1. Responsibility and role mapping: MRO-01, MRO-03, MRO-06.
- 2. Authority and action boundaries: MRO-02, MRO-05, MRO-07.
- 3. Outcome and closure: MRO-04, MRO-16.
- 4. Evidence partitioning and reuse: MRO-08, MRO-09.
- 5. Privacy and disclosure: MRO-10, MRO-11, MRO-12, MRO-13.
- 6. Vendor, processor, and substitution chains: MRO-14, MRO-15.

This cluster view helps readers see that auditability is not one control. It is the ability to reconstruct a lifecycle. A work unit can fail auditability at many points: no role map, no authority record, no accepted outcome state, no evidence partition, no privacy treatment, no substitution record, or no remediation closure.

Table 3: MRO-to-Audit-Evidence Mapping

Table 3: MRO-to-Audit-Evidence Mapping

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

MRO-01 Human Role to MAS Responsibility Mapping

|                         |                                                                                                 |
|-------------------------|-------------------------------------------------------------------------------------------------|
| AUDIT EVIDENCE OBJECT   | Human-role responsibility map                                                                   |
| SAMPLE EVIDENCE REQUEST | Show accountable role for intent, authorization, review, acceptance, exception, and remediation |
| SOURCE SUPPORT          | AUD-06, AUD-07, AI-01                                                                           |
| GUIDE 1 RELATION        | Role/evidence architecture                                                                      |
| GUIDE 2 RELATION        | Responsibility governance                                                                       |
| BOUNDARY                | No legal liability assignment                                                                   |

MRO-02 Delegated Authority Boundary

|                         |                                                                         |
|-------------------------|-------------------------------------------------------------------------|
| AUDIT EVIDENCE OBJECT   | Delegated authority record                                              |
| SAMPLE EVIDENCE REQUEST | Show scope, condition, expiry, escalation, and reauthorization evidence |
| SOURCE SUPPORT          | AI-01, AI-06, AI-08, AUD-06                                             |
| GUIDE 1 RELATION        | Authority control implementation                                        |
| GUIDE 2 RELATION        | Delegation policy-to-evidence                                           |
| BOUNDARY                | Not equal to IAM permission                                             |

**MRO-03 Agent Role is not Human Role**

|                                |                                                                                 |
|--------------------------------|---------------------------------------------------------------------------------|
| <b>AUDIT EVIDENCE OBJECT</b>   | Agent role boundary object                                                      |
| <b>SAMPLE EVIDENCE REQUEST</b> | Show agent role, capability limits, constraints, owner, and escalation boundary |
| <b>SOURCE SUPPORT</b>          | AI-08, EVID-01, AI-01                                                           |
| <b>GUIDE 1 RELATION</b>        | Agent role design                                                               |
| <b>GUIDE 2 RELATION</b>        | Accountability model                                                            |
| <b>BOUNDARY</b>                | Do not anthropomorphize agents                                                  |

**MRO-04 Accepted Outcome Compliance**

|                                |                                                                                           |
|--------------------------------|-------------------------------------------------------------------------------------------|
| <b>AUDIT EVIDENCE OBJECT</b>   | Accepted outcome record                                                                   |
| <b>SAMPLE EVIDENCE REQUEST</b> | Show acceptance criteria, reviewer, acceptance/rejection/dispute state, and evidence link |
| <b>SOURCE SUPPORT</b>          | AUD-01, AUD-02, AUD-04                                                                    |
| <b>GUIDE 1 RELATION</b>        | Outcome state capture                                                                     |
| <b>GUIDE 2 RELATION</b>        | Review/acceptance governance                                                              |
| <b>BOUNDARY</b>                | Acceptance is not compliance proof                                                        |

**MRO-05 Tool-Action Liability Boundary**

|                                |                                                                                                        |
|--------------------------------|--------------------------------------------------------------------------------------------------------|
| <b>AUDIT EVIDENCE OBJECT</b>   | Tool-action evidence object                                                                            |
| <b>SAMPLE EVIDENCE REQUEST</b> | Show tool identity, action, affected system/data, authority basis, reversibility, rollback path, owner |
| <b>SOURCE SUPPORT</b>          | EVID-02, EVID-03, EVID-04, AUD-05                                                                      |
| <b>GUIDE 1 RELATION</b>        | Tool evidence capture                                                                                  |
| <b>GUIDE 2 RELATION</b>        | Tool-action approval policy                                                                            |
| <b>BOUNDARY</b>                | No legal liability conclusion                                                                          |

**MRO-06 Responsibility Transfer Across Agents**

|                                |                                                                                                          |
|--------------------------------|----------------------------------------------------------------------------------------------------------|
| <b>AUDIT EVIDENCE OBJECT</b>   | Responsibility transfer record                                                                           |
| <b>SAMPLE EVIDENCE REQUEST</b> | Show source/target agent, transferred scope, retained scope, inherited constraints, acceptance/rejection |
| <b>SOURCE SUPPORT</b>          | EVID-01, AI-08, AUD-06                                                                                   |
| <b>GUIDE 1 RELATION</b>        | Handoff record design                                                                                    |
| <b>GUIDE 2 RELATION</b>        | Transfer governance                                                                                      |
| <b>BOUNDARY</b>                | Not legal transfer                                                                                       |

**MRO-07 Authority Drift**

|                                |                                                                                               |
|--------------------------------|-----------------------------------------------------------------------------------------------|
| <b>AUDIT EVIDENCE OBJECT</b>   | Authority drift exception record                                                              |
| <b>SAMPLE EVIDENCE REQUEST</b> | Show baseline authority, observed action, trigger, escalation, stop/downgrade/reauthorization |
| <b>SOURCE SUPPORT</b>          | AI-01, AI-02, EVID-04, AUD-04                                                                 |
| <b>GUIDE 1 RELATION</b>        | Drift monitoring                                                                              |
| <b>GUIDE 2 RELATION</b>        | Exception governance                                                                          |
| <b>BOUNDARY</b>                | Not every drift is a legal breach                                                             |

**MRO-08 MAS Evidence Partitioning**

|                                |                                                                                          |
|--------------------------------|------------------------------------------------------------------------------------------|
| <b>AUDIT EVIDENCE OBJECT</b>   | Partitioned evidence chain                                                               |
| <b>SAMPLE EVIDENCE REQUEST</b> | Show partition keys, evidence pointers, link identifiers, integrity hashes, access rules |
| <b>SOURCE SUPPORT</b>          | EVID-01, EVID-02, EVID-03, AUD-01                                                        |
| <b>GUIDE 1 RELATION</b>        | Evidence partitioning architecture                                                       |
| <b>GUIDE 2 RELATION</b>        | Evidence access governance                                                               |
| <b>BOUNDARY</b>                | No blanket retention/disclosure                                                          |

**MRO-09 Cross-Project Reuse Compliance**

|                                |                                                                                              |
|--------------------------------|----------------------------------------------------------------------------------------------|
| <b>AUDIT EVIDENCE OBJECT</b>   | Reuse context validation record                                                              |
| <b>SAMPLE EVIDENCE REQUEST</b> | Show source context, target context, reset validation, privacy review, authorization refresh |
| <b>SOURCE SUPPORT</b>          | AI-01, AI-05, PRIV-02, AUD-04                                                                |
| <b>GUIDE 1 RELATION</b>        | Reuse validation workflow                                                                    |
| <b>GUIDE 2 RELATION</b>        | Reuse governance                                                                             |
| <b>BOUNDARY</b>                | No legal reuse clearance                                                                     |

**MRO-10 Privacy / GDPR Lifecycle Mapping**

|                                |                                                                                             |
|--------------------------------|---------------------------------------------------------------------------------------------|
| <b>AUDIT EVIDENCE OBJECT</b>   | Privacy lifecycle evidence object                                                           |
| <b>SAMPLE EVIDENCE REQUEST</b> | Show data flow, processing purpose, retention rule, rights workflow, minimization treatment |
| <b>SOURCE SUPPORT</b>          | PRIV-01 to PRIV-05, AI-09                                                                   |
| <b>GUIDE 1 RELATION</b>        | Privacy-aware evidence capture                                                              |
| <b>GUIDE 2 RELATION</b>        | Privacy governance                                                                          |
| <b>BOUNDARY</b>                | No legal advice or GDPR proof                                                               |

**MRO-11 Privacy-Preserving Third-Party Validation**

|                                |                                                                                             |
|--------------------------------|---------------------------------------------------------------------------------------------|
| <b>AUDIT EVIDENCE OBJECT</b>   | Validation disclosure profile and verdict record                                            |
| <b>SAMPLE EVIDENCE REQUEST</b> | Show ruleset identity, redaction profile, evidence pointer, verdict hash, reviewer boundary |
| <b>SOURCE SUPPORT</b>          | PRIV-02, PRIV-03, BOUND-01, BOUND-02, AUD-03                                                |
| <b>GUIDE 1 RELATION</b>        | Evidence export/replay                                                                      |
| <b>GUIDE 2 RELATION</b>        | Third-party review governance                                                               |
| <b>BOUNDARY</b>                | No certification claim                                                                      |

**MRO-12 Evidence Minimization and Selective Disclosure**

|                                |                                                                                         |
|--------------------------------|-----------------------------------------------------------------------------------------|
| <b>AUDIT EVIDENCE OBJECT</b>   | Evidence minimization and disclosure object                                             |
| <b>SAMPLE EVIDENCE REQUEST</b> | Show minimization rule, disclosure justification, redaction profile, access log, expiry |
| <b>SOURCE SUPPORT</b>          | PRIV-01 to PRIV-05                                                                      |
| <b>GUIDE 1 RELATION</b>        | Selective disclosure design                                                             |
| <b>GUIDE 2 RELATION</b>        | Disclosure policy                                                                       |
| <b>BOUNDARY</b>                | No privacy compliance proof                                                             |

**MRO-13 Data Subject Rights vs Evidence Retention**

|                                |                                                                                   |
|--------------------------------|-----------------------------------------------------------------------------------|
| <b>AUDIT EVIDENCE OBJECT</b>   | Rights-retention reconciliation record                                            |
| <b>SAMPLE EVIDENCE REQUEST</b> | Show request type, retention basis, hold status, deletion/redaction action, owner |
| <b>SOURCE SUPPORT</b>          | PRIV-01, PRIV-02, PRIV-05, AI-09                                                  |
| <b>GUIDE 1 RELATION</b>        | Rights-aware retention                                                            |
| <b>GUIDE 2 RELATION</b>        | Legal/privacy governance                                                          |
| <b>BOUNDARY</b>                | Requires legal review                                                             |

**MRO-14 Third-Party Processor / Subprocessor Chain**

|                                |                                                                                       |
|--------------------------------|---------------------------------------------------------------------------------------|
| <b>AUDIT EVIDENCE OBJECT</b>   | Processor/subprocessor chain evidence                                                 |
| <b>SAMPLE EVIDENCE REQUEST</b> | Show processor role, subprocessor chain, data agreement pointer, responsibility owner |
| <b>SOURCE SUPPORT</b>          | AI-09, PRIV-01, PRIV-02, AUD-04                                                       |
| <b>GUIDE 1 RELATION</b>        | Vendor evidence partitioning                                                          |
| <b>GUIDE 2 RELATION</b>        | Processor governance                                                                  |
| <b>BOUNDARY</b>                | No vendor ranking                                                                     |

**MRO-15 Vendor / Model / Runtime Substitution Conformance**

|                                |                                                                                                     |
|--------------------------------|-----------------------------------------------------------------------------------------------------|
| <b>AUDIT EVIDENCE OBJECT</b>   | Substitution conformance record                                                                     |
| <b>SAMPLE EVIDENCE REQUEST</b> | Show prior/new component, control evidence, evidence-integrity check, reauthorization, revalidation |
| <b>SOURCE SUPPORT</b>          | AI-01, AI-04, AI-05, AUD-04, BOUND-02                                                               |
| <b>GUIDE 1 RELATION</b>        | Runtime/model change evidence                                                                       |
| <b>GUIDE 2 RELATION</b>        | Change governance                                                                                   |
| <b>BOUNDARY</b>                | No conformance certification                                                                        |

**MRO-16 Incident, Dispute, and Remediation Closure**

|                                |                                                                                                       |
|--------------------------------|-------------------------------------------------------------------------------------------------------|
| <b>AUDIT EVIDENCE OBJECT</b>   | Incident/remediation closure record                                                                   |
| <b>SAMPLE EVIDENCE REQUEST</b> | Show event, affected outcome, owner, corrective action, recheck, closure acceptance, reopen condition |
| <b>SOURCE SUPPORT</b>          | EVID-04, AI-02, AUD-04, AUD-05, AUD-06                                                                |
| <b>GUIDE 1 RELATION</b>        | Remediation workflow evidence                                                                         |
| <b>GUIDE 2 RELATION</b>        | Closure governance                                                                                    |
| <b>BOUNDARY</b>                | No legal closure or regulator acceptance                                                              |

**Table note:** GAIC-derived and author-synthesis mapping grounded in audit/control/governance source language. It is not a legal checklist, audit procedure, certification criterion, or procurement tool.

**Cross-Links**

- Appendix C contains the expanded mapping.
- Chapter 13 maps the same objects into AARM dimensions.
- Guide 1 will later translate this table into implementation planning.
- Guide 2 will later translate it into compliance operating-model planning.

**Boundary Note**

This chapter does not invent new MRO numbering. It uses current GAIC source truth for MRO-01 through MRO-16.

---

## 7. Evidence Request List for Agentic Systems

**Chapter purpose:** Define the logic of evidence requests for agentic systems without turning them into formal audit procedures. **Reader question:** What should reviewers ask to reconstruct agentic lifecycle work?

---

### Key Claims

- Evidence requests should be organized by lifecycle stage and review objective.
- Agentic evidence requests must cover role, authority, agent/tool action, outcomes, exceptions, privacy, partitioning, and closure.
- Evidence requests are readiness architecture, not universal legal demands or formal audit procedures.

---

### Source Grounding Note

Audit evidence and evidence sufficiency language is grounded in AUD-01 and AUD-02. Control and AI audit practice context comes from AUD-04, AUD-05, and AUD-06. AI governance/logging context may use AI-06.

---

### Author Synthesis Note

The evidence request taxonomy is author synthesis in this paper derived from GAIC MROs and professional audit/control language.

---

### Main Text

Agentic auditability becomes practical when it is expressed as evidence requests. A high-level statement that "the system is logged" is not enough. A reviewer needs to know what evidence should exist, where it lives, which lifecycle stage it supports, what role owns it, how sensitive fields are treated, and what boundary applies.

The evidence request list should begin with the lifecycle work unit. The reviewer should be able to ask: What work was initiated? What purpose and scope did it have? Who initiated it? Which agent or workflow received it? Which policy, authority, or business context applied? Without a work unit, later evidence cannot be connected.

The second request category is authority. What delegated authority allowed the agent to act? Was the action within scope? Did the authority have conditions, risk class, expiry, confirmation, or escalation requirements? If authority changed, was the change recorded? If authority drift occurred, was it escalated and closed?

The third category is responsibility. Which human role owned intent, review, acceptance, exception handling, and remediation? Which agent role executed the work? Was a responsibility transfer recorded across agents? Was a tool action tied back to a human or organizational owner? This category distinguishes accountability from technical execution.

The fourth category is tool-action evidence. Tool calls are where agent behavior often becomes operational consequence. Evidence should identify the tool, action, target system, affected data or process, response, reversibility, rollback path, and owner. The request is not asking for legal liability conclusions. It is asking for reconstructability.

The fifth category is outcome evidence. Reviewers should distinguish produced output from accepted outcome. Was the output accepted, rejected, disputed, escalated, or remediated? Who reviewed it? Which criteria or policy were applied? What evidence supports the final state?

The sixth category is exception and remediation. What event occurred? What lifecycle object was affected? Who owned the exception? What corrective action was taken? Was the action rechecked? Who accepted closure? What would reopen the issue?

The seventh category is privacy and selective disclosure. Evidence should not be over-collected or over-disclosed. Reviewers should ask for data categories, redaction profiles, retention rules, disclosure scope, access logs, and minimization rationale. Legal interpretation remains outside this paper, but the evidence architecture should make privacy-aware review possible.

Finally, evidence requests should include boundary statements. Each request should clarify whether it supports readiness, reconstruction, control review, third-party review, or assurance planning. It should also clarify what it does not prove.

**Table 4: Lifecycle Stage -> Audit Evidence Request**

**Table 4: Lifecycle Stage → Audit Evidence Request**

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

**Intent**

|                             |                                                  |
|-----------------------------|--------------------------------------------------|
| <b>CORE REVIEW QUESTION</b> | What work was initiated and why?                 |
| <b>EVIDENCE REQUESTED</b>   | Work unit ID, purpose, scope, initiating context |
| <b>HUMAN ROLE</b>           | Request owner                                    |
| <b>AGENT/TOOL RECORD</b>    | Initial agent/workflow assignment                |
| <b>OUTCOME/CLOSURE</b>      | No accepted outcome yet                          |
| <b>PRIVACY TREATMENT</b>    | Purpose and data categories                      |

**Delegation**

|                             |                                                            |
|-----------------------------|------------------------------------------------------------|
| <b>CORE REVIEW QUESTION</b> | Was action authorized within scope?                        |
| <b>EVIDENCE REQUESTED</b>   | Authority grant, conditions, expiry, confirmation boundary |
| <b>HUMAN ROLE</b>           | Delegating/approving role                                  |
| <b>AGENT/TOOL RECORD</b>    | Agent permission and tool scope                            |
| <b>OUTCOME/CLOSURE</b>      | Escalation path                                            |
| <b>PRIVACY TREATMENT</b>    | Minimum necessary authority evidence                       |

**Planning**

|                             |                                         |
|-----------------------------|-----------------------------------------|
| <b>CORE REVIEW QUESTION</b> | What plan or route was selected?        |
| <b>EVIDENCE REQUESTED</b>   | Plan version, constraints, review flags |
| <b>HUMAN ROLE</b>           | Plan reviewer if required               |
| <b>AGENT/TOOL RECORD</b>    | Agent plan trace                        |
| <b>OUTCOME/CLOSURE</b>      | Plan accepted/rejected state            |
| <b>PRIVACY TREATMENT</b>    | Redacted sensitive prompts if needed    |

**Tool action**

|                             |                                                                       |
|-----------------------------|-----------------------------------------------------------------------|
| <b>CORE REVIEW QUESTION</b> | What external consequence occurred?                                   |
| <b>EVIDENCE REQUESTED</b>   | Tool identity, action, target, response, affected data, reversibility |
| <b>HUMAN ROLE</b>           | Tool-action owner                                                     |
| <b>AGENT/TOOL RECORD</b>    | Tool-call record, service trace                                       |
| <b>OUTCOME/CLOSURE</b>      | Action status and rollback path                                       |
| <b>PRIVACY TREATMENT</b>    | Data minimization and access scope                                    |

**Handoff**

|                             |                                                                              |
|-----------------------------|------------------------------------------------------------------------------|
| <b>CORE REVIEW QUESTION</b> | Was responsibility transferred?                                              |
| <b>EVIDENCE REQUESTED</b>   | Source/target agent, transferred scope, retained scope, acceptance/rejection |
| <b>HUMAN ROLE</b>           | Handoff owner                                                                |
| <b>AGENT/TOOL RECORD</b>    | Agent transfer record                                                        |
| <b>OUTCOME/CLOSURE</b>      | Transfer state                                                               |
| <b>PRIVACY TREATMENT</b>    | Partition and disclosure rule                                                |

**Review**

|                             |                                           |
|-----------------------------|-------------------------------------------|
| <b>CORE REVIEW QUESTION</b> | Was work reviewed before acceptance?      |
| <b>EVIDENCE REQUESTED</b>   | Review record, criteria, evidence pointer |
| <b>HUMAN ROLE</b>           | Reviewer role                             |
| <b>AGENT/TOOL RECORD</b>    | Supporting traces/outputs                 |
| <b>OUTCOME/CLOSURE</b>      | Accepted/rejected/disputed                |
| <b>PRIVACY TREATMENT</b>    | Redaction profile                         |

|                             |                                             |
|-----------------------------|---------------------------------------------|
| <b>Exception</b>            |                                             |
| <b>CORE REVIEW QUESTION</b> | What deviation occurred?                    |
| <b>EVIDENCE REQUESTED</b>   | Exception record, baseline, trigger, impact |
| <b>HUMAN ROLE</b>           | Escalation owner                            |
| <b>AGENT/TOOL RECORD</b>    | Incident trace or alert                     |
| <b>OUTCOME/CLOSURE</b>      | Open/remediated/closed                      |
| <b>PRIVACY TREATMENT</b>    | Sensitive-data handling                     |

|                             |                                                                     |
|-----------------------------|---------------------------------------------------------------------|
| <b>Closure</b>              |                                                                     |
| <b>CORE REVIEW QUESTION</b> | Was remediation complete?                                           |
| <b>EVIDENCE REQUESTED</b>   | Corrective action, recheck evidence, closure owner, reopen criteria |
| <b>HUMAN ROLE</b>           | Closure owner                                                       |
| <b>AGENT/TOOL RECORD</b>    | Remediation workflow trace                                          |
| <b>OUTCOME/CLOSURE</b>      | Closure state                                                       |
| <b>PRIVACY TREATMENT</b>    | Retention/disclosure rule                                           |

**Table note:** Mixed source-grounded and author-synthesis table. It is a readiness evidence architecture, not a formal audit procedure, legal demand, or certification criterion.

- Cross-Links
- Appendix A expands the evidence request list.
  - Appendix B turns these categories into a walkthrough template.
  - Appendix E provides exception/remediation checklist detail.

Boundary Note

Evidence requests do not prove legal compliance or audit sufficiency. A qualified professional must evaluate sufficiency in context.

---

## 8. Lifecycle Walkthrough for AI Agent / MAS Work

**Chapter purpose:** Provide a reconstruction pattern for one agentic lifecycle run or work unit. **Reader question:** How is one agentic workflow walked from intent to accepted outcome?

---

### Key Claims

- A lifecycle walkthrough should connect intent, authority, planning, execution, handoff, evidence, review, outcome, exception, remediation, and closure.
- The walkthrough is illustrative and readiness-oriented, not a formal audit procedure.
- Walkthroughs should preserve human role and agent role separation.

---

### Source Grounding Note

Control and walkthrough-adjacent language uses AUD-04, AUD-05, and AUD-06. Provenance and trace context uses EVID-01 and EVID-02. GAIC source truth provides MRO lifecycle objects.

---

### Author Synthesis Note

The agentic lifecycle walkthrough is author synthesis in this paper. It translates GAIC MROs and evidence request categories into a reconstruction pattern.

---

### Main Text

A lifecycle walkthrough tests whether a single unit of agentic work can be reconstructed. It begins with a concrete work unit: a task, run, agent workflow, multi-agent handoff, tool action, or remediation event. The goal is not to prove compliance. The goal is to determine whether the evidence chain can explain the work from intent to closure.

The walkthrough should begin with intent and scope. What was the system asked to do? Who initiated the work? What business process, policy, or governance context applied? What was out of scope? The purpose of this step is to avoid reviewing an isolated output without understanding the work it belongs to.

The second step is authority. What authority allowed the agent or workflow to act? Was authority delegated by role, policy, system configuration, human confirmation, or another control? Did that authority include limits, expiry, risk class, confirmation

requirements, or escalation paths? If the tool action exceeded scope, was that captured as exception or authority drift?

The third step is agent and tool identification. Which agent role acted? What constraints applied to that role? Which tools were available? Which tools were used? Which external systems, data stores, APIs, or processors were touched? The walkthrough should separate the agent role from the human role and the tool surface.

The fourth step is evidence pointer review. Where are logs, traces, workflow records, approvals, policies, tickets, evidence packs, or hash manifests stored? Are they linked by a stable work unit ID? Are there partition keys by agent, tool, role, vendor, project, data class, and lifecycle stage? Can an authorized reviewer reconstruct the chain without overexposing sensitive data?

The fifth step is outcome review. What output or action occurred? Was the outcome accepted, rejected, disputed, escalated, remediated, or closed? Who had authority to accept it? What evidence supports that state? An accepted outcome is a governance state, not just a completed workflow event.

The sixth step is exception and remediation review. If something deviated, did the system record the baseline, trigger, owner, corrective action, recheck, closure state, and reopen criteria? A lifecycle walkthrough is incomplete if it stops at detection and never reaches closure.

The seventh step is boundary review. What does the walkthrough support? It may support readiness discussion, internal review, evidence request scoping, or future assurance planning. It does not itself produce an audit opinion, certification, legal conclusion, or regulator-approved result.

Table 5: Human Role / Agent / Tool Responsibility Matrix

Table 5: Human Role / Agent / Tool Responsibility Matrix

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

Initiate work

|                  |                                 |
|------------------|---------------------------------|
| HUMAN ROLE       | Request owner                   |
| AGENT ROLE       | Intake agent                    |
| TOOL/SYSTEM      | Workflow system                 |
| AUTHORITY BASIS  | Business request policy         |
| EVIDENCE POINTER | Work unit record                |
| ESCALATION PATH  | Reject or request clarification |

**Approve external action**

|                         |                                         |
|-------------------------|-----------------------------------------|
| <b>HUMAN ROLE</b>       | Business owner or delegated approver    |
| <b>AGENT ROLE</b>       | Execution agent                         |
| <b>TOOL/SYSTEM</b>      | CRM/API/configuration tool              |
| <b>AUTHORITY BASIS</b>  | Delegated authority record              |
| <b>EVIDENCE POINTER</b> | Authority evidence pointer and tool log |
| <b>ESCALATION PATH</b>  | Human confirmation or stop              |

**Transfer task**

|                         |                                        |
|-------------------------|----------------------------------------|
| <b>HUMAN ROLE</b>       | Process owner                          |
| <b>AGENT ROLE</b>       | Source and target agents               |
| <b>TOOL/SYSTEM</b>      | Orchestration runtime                  |
| <b>AUTHORITY BASIS</b>  | Handoff rule and inherited constraints |
| <b>EVIDENCE POINTER</b> | Responsibility transfer record         |
| <b>ESCALATION PATH</b>  | Return/reject/escalate                 |

**Review output**

|                         |                                       |
|-------------------------|---------------------------------------|
| <b>HUMAN ROLE</b>       | Reviewer role                         |
| <b>AGENT ROLE</b>       | Drafting or analysis agent            |
| <b>TOOL/SYSTEM</b>      | Document/review system                |
| <b>AUTHORITY BASIS</b>  | Review policy                         |
| <b>EVIDENCE POINTER</b> | Output, review note, evidence pointer |
| <b>ESCALATION PATH</b>  | Dispute or remediation                |

**Handle exception**

|                         |                                           |
|-------------------------|-------------------------------------------|
| <b>HUMAN ROLE</b>       | Incident/remediation owner                |
| <b>AGENT ROLE</b>       | Monitoring or remediation agent           |
| <b>TOOL/SYSTEM</b>      | Ticketing/incident system                 |
| <b>AUTHORITY BASIS</b>  | Escalation rule                           |
| <b>EVIDENCE POINTER</b> | Exception record and remediation evidence |
| <b>ESCALATION PATH</b>  | Reopen, escalate, or close                |

**Table note:** GAIC-derived and author-synthesis table grounded in internal audit, control, provenance, and observability context. It assigns review responsibility for auditability purposes only and does not assign legal liability.

---

## Walkthrough Output

A completed walkthrough should produce:

- Work unit identity and scope.
- Authority boundary and confirmation history.
- Human role and agent role map.
- Tool-action evidence and side-effect summary.
- Evidence chain with pointers and partition keys.
- Accepted outcome or dispute state.
- Exception/remediation closure state if applicable.
- Privacy treatment and disclosure profile.
- Boundary statement describing what the walkthrough does and does not prove.

---

## Cross-Links

- Appendix B provides the walkthrough template.
- Chapter 9 addresses partitioning across agents, tools, vendors, and projects.
- Chapter 11 addresses exceptions and remediation closure.

---

## Boundary Note

The walkthrough is a reconstruction pattern for readiness and evidence design. It is not a formal audit procedure, assurance engagement step, or legal compliance test.

---

---

## 9. Evidence Partitioning Across Agents, Tools, Roles, Vendors, and Projects

**Chapter purpose:** Define how evidence is separated and linked across lifecycle boundaries. **Reader question:** How do reviewers avoid flattened responsibility in multi-agent and tool-mediated work?

---

### Key Claims

- Evidence must be partitioned and linkable across agent, tool, role, vendor, project, data class, jurisdiction, and lifecycle state where relevant.
- Partitioning enables reconstruction without blanket disclosure.
- Cross-project reuse, third-party processor chains, and vendor/model/runtime substitution require explicit evidence boundaries.

---

### Source Grounding Note

Provenance and observability/logging context uses EVID-01, EVID-02, and EVID-03. Privacy context uses PRIV-01, PRIV-02, PRIV-03, and AI-09. GAIC source truth provides MRO-08, MRO-09, MRO-10, MRO-14, and MRO-15.

---

### Author Synthesis Note

The evidence partitioning model is author synthesis in this paper grounded in provenance, observability, privacy, and GAIC MROs.

---

### Main Text

Agentic evidence becomes difficult to review when everything is flattened into one trace, one log export, one transcript, or one vendor report. Multi-agent systems distribute work across agents, tools, human roles, vendors, processors, projects, and data boundaries. Auditability requires those boundaries to be visible.

Partitioning is the discipline of separating evidence by review-relevant boundary while preserving the links needed for reconstruction. It is not the same as hiding evidence. It is not the same as retaining everything. It is a way to make evidence reviewable, scoped, and privacy-aware.

The first partition is the work unit. Every evidence item should connect to a lifecycle work unit or reconstruction path. Without a work unit, reviewers may know that events happened but not which lifecycle task they belong to.

The second partition is the agent. Agent roles may differ in authority, capability, tool access, evidence duties, and escalation requirements. If evidence only shows that "the system" acted, responsibility surfaces become blurred. Reviewers should be able to distinguish planning agents, execution agents, review agents, monitoring agents, and remediation agents where those roles exist.

The third partition is the tool. Tool actions often create external consequences. Evidence should separate model output from tool execution, tool execution from downstream system state, and reversible actions from irreversible or high-impact actions. Tool-specific evidence can include request/response logs, target system records, affected data categories, rollback records, and owner information.

The fourth partition is the human role. A user account is not enough. Auditability requires role responsibility: request owner, delegating role, reviewer, acceptor, exception owner, remediation owner, privacy owner, or internal audit reviewer. These roles should connect to evidence without implying legal liability.

The fifth partition is vendor, processor, and runtime boundary. Agentic systems may call third-party tools, external APIs, model providers, orchestration runtimes, vector stores, workflow systems, or processors/subprocessors. Evidence should identify where responsibility, data processing, and review access shift. This is not vendor ranking. It is boundary mapping.

The sixth partition is project and reuse context. Cross-project reuse can move prompts, agents, workflows, memory, tools, or policies into a new context. Reuse may be efficient, but auditability requires evidence that scope, authority, privacy treatment, and evidence obligations were reset or revalidated for the new context.

The seventh partition is privacy and disclosure class. Some evidence can be disclosed directly. Some should be redacted. Some may require evidence pointers, hashes, summaries, or reviewer-specific access. The partitioning model should make it possible to review what is necessary without exposing more than the review requires.

Table 6: Evidence Partitioning Matrix

Table 6: Evidence Partitioning Matrix

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

## Work unit

|                               |                                  |
|-------------------------------|----------------------------------|
| <b>WHY IT MATTERS</b>         | Anchors reconstruction           |
| <b>EVIDENCE BOUNDARY</b>      | Work unit ID and lifecycle stage |
| <b>ACCESS/DISCLOSURE RULE</b> | Reviewer access tied to scope    |
| <b>FAILURE IF MISSING</b>     | Evidence cannot be connected     |
| <b>RELATED MRO</b>            | MRO-08                           |

## Agent

|                               |                                           |
|-------------------------------|-------------------------------------------|
| <b>WHY IT MATTERS</b>         | Separates responsibility surfaces         |
| <b>EVIDENCE BOUNDARY</b>      | Agent ID, role, constraints, work segment |
| <b>ACCESS/DISCLOSURE RULE</b> | Role-based evidence disclosure            |
| <b>FAILURE IF MISSING</b>     | Flattened agent responsibility            |
| <b>RELATED MRO</b>            | MRO-03, MRO-06, MRO-08                    |

## Tool

|                               |                                                    |
|-------------------------------|----------------------------------------------------|
| <b>WHY IT MATTERS</b>         | Tracks external consequence                        |
| <b>EVIDENCE BOUNDARY</b>      | Tool ID, target system, action type, reversibility |
| <b>ACCESS/DISCLOSURE RULE</b> | Tool-specific disclosure and redaction             |
| <b>FAILURE IF MISSING</b>     | Unclear side effects                               |
| <b>RELATED MRO</b>            | MRO-05, MRO-08                                     |

## Human role

|                               |                                                 |
|-------------------------------|-------------------------------------------------|
| <b>WHY IT MATTERS</b>         | Connects governance ownership                   |
| <b>EVIDENCE BOUNDARY</b>      | Role map, approval, review, acceptance, closure |
| <b>ACCESS/DISCLOSURE RULE</b> | Need-to-know reviewer access                    |
| <b>FAILURE IF MISSING</b>     | User identity mistaken for responsibility       |
| <b>RELATED MRO</b>            | MRO-01                                          |

**Vendor / processor**

|                               |                                                      |
|-------------------------------|------------------------------------------------------|
| <b>WHY IT MATTERS</b>         | Maps third-party boundary                            |
| <b>EVIDENCE BOUNDARY</b>      | Vendor/processor role, subprocessor chain, data path |
| <b>ACCESS/DISCLOSURE RULE</b> | Contract/privacy-aware disclosure                    |
| <b>FAILURE IF MISSING</b>     | Responsibility chain disappears                      |
| <b>RELATED MRO</b>            | MRO-14                                               |

**Project / reuse context**

|                               |                                                  |
|-------------------------------|--------------------------------------------------|
| <b>WHY IT MATTERS</b>         | Prevents context drift                           |
| <b>EVIDENCE BOUNDARY</b>      | Source context, target context, reset validation |
| <b>ACCESS/DISCLOSURE RULE</b> | Reuse-specific review access                     |
| <b>FAILURE IF MISSING</b>     | Old authority/policy reused incorrectly          |
| <b>RELATED MRO</b>            | MRO-09                                           |

**Runtime/model substitution**

|                               |                                            |
|-------------------------------|--------------------------------------------|
| <b>WHY IT MATTERS</b>         | Preserves evidence continuity after change |
| <b>EVIDENCE BOUNDARY</b>      | Prior/new component, revalidation evidence |
| <b>ACCESS/DISCLOSURE RULE</b> | Change-scoped disclosure                   |
| <b>FAILURE IF MISSING</b>     | Evidence chain breaks after substitution   |
| <b>RELATED MRO</b>            | MRO-15                                     |

**Privacy class**

|                               |                                                  |
|-------------------------------|--------------------------------------------------|
| <b>WHY IT MATTERS</b>         | Limits overexposure                              |
| <b>EVIDENCE BOUNDARY</b>      | Data category, redaction profile, retention rule |
| <b>ACCESS/DISCLOSURE RULE</b> | Minimized disclosure                             |
| <b>FAILURE IF MISSING</b>     | Privacy leakage or evidence hoarding             |
| <b>RELATED MRO</b>            | MRO-10, MRO-12, MRO-13                           |

**Table note:** Mixed source-grounded and author-synthesis table. It does not create universal disclosure or retention rules, vendor rankings, or procurement

recommendations.

---

## Cross-Links

- Chapter 10 covers privacy and selective disclosure.
- Appendix A includes partition fields for evidence requests.
- Appendix C maps partitioning to MRO-08, MRO-09, MRO-14, and MRO-15.

---

## Boundary Note

Partitioning is a review architecture. It does not decide legal retention, data subject rights, processor obligations, or vendor liability.

---

---

## 10. Privacy, Selective Disclosure, and Audit Evidence Minimization

**Chapter purpose:** Balance auditability with privacy, confidentiality, minimization, and retention discipline. **Reader question:** How can evidence be reviewable without overexposure?

---

### Key Claims

- Auditability should not become evidence hoarding.
- Privacy-aware auditability uses minimization, redaction, retention discipline, evidence pointers, partitioning, and selective disclosure.
- this paper frames privacy/evidence tension; it does not provide legal advice or GDPR compliance proof.

---

### Source Grounding Note

Privacy and minimization context uses PRIV-01, PRIV-02, PRIV-03, PRIV-04, PRIV-05, and AI-09. GAIC source truth provides MRO-10 through MRO-13 and the Validation Lab selective disclosure boundary.

---

### Author Synthesis Note

The privacy-preserving audit evidence controls are author synthesis in this paper grounded in public privacy guidance and GAIC MROs.

---

### Main Text

Auditability can create a privacy problem if it is misunderstood. The answer to weak evidence is not to collect everything forever. Agentic systems may process personal data, confidential business data, privileged material, regulated data, customer records, employee records, prompts, outputs, tool responses, logs, memory, and third-party processor data. If auditability becomes indiscriminate evidence capture, it may increase exposure and reduce trust.

The better question is: what evidence is necessary for the review objective, and how can it be disclosed safely? This requires minimization, partitioning, retention discipline, access control, redaction, evidence pointers, and reviewer-specific disclosure profiles.

Privacy-aware auditability begins with purpose. Evidence should be collected and retained for defined review purposes: reconstruction, control review, exception handling, third-party validation, remediation, or assurance planning. A generic desire to "keep all logs just in case" is not a mature evidence strategy.

The second principle is minimization. Evidence objects should identify what fields are needed for reconstruction and which fields can be redacted, tokenized, summarized, hashed, or replaced with pointers. A reviewer may need to know that a tool action affected a customer data category without seeing every raw value. Another reviewer may need access under a narrower approved scope. The evidence architecture should support both.

The third principle is selective disclosure. Different reviewers need different levels of evidence. Internal engineering may need logs and traces. Internal audit may need work unit chains and evidence pointers. Privacy teams may need data category and retention records. Third-party reviewers may need redacted evidence packs, ruleset identity, verdict hashes, or replay packages. The system should make disclosure explicit, scoped, and recorded.

The fourth principle is retention discipline. Evidence needs may conflict with storage limitation and deletion expectations. This paper does not decide legal retention periods. It asks that systems record retention rules, expiry, holds, deletion/redaction actions, and rights-retention reconciliation where relevant. A reviewable retention decision is stronger than a silent default.

The fifth principle is boundary clarity. Privacy sources provide context, not legal conclusions. A privacy-preserving evidence architecture does not prove GDPR compliance, UK GDPR compliance, EU AI Act compliance, or any jurisdiction-specific legal compliance. It creates the object layer that legal, privacy, audit, and governance professionals can review.

## Table 7: Privacy-Preserving Audit Evidence Controls

Table 7: Privacy-Preserving Audit Evidence Controls

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

|                                       |                                                |
|---------------------------------------|------------------------------------------------|
| <b>Evidence minimization</b>          |                                                |
| <b>PURPOSE</b>                        | Avoid unnecessary collection                   |
| <b>EVIDENCE ARTIFACT</b>              | Minimization rule, field inventory             |
| <b>SELECTIVE DISCLOSURE MECHANISM</b> | Request only review-relevant fields            |
| <b>PRIVACY RISK REDUCED</b>           | Overcollection                                 |
| <b>RESIDUAL REVIEW NEED</b>           | Legal/privacy review for purpose and necessity |

**Redaction profile**

|                                               |                                   |
|-----------------------------------------------|-----------------------------------|
| <b>PURPOSE</b>                                | Remove or mask sensitive fields   |
| <b>EVIDENCE<br/>ARTIFACT</b>                  | Redaction rule and redaction log  |
| <b>SELECTIVE<br/>DISCLOSURE<br/>MECHANISM</b> | Redacted evidence pointer         |
| <b>PRIVACY RISK<br/>REDUCED</b>               | Overexposure to reviewer          |
| <b>RESIDUAL REVIEW<br/>NEED</b>               | Sufficiency check after redaction |

**Evidence pointer**

|                                               |                                             |
|-----------------------------------------------|---------------------------------------------|
| <b>PURPOSE</b>                                | Reference evidence without copying raw data |
| <b>EVIDENCE<br/>ARTIFACT</b>                  | Pointer, hash, manifest entry               |
| <b>SELECTIVE<br/>DISCLOSURE<br/>MECHANISM</b> | Reviewer-scoped access or replay            |
| <b>PRIVACY RISK<br/>REDUCED</b>               | Data duplication                            |
| <b>RESIDUAL REVIEW<br/>NEED</b>               | Integrity and access review                 |

**Retention rule**

|                                               |                                            |
|-----------------------------------------------|--------------------------------------------|
| <b>PURPOSE</b>                                | Limit storage duration or define hold      |
| <b>EVIDENCE<br/>ARTIFACT</b>                  | Retention schedule, expiry, hold indicator |
| <b>SELECTIVE<br/>DISCLOSURE<br/>MECHANISM</b> | Time-bound access                          |
| <b>PRIVACY RISK<br/>REDUCED</b>               | Excessive retention                        |
| <b>RESIDUAL REVIEW<br/>NEED</b>               | Legal/privacy retention analysis           |

**Disclosure profile**

|                                               |                               |
|-----------------------------------------------|-------------------------------|
| <b>PURPOSE</b>                                | Define who can see what       |
| <b>EVIDENCE<br/>ARTIFACT</b>                  | Disclosure matrix, access log |
| <b>SELECTIVE<br/>DISCLOSURE<br/>MECHANISM</b> | Role-based disclosure package |
| <b>PRIVACY RISK<br/>REDUCED</b>               | Unscoped sharing              |
| <b>RESIDUAL REVIEW<br/>NEED</b>               | Reviewer authorization review |

**Rights-retention reconciliation**

|                                               |                                                                |
|-----------------------------------------------|----------------------------------------------------------------|
| <b>PURPOSE</b>                                | Address tension between rights requests and evidence retention |
| <b>EVIDENCE<br/>ARTIFACT</b>                  | Request record, action, hold status, rationale                 |
| <b>SELECTIVE<br/>DISCLOSURE<br/>MECHANISM</b> | Redaction/deletion/hold decision path                          |
| <b>PRIVACY RISK<br/>REDUCED</b>               | Silent conflict between evidence and rights                    |
| <b>RESIDUAL REVIEW<br/>NEED</b>               | Legal review                                                   |

**Third-party validation package**

|                                               |                                                  |
|-----------------------------------------------|--------------------------------------------------|
| <b>PURPOSE</b>                                | Support review without raw exposure              |
| <b>EVIDENCE<br/>ARTIFACT</b>                  | Redacted evidence pack, ruleset ID, verdict hash |
| <b>SELECTIVE<br/>DISCLOSURE<br/>MECHANISM</b> | Non-raw evidence adjudication                    |
| <b>PRIVACY RISK<br/>REDUCED</b>               | Third-party overexposure                         |
| <b>RESIDUAL REVIEW<br/>NEED</b>               | Boundary and conflict review                     |

**Table note:** Source-grounded and GAIC-derived table. It is not legal advice, privacy compliance proof, or universal retention guidance.

---

## Cross-Links

- Chapter 9 addresses evidence partitioning.
- Chapter 12 addresses third-party validation boundaries.
- Appendix F provides privacy/legal boundary language.

---

## Boundary Note

This chapter does not provide legal advice, interpret specific legal obligations, or claim that any architecture proves compliance with GDPR, UK GDPR, EU AI Act, or other privacy laws.

---

---

# 11. Exception, Dispute, and Remediation Closure

**Chapter purpose:** Define closure evidence for failed, disputed, drifted, or remediated agentic work. **Reader question:** How does auditability handle exceptions and closure?

---

## Key Claims

- Exceptions, disputes, authority drift, outcome rejection, and remediation must close as evidence-backed lifecycle states.
- Incident logs alone do not prove closure.
- Closure evidence should identify owner, corrective action, recheck, accepted closure state, and reopen criteria.

---

## Source Grounding Note

Incident/governance context uses EVID-04 and AI-02. Control and AI audit practice context uses AUD-04, AUD-05, and AUD-06. GAIC source truth provides MRO-04, MRO-07, and MRO-16.

---

## Author Synthesis Note

Remediation closure as an audit evidence state is GAIC-derived author synthesis. It is not a legal settlement, regulator closure, or assurance conclusion.

---

## Main Text

Agentic auditability is incomplete if it only reconstructs successful work. Many of the most important review questions arise when an agent acts outside expected authority, a tool action creates unintended consequence, an outcome is disputed, a privacy issue appears, a handoff fails, or remediation is required.

In ordinary operations, exceptions often become tickets. Tickets are useful, but they are not automatically lifecycle evidence. A ticket may describe an issue, assign an owner, record comments, and mark a status. For auditability, the ticket must connect to the agentic work unit, authority boundary, role map, tool action, evidence pointer, accepted outcome state, privacy treatment, corrective action, and closure state.

Closure is the critical concept. An incident can be detected without being remediated. A remediation can be attempted without being rechecked. A recheck can occur without being accepted. An accepted closure can later be reopened. Auditability requires those state transitions to be evidence-backed.

For agentic systems, exceptions should be categorized by lifecycle object. Was the issue an authority exception, where observed behavior exceeded delegated scope? Was it a responsibility exception, where no human owner could be identified? Was it a tool-action exception, where a tool created unexpected side effects? Was it an accepted-outcome dispute, where a result was produced but not accepted? Was it a privacy exception, where evidence collection or disclosure exceeded the intended boundary? Was it a substitution exception, where evidence integrity broke after runtime/model/tool change?

The closure record should answer five questions:

- 1. What happened and which lifecycle object was affected?
- 2. Who owned the exception and remediation?
- 3. What corrective action was taken?
- 4. What evidence shows recheck or review?
- 5. Who accepted closure, and what would reopen the issue?

None of these questions determines legal fault. None proves regulatory acceptance. None settles liability. They make the remediation lifecycle reconstructable.

Table 10: Exception / Dispute / Remediation Evidence Checklist

| Table 10: Exception / Dispute / Remediation Evidence Checklist                                                      |                                                                  |
|---------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------|
| Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces. |                                                                  |
| <b>Authority exception</b>                                                                                          |                                                                  |
| <b>EVIDENCE REQUIRED</b>                                                                                            | Baseline authority, observed action, trigger, affected work unit |
| <b>OWNER</b>                                                                                                        | Escalation owner                                                 |
| <b>ESCALATION PATH</b>                                                                                              | Stop, downgrade, reauthorize, or human confirmation              |
| <b>REMEDATION RECORD</b>                                                                                            | Reauthorization or stop-action record                            |
| <b>CLOSURE EVIDENCE</b>                                                                                             | Closure acceptance and new boundary                              |
| <b>BOUNDARY</b>                                                                                                     | Not legal determination                                          |

**Outcome dispute**

|                            |                                                             |
|----------------------------|-------------------------------------------------------------|
| <b>EVIDENCE REQUIRED</b>   | Output, acceptance criteria, dispute reason, reviewer notes |
| <b>OWNER</b>               | Review owner                                                |
| <b>ESCALATION PATH</b>     | Secondary review or remediation                             |
| <b>REMEDICATION RECORD</b> | Correction/review path                                      |
| <b>CLOSURE EVIDENCE</b>    | Accepted, rejected, or remediated state                     |
| <b>BOUNDARY</b>            | Not settlement proof                                        |

**Tool-action side effect**

|                            |                                                                |
|----------------------------|----------------------------------------------------------------|
| <b>EVIDENCE REQUIRED</b>   | Tool call, target system, affected data/process, reversibility |
| <b>OWNER</b>               | Tool-action owner                                              |
| <b>ESCALATION PATH</b>     | Rollback, disable, escalate                                    |
| <b>REMEDICATION RECORD</b> | Rollback or corrective action                                  |
| <b>CLOSURE EVIDENCE</b>    | Recheck and owner acceptance                                   |
| <b>BOUNDARY</b>            | Not liability finding                                          |

**Privacy event**

|                            |                                                            |
|----------------------------|------------------------------------------------------------|
| <b>EVIDENCE REQUIRED</b>   | Data category, disclosure path, redaction/retention record |
| <b>OWNER</b>               | Privacy owner                                              |
| <b>ESCALATION PATH</b>     | Privacy/legal review                                       |
| <b>REMEDICATION RECORD</b> | Redaction, deletion, hold, or access restriction           |
| <b>CLOSURE EVIDENCE</b>    | Review note and disclosure update                          |
| <b>BOUNDARY</b>            | Not legal advice                                           |

**Handoff failure**

|                           |                                                        |
|---------------------------|--------------------------------------------------------|
| <b>EVIDENCE REQUIRED</b>  | Source/target agent, transferred scope, retained scope |
| <b>OWNER</b>              | Process owner                                          |
| <b>ESCALATION PATH</b>    | Return, reassign, or escalate                          |
| <b>REMEDIATION RECORD</b> | Transfer correction                                    |
| <b>CLOSURE EVIDENCE</b>   | Accepted transfer or closed exception                  |
| <b>BOUNDARY</b>           | Not legal responsibility transfer                      |

**Remediation defect**

|                           |                                                        |
|---------------------------|--------------------------------------------------------|
| <b>EVIDENCE REQUIRED</b>  | Prior closure evidence, reopened issue, failed recheck |
| <b>OWNER</b>              | Remediation owner                                      |
| <b>ESCALATION PATH</b>    | Reopen and reassign                                    |
| <b>REMEDIATION RECORD</b> | Updated corrective action                              |
| <b>CLOSURE EVIDENCE</b>   | Recheck and closure reviewer                           |
| <b>BOUNDARY</b>           | Not regulator closure                                  |

**Table note:** Mixed source-grounded, GAIC-derived, and author-synthesis checklist. It is not an audit procedure, legal remedy checklist, or certification criterion.

**Cross-Links**

- Appendix E provides an expanded closure checklist.
- Chapter 13 includes exception traceability and remediation closure as AARM dimensions.
- Chapter 16 connects closure to future Guide 1, Guide 2, and the later insurability white paper.

**Boundary Note**

Closure in this chapter means evidence-backed governance closure. It does not mean legal closure, settlement, regulator acceptance, assurance opinion, or certification.

---

## 12. Third-Party Validation Without Certification Claims

**Chapter purpose:** Define how evidence review can involve third parties without becoming certification, assurance opinion, conformity assessment, legal proof, or regulator approval. **Reader question:** What can third-party validation support without becoming certification?

---

### Key Claims

- Third-party validation can support evidence review, replay, adjudication, and readiness discussion.
- Third-party validation does not certify compliance or issue assurance opinions unless a separate authorized engagement or conformity-assessment framework exists.
- Validation Lab is only a non-certifying evidence adjudication example.

---

### Source Grounding Note

Assurance and attestation boundaries are grounded in AUD-03 and BOUND-03. Conformity-assessment and validation/verification boundaries are grounded in BOUND-01 and BOUND-02. Validation Lab boundary is GAIC source truth.

---

### Author Synthesis Note

this paper's third-party validation boundary is author synthesis that uses professional boundary sources to avoid overclaim. It does not define a conformity assessment program.

---

### Main Text

Agentic auditability often benefits from independent review. A third party, internal audit team, governance function, or specialized review group may inspect evidence packs, replay work units, test whether evidence objects exist, adjudicate ruleset conformance, or evaluate whether evidence chains are reconstructable. These activities can be valuable.

They are also easy to overstate. Evidence review is not automatically certification. Validation is not automatically an assurance opinion. A verdict hash is not regulator approval. A ruleset check is not legal compliance proof. A third-party review note is not an audit report unless it is performed under an appropriate professional engagement, criteria, scope, independence framework, and reporting structure.

This chapter exists to keep that boundary visible. This paper uses "third-party validation" in a narrow sense: review of evidence against defined rulesets or reconstruction objectives. The review may support readiness, evidence quality, selective disclosure, replay, exception closure, or future assurance planning. It does not itself create certification or assurance.

Validation Lab is a useful boundary example because GAIC already frames it as non-certifying evidence adjudication. In this paper, Validation Lab should only be described as one example of how evidence-based validation might be operationalized. It must not be described as the only validation path, a certification body, a regulator, an audit firm, an assurance provider, a conformity assessment body, or proof of enterprise readiness.

The distinction matters for enterprises. If an organization receives a third-party evidence review, it should know exactly what was reviewed, against what ruleset, under what scope, with what evidence, and with what limitations. A review may show that evidence exists and can be replayed. It may show that a ruleset passed. It may show that disclosure was minimized. But it does not determine legal compliance or professional audit sufficiency.

The same distinction matters for reviewers. A review team should state its role: internal readiness review, independent evidence review, validation, verification, assurance engagement, attestation engagement, certification, or legal analysis. Each role has different requirements. this paper only defines the auditability object layer that may support those roles.

Table 8: Third-Party Validation Boundary Table

| Table 8: Third-Party Validation Boundary Table                                                                      |                                                   |
|---------------------------------------------------------------------------------------------------------------------|---------------------------------------------------|
| Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces. |                                                   |
| Evidence adjudication                                                                                               |                                                   |
| CAN SUPPORT                                                                                                         | Ruleset conformance review                        |
| CANNOT PROVE                                                                                                        | Legal compliance or audit opinion                 |
| EVIDENCE ARTIFACT                                                                                                   | Ruleset ID, evidence pointer, verdict/hash record |
| BOUNDARY LANGUAGE                                                                                                   | Non-certifying evidence review                    |
| SOURCE IDS                                                                                                          | GAIC-SOURCE, BOUND-01                             |

**Evidence replay**

|                              |                                                   |
|------------------------------|---------------------------------------------------|
| <b>CAN SUPPORT</b>           | Reconstructability check                          |
| <b>CANNOT PROVE</b>          | Operating effectiveness across all periods        |
| <b>EVIDENCE<br/>ARTIFACT</b> | Replay package, work unit chain, integrity record |
| <b>BOUNDARY<br/>LANGUAGE</b> | Scope-limited reconstruction                      |
| <b>SOURCE IDS</b>            | EVID-01, AUD-01                                   |

**Selective disclosure review**

|                              |                                                          |
|------------------------------|----------------------------------------------------------|
| <b>CAN SUPPORT</b>           | Privacy-aware review process                             |
| <b>CANNOT PROVE</b>          | Privacy law compliance                                   |
| <b>EVIDENCE<br/>ARTIFACT</b> | Redaction profile, disclosure log, reviewer access scope |
| <b>BOUNDARY<br/>LANGUAGE</b> | Privacy-aware evidence review, not legal advice          |
| <b>SOURCE IDS</b>            | PRIV-01 to PRIV-05                                       |

**Third-party readiness review**

|                              |                                           |
|------------------------------|-------------------------------------------|
| <b>CAN SUPPORT</b>           | Independent input to readiness discussion |
| <b>CANNOT PROVE</b>          | Assurance conclusion                      |
| <b>EVIDENCE<br/>ARTIFACT</b> | Scope statement, findings, limitations    |
| <b>BOUNDARY<br/>LANGUAGE</b> | Readiness input only                      |
| <b>SOURCE IDS</b>            | AUD-03, BOUND-03                          |

**Validation Lab example**

|                              |                                                                       |
|------------------------------|-----------------------------------------------------------------------|
| <b>CAN SUPPORT</b>           | Non-certifying evidence adjudication example                          |
| <b>CANNOT PROVE</b>          | Certification, regulator approval, conformity assessment, legal proof |
| <b>EVIDENCE<br/>ARTIFACT</b> | Verdict hash, ruleset identity, evidence pointer                      |
| <b>BOUNDARY<br/>LANGUAGE</b> | One optional example, not exclusive path                              |
| <b>SOURCE IDS</b>            | GAIC-SOURCE, BOUND-01, BOUND-02                                       |

**Assurance engagement**

|                          |                                                                     |
|--------------------------|---------------------------------------------------------------------|
| <b>CAN SUPPORT</b>       | Professional assurance under separate criteria and engagement terms |
| <b>CANNOT PROVE</b>      | Anything outside engagement scope                                   |
| <b>EVIDENCE ARTIFACT</b> | Practitioner report, criteria, scope, evidence basis                |
| <b>BOUNDARY LANGUAGE</b> | Outside this paper unless separately authorized                     |
| <b>SOURCE IDS</b>        | AUD-03, BOUND-03                                                    |

**Table note:** Source-grounded boundary table with GAIC-derived example. It is not a certification model, audit procedure, or assurance engagement template.

**Cross-Links**

- Appendix F provides reusable boundary language.
- Appendix D uses third-party review boundary as an AARM dimension.
- Chapter 15 explains professional use by audit and assurance firms.

**Boundary Note**

This chapter must not be read as claiming that any third party, including Validation Lab, certifies agentic systems, issues legal opinions, provides assurance opinions, or has regulator approval.

---

## 13. Agentic Auditability Readiness Model

**Chapter purpose:** Present AARM dimensions and readiness levels as a proposed readiness model. **Reader question:** How can auditability readiness be described without claiming assurance?

---

### Key Claims

- AARM describes readiness for reconstructability, testability, evidence structure, privacy-aware review, and third-party review boundaries.
- AARM is proposed author synthesis in this paper.
- AARM is not a score, audit standard, certification, assurance opinion, legal compliance proof, regulator approval, procurement tool, or vendor ranking.

---

### Source Grounding Note

Audit evidence language uses AUD-01. AI audit practice context uses AUD-05 and AUD-06. AI governance context uses AI-01. Attestation boundary context uses BOUND-03. R4B/R1 AARM baseline and GAIC source truth provide the model structure.

---

### Author Synthesis Note

AARM is author synthesis in this paper. It is a readiness vocabulary, not an external maturity model or professional assurance framework.

---

### Main Text

Once evidence objects are defined, organizations need a way to discuss readiness. AARM provides that vocabulary. It asks whether agentic lifecycle work can be reconstructed, tested, evidenced, and reviewed across the object areas this paper has defined.

AARM deliberately avoids scoring. Scores invite overprecision and can be mistaken for certification, compliance grades, vendor rankings, or procurement guidance. AARM instead uses dimensions and levels. The dimensions ask what must be reviewable. The levels describe the degree to which the lifecycle work is observable, trace-linked, evidence-structured, auditability-ready, or assurance-ready for planning.

The word "Assurance-Ready" is sensitive. In AARM, it does not mean assurance has been issued. It means the evidence architecture may be mature enough to support qualified assurance planning or review within a defined scope. It still depends on

professional scope, criteria, independence, engagement acceptance, and judgment, and it does not guarantee a positive result.

The ten AARM dimensions are:

- 1. Audit object clarity.
- 2. Authority traceability.
- 3. Human / agent responsibility mapping.
- 4. Evidence sufficiency.
- 5. Agent/tool partitioning.
- 6. Exception traceability.
- 7. Accepted outcome evidence.
- 8. Remediation closure.
- 9. Privacy and selective disclosure.
- 10. Third-party review boundary.

The levels describe readiness states. L0 Unobservable means lifecycle work cannot be reconstructed beyond informal narrative or final output. L1 Log-Visible means logs exist but responsibility linkage is weak. L2 Trace-Linked means traces connect some actions to agents, tools, or workflows. L3 Evidence-Structured means lifecycle evidence objects exist. L4 Auditability-Ready means defined-scope evidence chains can reconstruct authority, responsibility, tools, outcomes, exceptions, privacy, and closure. L5 Assurance-Ready means evidence architecture may support professional assurance planning or review within a defined scope.

The main use of AARM is conversation discipline. It helps technology leaders avoid saying "we have logs, so we are audit-ready." It helps governance teams ask whether accepted outcomes and remediation closure are evidence states. It helps audit teams separate raw observability from reconstructable evidence. It helps privacy teams see whether selective disclosure is built in. It helps third-party reviewers clarify what their review does and does not mean.

## Table 9: Auditability Readiness Levels

Table 9: Auditability Readiness Levels

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

## LO Unobservable

|                                |                                                                                      |
|--------------------------------|--------------------------------------------------------------------------------------|
| <b>DEFINITION</b>              | Lifecycle work cannot be reconstructed beyond informal narrative or isolated outputs |
| <b>OBSERVABLE TRAITS</b>       | Missing lifecycle IDs, no authority linkage, no role map, no accepted outcome state  |
| <b>MINIMUM EVIDENCE</b>        | None beyond ad hoc notes or final output                                             |
| <b>WHAT IT DOES NOT PROVE</b>  | Absence of risk, absence of non-compliance, or system safety                         |
| <b>SOURCE / SYNTHESIS NOTE</b> | Author synthesis grounded in AUD-01/AUD-02 expectations                              |

## L1 Log-Visible

|                                |                                                                           |
|--------------------------------|---------------------------------------------------------------------------|
| <b>DEFINITION</b>              | Activity logs exist without responsibility-linked evidence chains         |
| <b>OBSERVABLE TRAITS</b>       | Logs, traces, timestamps, tool events                                     |
| <b>MINIMUM EVIDENCE</b>        | Raw logs or trace exports                                                 |
| <b>WHAT IT DOES NOT PROVE</b>  | Authority, responsibility, acceptance, exception closure, audit readiness |
| <b>SOURCE / SYNTHESIS NOTE</b> | EVID-02/EVID-03 + author synthesis in this paper                          |

## L2 Trace-Linked

|                                |                                                                     |
|--------------------------------|---------------------------------------------------------------------|
| <b>DEFINITION</b>              | Logs/traces link some actions to agents, tools, users, or workflows |
| <b>OBSERVABLE TRAITS</b>       | Workflow IDs, agent/tool traces, partial actor linkage              |
| <b>MINIMUM EVIDENCE</b>        | Trace-linked run records and partial evidence map                   |
| <b>WHAT IT DOES NOT PROVE</b>  | Complete lifecycle responsibility or sufficient evidence            |
| <b>SOURCE / SYNTHESIS NOTE</b> | EVID-01/EVID-02 + author synthesis in this paper                    |

**L3 Evidence-Structured**

|                                |                                                                                         |
|--------------------------------|-----------------------------------------------------------------------------------------|
| <b>DEFINITION</b>              | Lifecycle evidence is organized into objects that support reconstruction                |
| <b>OBSERVABLE TRAITS</b>       | Role maps, authority objects, accepted outcome records, partitioning, exception records |
| <b>MINIMUM EVIDENCE</b>        | MRO-aligned evidence objects and sampled walkthroughs                                   |
| <b>WHAT IT DOES NOT PROVE</b>  | Assurance readiness, legal compliance, operating effectiveness                          |
| <b>SOURCE / SYNTHESIS NOTE</b> | GAIC-derived + AUD-01/AUD-04/AUD-05                                                     |

**L4 Auditability-Ready**

|                                |                                                                             |
|--------------------------------|-----------------------------------------------------------------------------|
| <b>DEFINITION</b>              | Evidence chains can reconstruct defined-scope agentic lifecycle work        |
| <b>OBSERVABLE TRAITS</b>       | Walkthroughs, evidence request response, privacy controls, closure evidence |
| <b>MINIMUM EVIDENCE</b>        | Complete evidence chain for defined scope plus boundary statement           |
| <b>WHAT IT DOES NOT PROVE</b>  | Audit opinion, certification, regulator approval, legal compliance          |
| <b>SOURCE / SYNTHESIS NOTE</b> | author synthesis in this paper grounded in audit/provenance/privacy sources |

**L5 Assurance-Ready**

|                                |                                                                                                    |
|--------------------------------|----------------------------------------------------------------------------------------------------|
| <b>DEFINITION</b>              | Evidence architecture may support professional assurance planning or review within a defined scope |
| <b>OBSERVABLE TRAITS</b>       | Repeatable evidence packs, ruleset identity, reviewer scope, remediation/recheck procedure         |
| <b>MINIMUM EVIDENCE</b>        | Reviewable evidence packs, integrity records, boundary statement                                   |
| <b>WHAT IT DOES NOT PROVE</b>  | Actual assurance opinion, certification, guaranteed compliance, regulator acceptance               |
| <b>SOURCE / SYNTHESIS NOTE</b> | author synthesis in this paper bounded by AUD-03/BOUND-03                                          |

**Table note:** AARM is author synthesis grounded in source language. It is not a score, standard, certification, assurance conclusion, procurement benchmark, or vendor ranking.

## AARM Dimension Summary

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

### Audit object clarity

|                         |                                                     |
|-------------------------|-----------------------------------------------------|
| AUDIT QUESTION          | Can reviewers name the lifecycle unit under review? |
| MINIMUM EVIDENCE FAMILY | Work unit ID, scope, object inventory               |
| PRIMARY MRO LINKS       | MRO-01, MRO-03, MRO-04                              |
| BOUNDARY                | Not legal category                                  |

### Authority traceability

|                         |                                        |
|-------------------------|----------------------------------------|
| AUDIT QUESTION          | Was the action authorized under scope? |
| MINIMUM EVIDENCE FAMILY | Authority grant, expiry, escalation    |
| PRIMARY MRO LINKS       | MRO-02, MRO-07                         |
| BOUNDARY                | Not legal delegation proof             |

### Human / agent responsibility mapping

|                         |                                                               |
|-------------------------|---------------------------------------------------------------|
| AUDIT QUESTION          | Who owned intent, execution, review, acceptance, remediation? |
| MINIMUM EVIDENCE FAMILY | Role map and transfer record                                  |
| PRIMARY MRO LINKS       | MRO-01, MRO-03, MRO-06                                        |
| BOUNDARY                | No anthropomorphism                                           |

### Evidence sufficiency

|                         |                                                          |
|-------------------------|----------------------------------------------------------|
| AUDIT QUESTION          | Can reviewers reconstruct work without narrative memory? |
| MINIMUM EVIDENCE FAMILY | Evidence chain and pointers                              |
| PRIMARY MRO LINKS       | MRO-04, MRO-08, MRO-16                                   |
| BOUNDARY                | Not universal sufficiency                                |

**Agent/tool partitioning**

|                                |                                                                |
|--------------------------------|----------------------------------------------------------------|
| <b>AUDIT QUESTION</b>          | Can actions be separated across agents/tools/vendors/projects? |
| <b>MINIMUM EVIDENCE FAMILY</b> | Partition keys, tool evidence, substitution record             |
| <b>PRIMARY MRO LINKS</b>       | MRO-05, MRO-08, MRO-14, MRO-15                                 |
| <b>BOUNDARY</b>                | No vendor ranking                                              |

**Exception traceability**

|                                |                                                   |
|--------------------------------|---------------------------------------------------|
| <b>AUDIT QUESTION</b>          | Can deviation be linked to owner and next action? |
| <b>MINIMUM EVIDENCE FAMILY</b> | Exception record and impact scope                 |
| <b>PRIMARY MRO LINKS</b>       | MRO-07, MRO-16                                    |
| <b>BOUNDARY</b>                | Not legal violation                               |

**Accepted outcome evidence**

|                                |                                             |
|--------------------------------|---------------------------------------------|
| <b>AUDIT QUESTION</b>          | Who accepted or disputed the result?        |
| <b>MINIMUM EVIDENCE FAMILY</b> | Outcome state, reviewer role, evidence link |
| <b>PRIMARY MRO LINKS</b>       | MRO-04                                      |
| <b>BOUNDARY</b>                | Not compliance proof                        |

**Remediation closure**

|                                |                                               |
|--------------------------------|-----------------------------------------------|
| <b>AUDIT QUESTION</b>          | What was remediated and who accepted closure? |
| <b>MINIMUM EVIDENCE FAMILY</b> | Corrective action, recheck, closure state     |
| <b>PRIMARY MRO LINKS</b>       | MRO-16                                        |
| <b>BOUNDARY</b>                | Not settlement                                |

**Privacy and selective disclosure**

|                                |                                                |
|--------------------------------|------------------------------------------------|
| <b>AUDIT QUESTION</b>          | Can review occur without unnecessary exposure? |
| <b>MINIMUM EVIDENCE FAMILY</b> | Redaction, retention, disclosure profile       |
| <b>PRIMARY MRO LINKS</b>       | MRO-10 to MRO-13                               |
| <b>BOUNDARY</b>                | Not legal advice                               |

**Third-party review boundary**

|                                |                                             |
|--------------------------------|---------------------------------------------|
| <b>AUDIT QUESTION</b>          | What can a reviewer inspect and claim?      |
| <b>MINIMUM EVIDENCE FAMILY</b> | Scope, ruleset, verdict, boundary statement |
| <b>PRIMARY MRO LINKS</b>       | MRO-11, MRO-12, MRO-16                      |
| <b>BOUNDARY</b>                | No certification/assurance opinion          |

**Cross-Links**

- Appendix D provides the expanded readiness matrix.
- Chapter 14 connects AARM to enterprise readiness.
- Chapter 15 connects AARM to professional use boundaries.

**Boundary Note**

AARM is a readiness model only. It does not prove legal compliance, issue assurance, certify systems, rank vendors, bind regulators, or require MPLP.

---

## 14. Enterprise Implementation: CIO / CTO / CCO Readiness

**Chapter purpose:** Connect this paper to enterprise readiness without replacing Guide 1 or Guide 2. **Reader question:** What must leaders prepare before auditability or assurance conversations become credible?

---

### Key Claims

- Audit-ready agentic systems require architecture, governance, evidence ownership, privacy controls, exception closure, and role clarity.
- CIO/CTO readiness focuses on technical evidence architecture.
- CCO readiness focuses on policy-to-evidence governance.
- This paper bridges to Guide 1 and Guide 2 but does not replace them.

---

### Source Grounding Note

Enterprise control context uses AUD-04 and AUD-07. AI governance context uses AI-01 and PRIV-03. Big Four sources BF-02 and BF-04 provide market context only. GAIC source truth provides Guide 1/Guide 2 boundaries.

---

### Author Synthesis Note

The CIO/CTO/CCO readiness split is author synthesis in this paper derived from R4C Guide planning and audit evidence architecture in this paper.

---

### Main Text

Enterprises do not become audit-ready by exporting logs at the end of an agentic workflow. Auditability must be designed into architecture, governance, review routines, and evidence ownership. This paper is not an implementation guide, but it defines the readiness questions leaders should ask before Guide 1 and Guide 2 translate the framework into practice.

For CIOs and CTOs, the first question is whether the system can produce evidence by design. Agent runtimes, orchestration layers, tool integrations, workflow engines, ticketing systems, approval systems, data stores, and logging platforms must be able to preserve the relationships that matter: work unit, authority, role, tool action, evidence pointer, outcome, exception, privacy treatment, and closure. If these relationships are not captured at runtime, auditability becomes after-the-fact reconstruction.

The second technology question is partitioning. Evidence should be separable by agent, tool, role, vendor, project, data class, lifecycle state, and review purpose. This requires

stable identifiers, metadata, access controls, retention logic, evidence export, and selective disclosure. It may also require integration with existing observability, security, data governance, and workflow systems.

The third technology question is change. Agentic systems change as models, prompts, tools, vendors, runtimes, policies, workflows, and memory stores change. Auditability requires substitution records, evidence continuity checks, regression evidence, reauthorization triggers, and revalidation boundaries. A system that loses evidence continuity after a runtime change is not auditability-ready for that scope.

For CCOs and governance leaders, the first question is policy-to-evidence translation. Policies should not only describe desired behavior; they should identify the evidence object that proves the policy was operationalized. A delegated authority policy should map to authority records. A human oversight policy should map to role and confirmation records. A remediation policy should map to closure evidence.

The second governance question is role operating model. The IIA Three Lines context is useful here as governance vocabulary, not as a this paper mandate. Management, risk/compliance, internal audit, privacy, security, legal, and technology teams need different evidence views and responsibilities. this paper helps name the evidence layer they must share.

The third governance question is review cadence. Agentic auditability is not a one-time readiness exercise. Evidence objects should be sampled, walked through, rechecked after change, reviewed after exceptions, and re-evaluated when authority, vendor, model, runtime, or privacy conditions change.

Guide 1 will later turn these questions into technical architecture planning. Guide 2 will later turn them into compliance operating-model routines. This paper should not absorb those details. It should define the evidence object layer that both guides will implement.

## Enterprise Readiness Crosswalk

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

|                                  |                                            |
|----------------------------------|--------------------------------------------|
| <b>Work unit identity</b>        |                                            |
| <b>CIO / CTO QUESTION</b>        | Can systems assign stable lifecycle IDs?   |
| <b>CCO / GOVERNANCE QUESTION</b> | Which work types require evidence?         |
| <b>EVIDENCE OBJECT</b>           | Work unit record                           |
| <b>GUIDE RELATION</b>            | Guide 1 architecture; Guide 2 policy scope |
| <b>BOUNDARY</b>                  | Not audit procedure                        |

**Authority**

|                                          |                                                                        |
|------------------------------------------|------------------------------------------------------------------------|
| <b>CIO / CTO<br/>QUESTION</b>            | Can delegated authority be enforced and logged as business scope?      |
| <b>CCO /<br/>GOVERNANCE<br/>QUESTION</b> | Which policies define authority, expiry, confirmation, and escalation? |
| <b>EVIDENCE OBJECT</b>                   | Delegated authority record                                             |
| <b>GUIDE RELATION</b>                    | Guide 1 controls; Guide 2 governance                                   |
| <b>BOUNDARY</b>                          | Not legal delegation proof                                             |

**Role mapping**

|                                          |                                                   |
|------------------------------------------|---------------------------------------------------|
| <b>CIO / CTO<br/>QUESTION</b>            | Can human/agent roles be captured separately?     |
| <b>CCO /<br/>GOVERNANCE<br/>QUESTION</b> | Who owns intent, review, acceptance, and closure? |
| <b>EVIDENCE OBJECT</b>                   | Responsibility map                                |
| <b>GUIDE RELATION</b>                    | Both guides                                       |
| <b>BOUNDARY</b>                          | Not legal liability assignment                    |

**Tool actions**

|                                          |                                                    |
|------------------------------------------|----------------------------------------------------|
| <b>CIO / CTO<br/>QUESTION</b>            | Can consequential tool actions generate evidence?  |
| <b>CCO /<br/>GOVERNANCE<br/>QUESTION</b> | Which tool actions require review or confirmation? |
| <b>EVIDENCE OBJECT</b>                   | Tool-action evidence                               |
| <b>GUIDE RELATION</b>                    | Guide 1 high relevance                             |
| <b>BOUNDARY</b>                          | No liability conclusion                            |

|                                  |                                                                           |
|----------------------------------|---------------------------------------------------------------------------|
| <b>Privacy</b>                   |                                                                           |
| <b>CIO / CTO QUESTION</b>        | Can evidence be minimized, redacted, retained, and selectively disclosed? |
| <b>CCO / GOVERNANCE QUESTION</b> | Which privacy rules and approvals govern review?                          |
| <b>EVIDENCE OBJECT</b>           | Disclosure and retention profile                                          |
| <b>GUIDE RELATION</b>            | Both guides                                                               |
| <b>BOUNDARY</b>                  | No legal advice                                                           |

|                                  |                                                       |
|----------------------------------|-------------------------------------------------------|
| <b>Exceptions</b>                |                                                       |
| <b>CIO / CTO QUESTION</b>        | Can exceptions link to lifecycle objects and closure? |
| <b>CCO / GOVERNANCE QUESTION</b> | Who governs dispute/remediation workflow?             |
| <b>EVIDENCE OBJECT</b>           | Exception/remediation record                          |
| <b>GUIDE RELATION</b>            | Both guides                                           |
| <b>BOUNDARY</b>                  | No legal remedy claim                                 |

|                                  |                                            |
|----------------------------------|--------------------------------------------|
| <b>Third-party review</b>        |                                            |
| <b>CIO / CTO QUESTION</b>        | Can evidence be exported/replayed safely?  |
| <b>CCO / GOVERNANCE QUESTION</b> | What can third parties claim or not claim? |
| <b>EVIDENCE OBJECT</b>           | Validation scope and boundary statement    |
| <b>GUIDE RELATION</b>            | Guide 1 export; Guide 2 review boundary    |
| <b>BOUNDARY</b>                  | No certification                           |

**Table note:** Author-synthesis crosswalk grounded in controls, governance, privacy, and GAIC Guide planning. It is not procurement guidance or readiness guarantee.

Cross-Links

- Guide 1 will own technical implementation details.
- Guide 2 will own compliance operating-model details.
- Appendix D can be used as a readiness conversation input.

---

## Boundary Note

This chapter does not claim that following this paper makes an enterprise compliant, audit-ready in a professional engagement, certified, regulator-approved, or procurement-preferred.

---

---

# 15. How Audit and Assurance Firms Can Use This Framework

**Chapter purpose:** Define cautious professional use without claiming endorsement, replacement methodology, assurance opinion, or standard-setting status. **Reader question:** How might audit and assurance firms use this paper responsibly?

---

## Key Claims

- Audit and assurance firms may use this paper as a conceptual object/evidence request and readiness framing tool.
- This framework can support discussion, scoping, evidence request design, walkthrough planning, and methodology exploration.
- This paper does not replace professional methodology, engagement criteria, independence, judgment, standards, or reporting.

---

## Source Grounding Note

Big Four context uses BF-02 and BF-03 as market context only. Assurance boundary uses AUD-03 and BOUND-03. AI audit practice context uses AUD-05 and AUD-06.

---

## Author Synthesis Note

The professional-use framing is author synthesis in this paper. It is intentionally cautious and non-endorsement based.

---

## Main Text

This framework is relevant to audit and assurance firms because agentic AI changes the evidence question. Firms already discuss AI assurance, trusted AI, audit transformation, technology-enabled assurance, governance, controls, and responsible AI. Agentic systems add another layer: lifecycle work that must be reconstructed across authority, responsibility, agents, tools, outcomes, exceptions, privacy, and remediation.

The safest professional use of this paper is as a framing tool. It can help teams ask what the review object is. It can help distinguish logs from audit evidence chains. It can help convert governance concerns into evidence requests. It can help structure walkthrough conversations. It can help identify whether an enterprise has evidence gaps before any formal assurance claim is discussed.

This paper may also help methodology teams think about agentic evidence architecture. A method team could use the Agentic Audit Object as a conceptual lens: What lifecycle work unit is under review? What authority record exists? What human role accepted the

outcome? What tool-action evidence exists? What exception state or remediation closure exists? What privacy treatment governs disclosure?

Internal audit teams may use this paper differently. They may use it to scope readiness reviews, design evidence request lists, test walkthrough templates, ask whether human oversight is evidence-backed, or evaluate whether remediation closure is reconstructable. Again, this use is readiness-oriented. It does not transform this paper into an internal audit standard.

Assurance teams must be especially careful. AARM's "Assurance-Ready" level does not mean assurance has been issued or is guaranteed. It means that an evidence architecture may be mature enough to support qualified assurance planning or review, subject to the firm's professional standards, criteria, engagement scope, independence, risk assessment, evidence evaluation, and reporting requirements.

This paper should never be represented as a Big Four methodology. It should never be described as endorsed by Deloitte, PwC, EY, KPMG, ISACA, IIA, IAASB, PCAOB, AICPA, ISO, NIST, OECD, EU institutions, Singapore agencies, ICO, or any source organization. Sources provide context and terminology. They do not approve this framework.

Professional Use Matrix

| Professional Use Matrix    |                                                                                  |                                                 |                                      |                          |
|----------------------------|----------------------------------------------------------------------------------|-------------------------------------------------|--------------------------------------|--------------------------|
| Use case                   | This frame-work can support                                                      | this paper can-not replace                      | Useful objects                       | Boundary                 |
| Readiness conversation     | Shared vocabu-lary for au-ditability gaps                                        | Professional judgment or en-gagement acceptance | AARM, evi-dence chain, ob-ject model | No assur-ance conclusion |
| Evidence re-request design | Categories for authority, roles, tools, out-comes, excep-tions, privacy, closure | Audit proce-dures or legal demands              | Appendix A request list              | Not formal procedure     |
| Walkthrough scoping        | Lifecycle recon-struction pattern                                                | Sampling methodology or engagement workplan     | Appendix B walk-through              | Illustrative only        |
| Internal audit review      | Auditability gap identification                                                  | Internal audit standards or fi-                 | MRO map-ping, clo-                   | No stan-dard-            |

| Use case                   | This frame-<br>work can<br>support       | this paper can-<br>not replace                                                     | Useful<br>objects                                         | Boundary                           |
|----------------------------|------------------------------------------|------------------------------------------------------------------------------------|-----------------------------------------------------------|------------------------------------|
|                            |                                          | nal reports                                                                        | sure<br>checklist                                         | setting                            |
| Methodology<br>exploration | Object model<br>for agentic<br>work      | Firm methodol-<br>ogy or regulato-<br>ry requirements                              | Agentic<br>Audit<br>Object,<br>Audit<br>Evidence<br>Chain | No Big<br>Four<br>endorse-<br>ment |
| Assurance<br>planning      | Evidence archi-<br>tecture<br>discussion | Assurance cri-<br>teria, indepen-<br>dence, evi-<br>dence evalua-<br>tion, opinion | AARM<br>L4/L5<br>boundary                                 | No assur-<br>ance<br>opinion       |

**Table note:** Professional-use matrix is author synthesis grounded in assurance and AI audit practice sources. It is not a firm methodology, engagement template, or endorsement claim.

Cross-Links

- Chapter 1 defines scope and boundary.
- Chapter 12 defines third-party validation boundaries.
- Appendix F provides safe language.

Boundary Note

This chapter does not claim that audit or assurance firms need this paper, endorse this paper, have adopted this paper, or should replace their methods with this paper.

---

## 16. Conclusion

**Chapter purpose:** Close this paper's argument and bridge to Guide 1, Guide 2, and the later insurability white paper without publication or insurability overclaim. **Reader question:** What must exist before agentic assurance and insurability conversations can mature?

---

### Key Claims

- Auditability begins when agentic lifecycle work becomes reconstructable.
- Logs record activity; evidence chains support responsibility review.
- This paper prepares the path for Guide 1, Guide 2, and the later insurability white paper.

---

### Source Grounding Note

Audit evidence and assurance boundary language uses AUD-01 and AUD-03. AI governance context uses AI-01. GAIC source truth supports the Global AI Compliance White Paper 2026 / this paper / later insurability white paper object-chain relationship and companion-paper boundary.

---

### Author Synthesis Note

The Compliance Object → Audit Evidence Object → Insurable Risk Object chain is GAIC/author synthesis in this paper. The later insurability white paper source research is deferred.

---

### Main Text

Agentic AI changes the auditability problem because agentic systems do not merely produce outputs. They perform lifecycle work. They act through roles, tools, delegated authority, handoffs, privacy boundaries, exceptions, and remediation workflows. If that work cannot be reconstructed, the system may be impressive, observable, and even well-governed in parts, but it is not audit-ready in the sense this paper defines.

The core thesis therefore remains simple: AI agent auditability cannot be built on logs alone. Logs record activity; evidence chains support responsibility review. The difference is the lifecycle responsibility layer. A log may show that an event occurred. An Audit Evidence Chain shows how that event relates to authority, human role, agent role, tool action, evidence pointer, accepted outcome, exception, privacy treatment, and closure.

This paper's contribution is the Audit Evidence Object layer. The Global AI Compliance White Paper 2026 defined the compliance object layer through Agentic Lifecycle Governance and MROs. This paper translates that layer into evidence requests, walkthroughs, partitioning rules, privacy-aware disclosure, validation boundaries, and readiness levels. It defines Agentic AI Auditability as the ability to reconstruct, test, and evidence agentic lifecycle work across authority, responsibility, tools, outcomes, exceptions, and remediation.

The paper also sets limits. It does not issue assurance. It does not certify systems. It does not prove legal compliance. It does not bind regulators. It does not replace professional audit or assurance methodology. It does not rank vendors. It does not require MPLP. It does not make Validation Lab a certification authority.

These limits are not disclaimers at the edge of the work; they are part of the architecture. Auditability is strongest when the evidence object and the claim boundary are both explicit.

For enterprises, the next step is design. CIOs and CTOs need evidence architecture that captures lifecycle responsibility by design, not after-the-fact logs by accident. CCOs and governance leaders need policy-to-evidence operating models that turn authority, oversight, privacy, exception, and remediation policies into reviewable objects. Internal audit and risk teams need walkthrough structures that test reconstructability before assurance claims appear.

For audit and assurance firms, the next step is method dialogue. This framework can support object-model discussions, readiness scoping, evidence request framing, and lifecycle walkthrough design. It cannot replace professional standards, engagement criteria, independence, evidence evaluation, or reporting.

For the series, this paper completes the second layer of the 3+3 front half. The Global AI Compliance White Paper 2026 established the compliance object layer. This paper establishes the audit evidence object layer. Guide 1 will translate auditability into technical architecture. Guide 2 will translate auditability into compliance operating model. The later insurability white paper will later ask what evidence and governance structures are needed before insurable risk can be evaluated.

The object chain is:

| Series layer                                         | Object layer          | Question                                                               | Boundary                                                               |
|------------------------------------------------------|-----------------------|------------------------------------------------------------------------|------------------------------------------------------------------------|
| Global AI Compliance White Paper 2026                | Compliance Object     | What lifecycle responsibility objects does agentic governance require? | Not legal compliance proof                                             |
| Agentic AI Auditability & Assurance White Paper 2026 | Audit Evidence Object | Can agentic lifecycle work be reconstructed through responsi-          | Not audit stan-<br>dard, certifica-<br>tion, or assur-<br>ance opinion |

| Series layer                   | Object layer                    | Question                                                    | Boundary                         |
|--------------------------------|---------------------------------|-------------------------------------------------------------|----------------------------------|
|                                |                                 | bility-linked evidence chains?                              |                                  |
| Later insurability white paper | Insurable Risk Object           | What risk evidence may later support insurability analysis? | Deferred; no insurance guarantee |
| Guide 1                        | Technical implementation path   | How can systems capture audit-ready evidence by design?     | Future guide, not drafted here   |
| Guide 2                        | Compliance operating model path | How can governance translate policy into evidence?          | Future guide, not drafted here   |

**Table note:** GAIC/author-synthesis table. It preserves companion-paper sequencing and does not claim conclusions for the later insurability white paper.

Auditability begins when agentic lifecycle work becomes reconstructable. Logs record activity; evidence chains support responsibility review. This paper prepares the path for Guide 1, Guide 2, and the later insurability white paper by defining the evidence object layer that must exist before agentic assurance or insurability conversations can be serious.

### Cross-Links

- Appendix F preserves conclusion-safe boundary language.
- Guide 1 and Guide 2 remain future practitioner guides.
- The later insurability white paper remains a future white paper with separate source research.

### Boundary Note

This chapter does not claim this paper is published, final, sealed, regulator-approved, adopted by firms, or sufficient for assurance, certification, legal compliance, procurement, or insurance.

# Appendix A — Agentic Audit Evidence Request List

**Purpose:** Provide a structured evidence request catalog for agentic systems. **Intended reader:** Audit/assurance teams, internal audit, AI governance, CIO/CTO/CCO teams, technology risk teams, and evidence architecture owners.

## Source Grounding

Audit evidence and control vocabulary: AUD-01, AUD-02, AUD-04, AUD-05, AUD-06. GAIC dependencies: MRO-01 through MRO-16, Evidence-Based Validation Pattern, Validation Lab boundary.

## Boundary

This appendix is an evidence request architecture. It is not a formal audit procedure, legal demand, certification criterion, procurement checklist, assurance engagement plan, or proof of legal compliance.

## Evidence Request Catalog

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

|                           |                                                                  |
|---------------------------|------------------------------------------------------------------|
| <b>ER-01</b>              |                                                                  |
| <b>LIFECYCLE STAGE</b>    | Work initiation                                                  |
| <b>EVIDENCE REQUESTED</b> | Work unit ID, task intent, request source, scope, excluded scope |
| <b>RELATED MRO</b>        | MRO-08                                                           |
| <b>OWNER</b>              | Request owner                                                    |
| <b>PRIVACY TREATMENT</b>  | Record only necessary purpose/context fields                     |
| <b>BOUNDARY NOTE</b>      | Does not prove business appropriateness                          |

**ER-02**

|                           |                                                                                  |
|---------------------------|----------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Role mapping                                                                     |
| <b>EVIDENCE REQUESTED</b> | Human role for intent, authorization, review, acceptance, exception, remediation |
| <b>RELATED MRO</b>        | MRO-01                                                                           |
| <b>OWNER</b>              | Governance/control owner                                                         |
| <b>PRIVACY TREATMENT</b>  | Role data minimized to function and authority                                    |
| <b>BOUNDARY NOTE</b>      | Does not assign legal liability                                                  |

**ER-03**

|                           |                                                                                    |
|---------------------------|------------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Delegated authority                                                                |
| <b>EVIDENCE REQUESTED</b> | Authority grant, conditions, expiry, risk class, confirmation and escalation rules |
| <b>RELATED MRO</b>        | MRO-02                                                                             |
| <b>OWNER</b>              | Delegating role                                                                    |
| <b>PRIVACY TREATMENT</b>  | Restrict authority evidence to review scope                                        |
| <b>BOUNDARY NOTE</b>      | Not legal delegation proof                                                         |

**ER-04**

|                           |                                                                                 |
|---------------------------|---------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Agent role                                                                      |
| <b>EVIDENCE REQUESTED</b> | Agent role, capability boundary, constraints, evidence duties, escalation route |
| <b>RELATED MRO</b>        | MRO-03                                                                          |
| <b>OWNER</b>              | Platform owner                                                                  |
| <b>PRIVACY TREATMENT</b>  | Avoid storing unnecessary prompt/user data                                      |
| <b>BOUNDARY NOTE</b>      | Agent is not a human role                                                       |

**ER-05**

|                           |                                                                                 |
|---------------------------|---------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Accepted outcome                                                                |
| <b>EVIDENCE REQUESTED</b> | Acceptance criteria, reviewer role, accepted/rejected/disputed/remediated state |
| <b>RELATED MRO</b>        | MRO-04                                                                          |
| <b>OWNER</b>              | Review owner                                                                    |
| <b>PRIVACY TREATMENT</b>  | Redact sensitive output fields where possible                                   |
| <b>BOUNDARY NOTE</b>      | Acceptance is not compliance proof                                              |

**ER-06**

|                           |                                                                                          |
|---------------------------|------------------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Tool action                                                                              |
| <b>EVIDENCE REQUESTED</b> | Tool ID, action type, target system, affected data/process, reversibility, rollback path |
| <b>RELATED MRO</b>        | MRO-05                                                                                   |
| <b>OWNER</b>              | Tool-action owner                                                                        |
| <b>PRIVACY TREATMENT</b>  | Classify affected data and redact as needed                                              |
| <b>BOUNDARY NOTE</b>      | Not legal liability finding                                                              |

**ER-07**

|                           |                                                                                      |
|---------------------------|--------------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Responsibility transfer                                                              |
| <b>EVIDENCE REQUESTED</b> | Source agent, target agent, transferred scope, retained scope, inherited constraints |
| <b>RELATED MRO</b>        | MRO-06                                                                               |
| <b>OWNER</b>              | Process owner                                                                        |
| <b>PRIVACY TREATMENT</b>  | Partition by work unit and agent role                                                |
| <b>BOUNDARY NOTE</b>      | Not legal responsibility transfer                                                    |

**ER-08**

|                           |                                                                                                |
|---------------------------|------------------------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Authority drift                                                                                |
| <b>EVIDENCE REQUESTED</b> | Baseline authority, observed action, drift trigger, escalation, stop/downgrade/reauthorization |
| <b>RELATED MRO</b>        | MRO-07                                                                                         |
| <b>OWNER</b>              | Escalation owner                                                                               |
| <b>PRIVACY TREATMENT</b>  | Limit exposure to event-relevant fields                                                        |
| <b>BOUNDARY NOTE</b>      | Not every drift is legal breach                                                                |

**ER-09**

|                           |                                                                                     |
|---------------------------|-------------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Evidence partitioning                                                               |
| <b>EVIDENCE REQUESTED</b> | Partition keys, links, evidence pointers, integrity hashes, access/disclosure rules |
| <b>RELATED MRO</b>        | MRO-08                                                                              |
| <b>OWNER</b>              | Evidence owner                                                                      |
| <b>PRIVACY TREATMENT</b>  | Use scoped access and redaction profiles                                            |
| <b>BOUNDARY NOTE</b>      | Not blanket retention                                                               |

**ER-10**

|                           |                                                                                         |
|---------------------------|-----------------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Cross-project reuse                                                                     |
| <b>EVIDENCE REQUESTED</b> | Source context, target context, reset validation, privacy review, authorization refresh |
| <b>RELATED MRO</b>        | MRO-09                                                                                  |
| <b>OWNER</b>              | Reuse owner                                                                             |
| <b>PRIVACY TREATMENT</b>  | Revalidate data categories and retention                                                |
| <b>BOUNDARY NOTE</b>      | Not legal reuse clearance                                                               |

**ER-11**

|                           |                                                                                        |
|---------------------------|----------------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Privacy lifecycle                                                                      |
| <b>EVIDENCE REQUESTED</b> | Data flow, processing purpose, retention rule, rights workflow, minimization treatment |
| <b>RELATED MRO</b>        | MRO-10                                                                                 |
| <b>OWNER</b>              | Privacy owner                                                                          |
| <b>PRIVACY TREATMENT</b>  | Apply privacy-by-design review                                                         |
| <b>BOUNDARY NOTE</b>      | Not legal advice                                                                       |

**ER-12**

|                           |                                                                                                   |
|---------------------------|---------------------------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Third-party validation                                                                            |
| <b>EVIDENCE REQUESTED</b> | Ruleset identity, evidence pointer, disclosure/redaction profile, verdict hash, reviewer boundary |
| <b>RELATED MRO</b>        | MRO-11                                                                                            |
| <b>OWNER</b>              | Validation/review owner                                                                           |
| <b>PRIVACY TREATMENT</b>  | Use selective disclosure package                                                                  |
| <b>BOUNDARY NOTE</b>      | Not certification                                                                                 |

**ER-13**

|                           |                                                                                    |
|---------------------------|------------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Evidence minimization                                                              |
| <b>EVIDENCE REQUESTED</b> | Minimization rule, disclosure justification, redaction profile, access log, expiry |
| <b>RELATED MRO</b>        | MRO-12                                                                             |
| <b>OWNER</b>              | Evidence/privacy owner                                                             |
| <b>PRIVACY TREATMENT</b>  | Minimize and document residual exposure                                            |
| <b>BOUNDARY NOTE</b>      | Not privacy compliance proof                                                       |

**ER-14**

|                           |                                                                                     |
|---------------------------|-------------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Rights vs retention                                                                 |
| <b>EVIDENCE REQUESTED</b> | Request type, retention basis, hold status, deletion/redaction action, review owner |
| <b>RELATED MRO</b>        | MRO-13                                                                              |
| <b>OWNER</b>              | Privacy/legal owner                                                                 |
| <b>PRIVACY TREATMENT</b>  | Requires jurisdiction-specific review                                               |
| <b>BOUNDARY NOTE</b>      | Not legal interpretation                                                            |

**ER-15**

|                           |                                                                                   |
|---------------------------|-----------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Processor chain                                                                   |
| <b>EVIDENCE REQUESTED</b> | Processor role, subprocessor chain, data processing pointer, responsibility owner |
| <b>RELATED MRO</b>        | MRO-14                                                                            |
| <b>OWNER</b>              | Vendor/privacy owner                                                              |
| <b>PRIVACY TREATMENT</b>  | Restrict processor evidence by need-to-know                                       |
| <b>BOUNDARY NOTE</b>      | Not vendor ranking                                                                |

**ER-16**

|                           |                                                                                   |
|---------------------------|-----------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Substitution conformance                                                          |
| <b>EVIDENCE REQUESTED</b> | Prior/new model, vendor, runtime, tool, evidence-integrity check, reauthorization |
| <b>RELATED MRO</b>        | MRO-15                                                                            |
| <b>OWNER</b>              | Change owner                                                                      |
| <b>PRIVACY TREATMENT</b>  | Preserve evidence continuity without over-retention                               |
| <b>BOUNDARY NOTE</b>      | Not conformance certification                                                     |

**ER-17**

|                           |                                                                                                 |
|---------------------------|-------------------------------------------------------------------------------------------------|
| <b>LIFECYCLE STAGE</b>    | Incident/remediation closure                                                                    |
| <b>EVIDENCE REQUESTED</b> | Event, affected outcome, owner, corrective action, recheck, closure acceptance, reopen criteria |
| <b>RELATED MRO</b>        | MRO-16                                                                                          |
| <b>OWNER</b>              | Remediation owner                                                                               |
| <b>PRIVACY TREATMENT</b>  | Redact sensitive incident details as needed                                                     |
| <b>BOUNDARY NOTE</b>      | Not legal closure                                                                               |

**Guide Relationship**

Guide 1 should translate this appendix into system output requirements, evidence export design, runtime metadata, and workflow architecture. Guide 2 should translate it into policy-to-evidence governance, audit preparation, review cadence, and closure routines.

**Placement Recommendation**

Retain a condensed version in the main artifact. A fuller operational checklist can be deferred to Guide 1 and Guide 2.

# Appendix B – Agentic System Walkthrough Template

**Purpose:** Provide a repeatable walkthrough structure for a single agentic lifecycle run.  
**Intended reader:** Internal audit, technology risk, AI governance committees, architecture review teams, platform owners, and compliance engineering teams.

## Source Grounding

Control and audit practice context: AUD-04, AUD-05, AUD-06. Provenance and observability context: EVID-01, EVID-02. GAIC dependencies: MRO-01 through MRO-08 and MRO-16.

## Boundary

This template is illustrative. It is not a formal audit procedure, assurance engagement program, legal compliance test, certification path, or regulator-approved walkthrough.

## Walkthrough Fields

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

|                         |                                                       |
|-------------------------|-------------------------------------------------------|
| <b>Work unit ID</b>     |                                                       |
| <b>DESCRIPTION</b>      | Stable identifier for the lifecycle work under review |
| <b>EVIDENCE POINTER</b> | Workflow/run/task record                              |
| <b>RELATED MRO</b>      | MRO-08                                                |
| <b>BOUNDARY NOTE</b>    | Not mandatory schema                                  |

|                         |                                                                |
|-------------------------|----------------------------------------------------------------|
| <b>Intent and scope</b> |                                                                |
| <b>DESCRIPTION</b>      | Purpose, initiating request, in-scope and out-of-scope actions |
| <b>EVIDENCE POINTER</b> | Request record, policy reference                               |
| <b>RELATED MRO</b>      | MRO-01, MRO-08                                                 |
| <b>BOUNDARY NOTE</b>    | Does not prove appropriateness                                 |

**Initiator role**

|                         |                                              |
|-------------------------|----------------------------------------------|
| <b>DESCRIPTION</b>      | Human or system role that initiated the work |
| <b>EVIDENCE POINTER</b> | Role map                                     |
| <b>RELATED MRO</b>      | MRO-01                                       |
| <b>BOUNDARY NOTE</b>    | Not legal liability assignment               |

**Delegated authority**

|                         |                                                               |
|-------------------------|---------------------------------------------------------------|
| <b>DESCRIPTION</b>      | Authority grant, conditions, expiry, confirmation, escalation |
| <b>EVIDENCE POINTER</b> | Authority record                                              |
| <b>RELATED MRO</b>      | MRO-02                                                        |
| <b>BOUNDARY NOTE</b>    | Not legal delegation proof                                    |

**Agent role**

|                         |                                                               |
|-------------------------|---------------------------------------------------------------|
| <b>DESCRIPTION</b>      | Agent identity, role boundary, constraints, capability limits |
| <b>EVIDENCE POINTER</b> | Agent registry/config record                                  |
| <b>RELATED MRO</b>      | MRO-03                                                        |
| <b>BOUNDARY NOTE</b>    | Agent role is not human role                                  |

**Tool/system used**

|                         |                                                |
|-------------------------|------------------------------------------------|
| <b>DESCRIPTION</b>      | Tool, external system, API, processor, runtime |
| <b>EVIDENCE POINTER</b> | Tool-action evidence                           |
| <b>RELATED MRO</b>      | MRO-05, MRO-14                                 |
| <b>BOUNDARY NOTE</b>    | Not vendor ranking                             |

**Evidence pointers**

|                         |                                                           |
|-------------------------|-----------------------------------------------------------|
| <b>DESCRIPTION</b>      | Logs, traces, approvals, tickets, records, hash manifests |
| <b>EVIDENCE POINTER</b> | Evidence chain references                                 |
| <b>RELATED MRO</b>      | MRO-08, MRO-12                                            |
| <b>BOUNDARY NOTE</b>    | Not blanket disclosure                                    |

**Human review**

|                         |                                               |
|-------------------------|-----------------------------------------------|
| <b>DESCRIPTION</b>      | Reviewer role, review criteria, review result |
| <b>EVIDENCE POINTER</b> | Review record                                 |
| <b>RELATED MRO</b>      | MRO-01, MRO-04                                |
| <b>BOUNDARY NOTE</b>    | Not assurance opinion                         |

**Accepted outcome**

|                         |                                                      |
|-------------------------|------------------------------------------------------|
| <b>DESCRIPTION</b>      | Produced/accepted/rejected/disputed/remediated state |
| <b>EVIDENCE POINTER</b> | Outcome record                                       |
| <b>RELATED MRO</b>      | MRO-04                                               |
| <b>BOUNDARY NOTE</b>    | Not compliance proof                                 |

**Exception state**

|                         |                                                          |
|-------------------------|----------------------------------------------------------|
| <b>DESCRIPTION</b>      | Drift, dispute, failure, privacy event, tool side effect |
| <b>EVIDENCE POINTER</b> | Exception record                                         |
| <b>RELATED MRO</b>      | MRO-07, MRO-16                                           |
| <b>BOUNDARY NOTE</b>    | Not legal breach finding                                 |

**Remediation closure**

|                         |                                                                 |
|-------------------------|-----------------------------------------------------------------|
| <b>DESCRIPTION</b>      | Corrective action, recheck, closure acceptance, reopen criteria |
| <b>EVIDENCE POINTER</b> | Remediation record                                              |
| <b>RELATED MRO</b>      | MRO-16                                                          |
| <b>BOUNDARY NOTE</b>    | Not legal settlement                                            |

**Privacy treatment**

|                         |                                                         |
|-------------------------|---------------------------------------------------------|
| <b>DESCRIPTION</b>      | Data category, redaction, retention, disclosure profile |
| <b>EVIDENCE POINTER</b> | Privacy/disclosure record                               |
| <b>RELATED MRO</b>      | MRO-10 to MRO-13                                        |
| <b>BOUNDARY NOTE</b>    | Not legal advice                                        |

**Third-party review boundary**

|                  |                                                   |
|------------------|---------------------------------------------------|
| DESCRIPTION      | Reviewer role, ruleset, scope, result, limitation |
| EVIDENCE POINTER | Validation/review record                          |
| RELATED MRO      | MRO-11                                            |
| BOUNDARY NOTE    | Not certification                                 |

**Walkthrough Sequence**

1. Identify the work unit and review scope.
2. Confirm the authority boundary and role ownership.
3. Reconstruct agent actions, tool actions, and handoffs.
4. Link logs/traces to evidence pointers and lifecycle objects.
5. Verify accepted outcome, dispute, or rejection state.
6. Review exceptions and remediation closure.
7. Confirm privacy/selective disclosure treatment.
8. Record what the walkthrough supports and what it does not prove.

**Guide Relationship**

Guide 1 should use this template as a technical evidence export and walkthrough architecture pattern. Guide 2 should use it as an audit preparation and governance review workflow.

**Placement Recommendation**

Include in the main artifact if space allows. Expand implementation details in Guide 1 and governance workflow details in Guide 2.

# Appendix C — MRO-to-Audit-Evidence Mapping

**Purpose:** Preserve the full mapping from GAIC MROs to audit evidence objects in this paper. **Intended reader:** Audit/assurance method teams, compliance engineering, internal audit, AI governance, Guide 1/2 authors.

## Source Grounding

GAIC source truth defines MRO-01 through MRO-16. Audit evidence and control language is grounded in AUD-01, AUD-05, and AI-01. This appendix is GAIC-derived plus author synthesis in this paper.

## Boundary

MROs are proposed GAIC governance objects. They are not legal mandates, audit standards, certification criteria, procurement criteria, or legal compliance proof.

## Full Mapping

| Full Mapping                                                                                                        |                                                                                  |
|---------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------|
| Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces. |                                                                                  |
| <b>MRO-01 Human Role to MAS Responsibility Mapping</b>                                                              |                                                                                  |
| <b>EVIDENCE OBJECT</b>                                                                                              | Human-role responsibility map                                                    |
| <b>AUDIT QUESTION</b>                                                                                               | Who owned intent, authorization, review, acceptance, exception, and remediation? |
| <b>EVIDENCE REQUEST</b>                                                                                             | Role map with lifecycle stage and evidence duty                                  |
| <b>AARM DIMENSION</b>                                                                                               | Human / agent responsibility mapping                                             |
| <b>GUIDE 1 RELATION</b>                                                                                             | Capture roles in workflows and evidence metadata                                 |
| <b>GUIDE 2 RELATION</b>                                                                                             | Define role accountability routines                                              |
| <b>BOUNDARY</b>                                                                                                     | No legal liability assignment                                                    |

**MRO-02 Delegated Authority Boundary**

|                         |                                                                    |
|-------------------------|--------------------------------------------------------------------|
| <b>EVIDENCE OBJECT</b>  | Delegated authority record                                         |
| <b>AUDIT QUESTION</b>   | Was the agent authorized for the action under scope and condition? |
| <b>EVIDENCE REQUEST</b> | Authority grant, expiry, risk class, escalation, reauthorization   |
| <b>AARM DIMENSION</b>   | Authority traceability                                             |
| <b>GUIDE 1 RELATION</b> | Implement authority checks and evidence records                    |
| <b>GUIDE 2 RELATION</b> | Govern delegation policy and approval                              |
| <b>BOUNDARY</b>         | Not IAM-only and not legal proof                                   |

**MRO-03 Agent Role is not Human Role**

|                         |                                                         |
|-------------------------|---------------------------------------------------------|
| <b>EVIDENCE OBJECT</b>  | Agent role boundary object                              |
| <b>AUDIT QUESTION</b>   | Was the agent role separated from human accountability? |
| <b>EVIDENCE REQUEST</b> | Agent role, constraints, owner, escalation              |
| <b>AARM DIMENSION</b>   | Human / agent responsibility mapping                    |
| <b>GUIDE 1 RELATION</b> | Define agent role registry and capability boundaries    |
| <b>GUIDE 2 RELATION</b> | Govern agent role ownership                             |
| <b>BOUNDARY</b>         | Do not anthropomorphize                                 |

**MRO-04 Accepted Outcome Compliance**

|                         |                                                             |
|-------------------------|-------------------------------------------------------------|
| <b>EVIDENCE OBJECT</b>  | Accepted outcome record                                     |
| <b>AUDIT QUESTION</b>   | Was the output accepted, rejected, disputed, or remediated? |
| <b>EVIDENCE REQUEST</b> | Acceptance criteria, reviewer, outcome state, evidence link |
| <b>AARM DIMENSION</b>   | Accepted outcome evidence                                   |
| <b>GUIDE 1 RELATION</b> | Capture outcome state transitions                           |
| <b>GUIDE 2 RELATION</b> | Govern review and acceptance                                |
| <b>BOUNDARY</b>         | Not compliance proof                                        |

**MRO-05 Tool-Action Liability Boundary**

|                         |                                                                                   |
|-------------------------|-----------------------------------------------------------------------------------|
| <b>EVIDENCE OBJECT</b>  | Tool-action evidence object                                                       |
| <b>AUDIT QUESTION</b>   | What tool created consequence, under what authority, and with what reversibility? |
| <b>EVIDENCE REQUEST</b> | Tool, target, action, affected data, rollback, owner                              |
| <b>AARM DIMENSION</b>   | Agent/tool partitioning                                                           |
| <b>GUIDE 1 RELATION</b> | Instrument tool-action evidence                                                   |
| <b>GUIDE 2 RELATION</b> | Define tool-action approval and review                                            |
| <b>BOUNDARY</b>         | No legal liability conclusion                                                     |

**MRO-06 Responsibility Transfer Across Agents**

|                         |                                                                        |
|-------------------------|------------------------------------------------------------------------|
| <b>EVIDENCE OBJECT</b>  | Responsibility transfer record                                         |
| <b>AUDIT QUESTION</b>   | Did responsibility transfer or remain with source role?                |
| <b>EVIDENCE REQUEST</b> | Source/target agent, transferred/retained scope, inherited constraints |
| <b>AARM DIMENSION</b>   | Human / agent responsibility mapping                                   |
| <b>GUIDE 1 RELATION</b> | Capture handoff records                                                |
| <b>GUIDE 2 RELATION</b> | Govern transfer and escalation                                         |
| <b>BOUNDARY</b>         | Not legal transfer                                                     |

**MRO-07 Authority Drift**

|                         |                                                         |
|-------------------------|---------------------------------------------------------|
| <b>EVIDENCE OBJECT</b>  | Authority drift exception record                        |
| <b>AUDIT QUESTION</b>   | Did observed action exceed delegated authority?         |
| <b>EVIDENCE REQUEST</b> | Baseline, observed action, trigger, escalation, closure |
| <b>AARM DIMENSION</b>   | Authority traceability; exception traceability          |
| <b>GUIDE 1 RELATION</b> | Detect and record drift                                 |
| <b>GUIDE 2 RELATION</b> | Govern drift escalation                                 |
| <b>BOUNDARY</b>         | Not automatic legal breach                              |

**MRO-08 MAS Evidence Partitioning**

|                         |                                                                |
|-------------------------|----------------------------------------------------------------|
| <b>EVIDENCE OBJECT</b>  | Partitioned evidence chain                                     |
| <b>AUDIT QUESTION</b>   | Can evidence be separated and linked for review?               |
| <b>EVIDENCE REQUEST</b> | Partition keys, links, hashes, evidence pointers, access rules |
| <b>AARM DIMENSION</b>   | Evidence sufficiency; agent/tool partitioning                  |
| <b>GUIDE 1 RELATION</b> | Build partitioned evidence architecture                        |
| <b>GUIDE 2 RELATION</b> | Govern evidence access                                         |
| <b>BOUNDARY</b>         | No blanket retention                                           |

**MRO-09 Cross-Project Reuse Compliance**

|                         |                                                                                |
|-------------------------|--------------------------------------------------------------------------------|
| <b>EVIDENCE OBJECT</b>  | Reuse context validation record                                                |
| <b>AUDIT QUESTION</b>   | Was reused agent/workflow context reset and revalidated?                       |
| <b>EVIDENCE REQUEST</b> | Source/target context, reset validation, privacy review, authorization refresh |
| <b>AARM DIMENSION</b>   | Agent/tool partitioning                                                        |
| <b>GUIDE 1 RELATION</b> | Implement reuse validation checks                                              |
| <b>GUIDE 2 RELATION</b> | Govern reuse policy                                                            |
| <b>BOUNDARY</b>         | No legal reuse clearance                                                       |

**MRO-10 Privacy / GDPR Lifecycle Mapping**

|                         |                                                                   |
|-------------------------|-------------------------------------------------------------------|
| <b>EVIDENCE OBJECT</b>  | Privacy lifecycle evidence object                                 |
| <b>AUDIT QUESTION</b>   | How does lifecycle work map to data processing and retention?     |
| <b>EVIDENCE REQUEST</b> | Data flow, purpose, retention rule, rights workflow, minimization |
| <b>AARM DIMENSION</b>   | Privacy and selective disclosure                                  |
| <b>GUIDE 1 RELATION</b> | Capture privacy metadata                                          |
| <b>GUIDE 2 RELATION</b> | Govern privacy review                                             |
| <b>BOUNDARY</b>         | Not legal advice                                                  |

**MRO-11 Privacy-Preserving Third-Party Validation**

|                         |                                                                       |
|-------------------------|-----------------------------------------------------------------------|
| <b>EVIDENCE OBJECT</b>  | Validation disclosure profile and verdict record                      |
| <b>AUDIT QUESTION</b>   | Can third-party review happen without raw over-disclosure?            |
| <b>EVIDENCE REQUEST</b> | Ruleset, redaction/disclosure profile, evidence pointer, verdict hash |
| <b>AARM DIMENSION</b>   | Third-party review boundary; privacy                                  |
| <b>GUIDE 1 RELATION</b> | Support export/replay packages                                        |
| <b>GUIDE 2 RELATION</b> | Govern third-party review                                             |
| <b>BOUNDARY</b>         | Not certification                                                     |

**MRO-12 Evidence Minimization and Selective Disclosure**

|                         |                                                                            |
|-------------------------|----------------------------------------------------------------------------|
| <b>EVIDENCE OBJECT</b>  | Evidence minimization and disclosure object                                |
| <b>AUDIT QUESTION</b>   | Is disclosed evidence limited to review purpose?                           |
| <b>EVIDENCE REQUEST</b> | Minimization rule, disclosure justification, redaction profile, access log |
| <b>AARM DIMENSION</b>   | Privacy and selective disclosure                                           |
| <b>GUIDE 1 RELATION</b> | Build redaction/selective disclosure                                       |
| <b>GUIDE 2 RELATION</b> | Govern disclosure approvals                                                |
| <b>BOUNDARY</b>         | Not privacy compliance proof                                               |

**MRO-13 Data Subject Rights vs Evidence Retention**

|                         |                                                                  |
|-------------------------|------------------------------------------------------------------|
| <b>EVIDENCE OBJECT</b>  | Rights-retention reconciliation record                           |
| <b>AUDIT QUESTION</b>   | How are rights requests reconciled with evidence retention?      |
| <b>EVIDENCE REQUEST</b> | Request, retention basis, hold status, redaction/deletion action |
| <b>AARM DIMENSION</b>   | Privacy and selective disclosure                                 |
| <b>GUIDE 1 RELATION</b> | Support retention/rights workflows                               |
| <b>GUIDE 2 RELATION</b> | Govern legal/privacy review                                      |
| <b>BOUNDARY</b>         | Requires counsel                                                 |

**MRO-14 Third-Party Processor / Subprocessor Chain**

|                         |                                                                                  |
|-------------------------|----------------------------------------------------------------------------------|
| <b>EVIDENCE OBJECT</b>  | Processor/subprocessor chain evidence                                            |
| <b>AUDIT QUESTION</b>   | Which third parties processed or affected lifecycle work?                        |
| <b>EVIDENCE REQUEST</b> | Processor role, subprocessor chain, data agreement pointer, responsibility owner |
| <b>AARM DIMENSION</b>   | Agent/tool partitioning                                                          |
| <b>GUIDE 1 RELATION</b> | Capture vendor/processor evidence                                                |
| <b>GUIDE 2 RELATION</b> | Govern processor chain                                                           |
| <b>BOUNDARY</b>         | No vendor ranking                                                                |

**MRO-15 Vendor / Model / Runtime Substitution Conformance**

|                         |                                                                     |
|-------------------------|---------------------------------------------------------------------|
| <b>EVIDENCE OBJECT</b>  | Substitution conformance record                                     |
| <b>AUDIT QUESTION</b>   | Did evidence continuity survive component change?                   |
| <b>EVIDENCE REQUEST</b> | Prior/new component, integrity check, reauthorization, revalidation |
| <b>AARM DIMENSION</b>   | Agent/tool partitioning; evidence sufficiency                       |
| <b>GUIDE 1 RELATION</b> | Implement substitution evidence controls                            |
| <b>GUIDE 2 RELATION</b> | Govern change review                                                |
| <b>BOUNDARY</b>         | Not conformance certification                                       |

**MRO-16 Incident, Dispute, and Remediation Closure**

|                         |                                                                     |
|-------------------------|---------------------------------------------------------------------|
| <b>EVIDENCE OBJECT</b>  | Incident/remediation closure record                                 |
| <b>AUDIT QUESTION</b>   | Did exception or dispute close with evidence?                       |
| <b>EVIDENCE REQUEST</b> | Event, affected outcome, owner, corrective action, recheck, closure |
| <b>AARM DIMENSION</b>   | Exception traceability; remediation closure                         |
| <b>GUIDE 1 RELATION</b> | Build closure workflow evidence                                     |
| <b>GUIDE 2 RELATION</b> | Govern remediation closure                                          |
| <b>BOUNDARY</b>         | Not legal closure                                                   |

**Guide Relationship**

Guide 1 should treat this appendix as an implementation-neutral architecture checklist.  
Guide 2 should treat it as a policy-to-evidence and compliance operating model checklist.

---

## Placement Recommendation

Include full mapping in the this paper appendix. Consider a shorter summary table in Chapter 6.

---

# Appendix D — Auditability Readiness Matrix

**Purpose:** Provide AARM dimension and readiness-level crosswalk. **Intended reader:** Audit readiness teams, CIO/CTO/CCO leaders, internal audit, governance committees, platform teams.

## Source Grounding

Audit evidence: AUD-01. AI audit practice: AUD-05, AUD-06. AI governance: AI-01. Boundary: BOUND-03. GAIC dependency: R4B/R1 AARM baseline, MROs, ALCS context, Validation Lab boundary.

## Boundary

AARM is not an audit standard, certification, assurance opinion, legal compliance proof, regulator approval, procurement recommendation, vendor ranking, or score.

## AARM Dimensions

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

| Audit object clarity  |                                                            |
|-----------------------|------------------------------------------------------------|
| DEFINITION            | System can identify lifecycle work under review            |
| AUDIT QUESTION        | Can reviewers name the lifecycle unit, scope, and outcome? |
| EVIDENCE EXPECTED     | Lifecycle ID, workflow scope, object inventory             |
| GAIC/MRO DEPENDENCY   | MRO-01, MRO-03, MRO-04                                     |
| AUTHOR-SYNTHESIS NOTE | AARM synthesis                                             |
| BOUNDARY              | Not legal category                                         |

**Authority traceability**

|                              |                                                      |
|------------------------------|------------------------------------------------------|
| <b>DEFINITION</b>            | Delegated authority can be traced to action          |
| <b>AUDIT QUESTION</b>        | Was the action authorized within scope?              |
| <b>EVIDENCE EXPECTED</b>     | Authority grant, expiry, escalation, reauthorization |
| <b>GAIC/MRO DEPENDENCY</b>   | MRO-02, MRO-07                                       |
| <b>AUTHOR-SYNTHESIS NOTE</b> | AARM synthesis                                       |
| <b>BOUNDARY</b>              | Not legal delegation proof                           |

**Human / agent responsibility mapping**

|                              |                                                               |
|------------------------------|---------------------------------------------------------------|
| <b>DEFINITION</b>            | Human and agent roles are separated and mapped                |
| <b>AUDIT QUESTION</b>        | Who owned intent, execution, review, acceptance, remediation? |
| <b>EVIDENCE EXPECTED</b>     | Role-responsibility map, transfer record                      |
| <b>GAIC/MRO DEPENDENCY</b>   | MRO-01, MRO-03, MRO-06                                        |
| <b>AUTHOR-SYNTHESIS NOTE</b> | AARM synthesis                                                |
| <b>BOUNDARY</b>              | No anthropomorphism                                           |

**Evidence sufficiency**

|                              |                                                     |
|------------------------------|-----------------------------------------------------|
| <b>DEFINITION</b>            | Evidence supports reconstruction under review scope |
| <b>AUDIT QUESTION</b>        | Can reviewers reconstruct without narrative memory? |
| <b>EVIDENCE EXPECTED</b>     | Evidence chain, pointers, integrity records         |
| <b>GAIC/MRO DEPENDENCY</b>   | MRO-04, MRO-08, MRO-16                              |
| <b>AUTHOR-SYNTHESIS NOTE</b> | AARM synthesis                                      |
| <b>BOUNDARY</b>              | Not professional sufficiency conclusion             |

**Agent/tool partitioning**

|                       |                                                                   |
|-----------------------|-------------------------------------------------------------------|
| DEFINITION            | Evidence separates agent, tool, vendor, processor, project        |
| AUDIT QUESTION        | Can reviewers identify which surface produced each evidence item? |
| EVIDENCE EXPECTED     | Agent/tool trace, partition keys, substitution record             |
| GAIC/MRO DEPENDENCY   | MRO-05, MRO-08, MRO-14, MRO-15                                    |
| AUTHOR-SYNTHESIS NOTE | AARM synthesis                                                    |
| BOUNDARY              | No vendor ranking                                                 |

**Exception traceability**

|                       |                                                                |
|-----------------------|----------------------------------------------------------------|
| DEFINITION            | Exceptions link to lifecycle objects and owners                |
| AUDIT QUESTION        | Can deviation be linked to authority, impact, and next action? |
| EVIDENCE EXPECTED     | Exception record, impact scope, escalation                     |
| GAIC/MRO DEPENDENCY   | MRO-07, MRO-16                                                 |
| AUTHOR-SYNTHESIS NOTE | GAIC-derived synthesis                                         |
| BOUNDARY              | Not legal violation finding                                    |

**Accepted outcome evidence**

|                       |                                                             |
|-----------------------|-------------------------------------------------------------|
| DEFINITION            | Outcome states are recorded with role and evidence linkage  |
| AUDIT QUESTION        | Who accepted, rejected, disputed, or remediated the result? |
| EVIDENCE EXPECTED     | Outcome state, reviewer role, evidence link                 |
| GAIC/MRO DEPENDENCY   | MRO-04                                                      |
| AUTHOR-SYNTHESIS NOTE | GAIC-derived synthesis                                      |
| BOUNDARY              | Acceptance is not compliance proof                          |

**Remediation closure**

|                       |                                                                |
|-----------------------|----------------------------------------------------------------|
| DEFINITION            | Corrective action and closure are evidence-backed              |
| AUDIT QUESTION        | What was remediated and who accepted closure?                  |
| EVIDENCE EXPECTED     | Remediation plan, correction evidence, recheck, closure status |
| GAIC/MRO DEPENDENCY   | MRO-16                                                         |
| AUTHOR-SYNTHESIS NOTE | GAIC-derived synthesis                                         |
| BOUNDARY              | Not legal settlement                                           |

**Privacy and selective disclosure**

|                       |                                                      |
|-----------------------|------------------------------------------------------|
| DEFINITION            | Evidence can be minimized and disclosed by scope     |
| AUDIT QUESTION        | Can review happen without exposing more than needed? |
| EVIDENCE EXPECTED     | Redaction, retention, disclosure profile             |
| GAIC/MRO DEPENDENCY   | MRO-10 to MRO-13                                     |
| AUTHOR-SYNTHESIS NOTE | author synthesis in this paper                       |
| BOUNDARY              | Not legal advice                                     |

**Third-party review boundary**

|                       |                                             |
|-----------------------|---------------------------------------------|
| DEFINITION            | Review role and claim boundary are explicit |
| AUDIT QUESTION        | What can reviewers inspect and not claim?   |
| EVIDENCE EXPECTED     | Scope, ruleset, verdict, boundary statement |
| GAIC/MRO DEPENDENCY   | MRO-11, MRO-12, MRO-16                      |
| AUTHOR-SYNTHESIS NOTE | author synthesis in this paper              |
| BOUNDARY              | No certification/opinion                    |

**Readiness Levels**

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

**L0 Unobservable**

|                                          |                                                                    |
|------------------------------------------|--------------------------------------------------------------------|
| <b>DEFINITION</b>                        | Lifecycle work cannot be reconstructed                             |
| <b>OBSERVABLE TRAITS</b>                 | Outputs only, informal memory, no lifecycle ID                     |
| <b>MINIMUM EVIDENCE</b>                  | None or ad hoc notes                                               |
| <b>WHAT IT DOES NOT PROVE</b>            | Absence of risk or non-compliance                                  |
| <b>SOURCE / SYNTHESIS CLASSIFICATION</b> | author synthesis in this paper grounded in audit evidence concepts |

**L1 Log-Visible**

|                                          |                                                       |
|------------------------------------------|-------------------------------------------------------|
| <b>DEFINITION</b>                        | Activity logs exist but are not responsibility-linked |
| <b>OBSERVABLE TRAITS</b>                 | Logs, traces, timestamps, event streams               |
| <b>MINIMUM EVIDENCE</b>                  | Raw logs or trace exports                             |
| <b>WHAT IT DOES NOT PROVE</b>            | Authority, responsibility, accepted outcome, closure  |
| <b>SOURCE / SYNTHESIS CLASSIFICATION</b> | EVID-02/EVID-03 + author synthesis in this paper      |

**L2 Trace-Linked**

|                                          |                                                          |
|------------------------------------------|----------------------------------------------------------|
| <b>DEFINITION</b>                        | Logs/traces connect some actions to workflows or actors  |
| <b>OBSERVABLE TRAITS</b>                 | Workflow IDs, agent/tool traces, partial actor link      |
| <b>MINIMUM EVIDENCE</b>                  | Trace-linked run records                                 |
| <b>WHAT IT DOES NOT PROVE</b>            | Complete lifecycle responsibility or sufficient evidence |
| <b>SOURCE / SYNTHESIS CLASSIFICATION</b> | EVID-01/EVID-02 + author synthesis in this paper         |

### L3 Evidence-Structured

|                                          |                                                                  |
|------------------------------------------|------------------------------------------------------------------|
| <b>DEFINITION</b>                        | Lifecycle objects exist for reconstruction                       |
| <b>OBSERVABLE TRAITS</b>                 | Role maps, authority objects, outcome records, exception records |
| <b>MINIMUM EVIDENCE</b>                  | MRO-aligned evidence objects                                     |
| <b>WHAT IT DOES NOT PROVE</b>            | Assurance readiness, legal compliance, operating effectiveness   |
| <b>SOURCE / SYNTHESIS CLASSIFICATION</b> | GAIC-derived + audit/control grounding                           |

### L4 Auditability-Ready

|                                          |                                                                    |
|------------------------------------------|--------------------------------------------------------------------|
| <b>DEFINITION</b>                        | Defined-scope evidence chains reconstruct lifecycle work           |
| <b>OBSERVABLE TRAITS</b>                 | Walkthroughs, evidence requests, privacy controls, closure records |
| <b>MINIMUM EVIDENCE</b>                  | Complete evidence chain for review scope                           |
| <b>WHAT IT DOES NOT PROVE</b>            | Audit opinion, certification, regulator approval                   |
| <b>SOURCE / SYNTHESIS CLASSIFICATION</b> | author synthesis in this paper                                     |

### L5 Assurance-Ready

|                                          |                                                                                                    |
|------------------------------------------|----------------------------------------------------------------------------------------------------|
| <b>DEFINITION</b>                        | Evidence architecture may support professional assurance planning or review within a defined scope |
| <b>OBSERVABLE TRAITS</b>                 | Repeatable evidence packs, reviewer scope, rulesets, recheck procedures                            |
| <b>MINIMUM EVIDENCE</b>                  | Reviewable evidence packs and boundary statement                                                   |
| <b>WHAT IT DOES NOT PROVE</b>            | Actual assurance, certification, legal compliance, regulator acceptance                            |
| <b>SOURCE / SYNTHESIS CLASSIFICATION</b> | author synthesis in this paper bounded by AUD-03/BOUND-03                                          |

# Minimum Evidence Crosswalk

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

## Audit object clarity

- L0/L1 CONCERN** No stable work unit
- L2/L3 TRANSITION** Work unit IDs appear in traces and evidence objects
- L4/L5 TRANSITION** Work units support scoped walkthroughs and evidence packs

## Authority traceability

- L0/L1 CONCERN** Access logs only
- L2/L3 TRANSITION** Partial authority metadata
- L4/L5 TRANSITION** Delegated authority records are reconstructable

## Responsibility mapping

- L0/L1 CONCERN** User/service IDs only
- L2/L3 TRANSITION** Partial role map
- L4/L5 TRANSITION** Human/agent responsibility chain is explicit

## Evidence sufficiency

- L0/L1 CONCERN** Narrative memory
- L2/L3 TRANSITION** Evidence pointers emerge
- L4/L5 TRANSITION** Evidence chain supports review objective

## Partitioning

- L0/L1 CONCERN** Flattened logs
- L2/L3 TRANSITION** Partial partitions
- L4/L5 TRANSITION** Agent/tool/vendor/project/privacy partitions are linked

## Exception traceability

- L0/L1 CONCERN** Incident tickets only
- L2/L3 TRANSITION** Exceptions linked to work units
- L4/L5 TRANSITION** Exceptions include owner, action, recheck, closure

**Accepted outcome**

- L0/L1 CONCERN**      Output completion only
- L2/L3 TRANSITION**      Partial review records
- L4/L5 TRANSITION**      Accepted/rejected/disputed/remediated states are evidence-backed

**Remediation closure**

- L0/L1 CONCERN**      Closed ticket only
- L2/L3 TRANSITION**      Corrective action tracked
- L4/L5 TRANSITION**      Closure is accepted and reopen criteria exist

**Privacy/selective disclosure**

- L0/L1 CONCERN**      Collect everything or disclose ad hoc
- L2/L3 TRANSITION**      Redaction/retention rules appear
- L4/L5 TRANSITION**      Selective disclosure packages support review

**Third-party boundary**

- L0/L1 CONCERN**      No scope statement
- L2/L3 TRANSITION**      Limited reviewer scope
- L4/L5 TRANSITION**      Ruleset, scope, verdict, and non-certifying boundary are explicit

# Appendix E – Exception / Remediation Closure Checklist

**Purpose:** Define closure evidence requirements for exceptions, disputes, and remediation. **Intended reader:** Internal audit, incident governance, compliance, risk, security, AI platform owners, remediation owners.

## Source Grounding

Incident/governance context: EVID-04, AI-02. Controls and audit practice: AUD-04, AUD-05, AUD-06. GAIC dependencies: MRO-04, MRO-07, MRO-16.

## Boundary

This checklist frames closure as evidence-backed governance state. It does not prove legal remedy, settlement, regulator acceptance, operational effectiveness, certification, or assurance.

## Closure Checklist

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

| Authority exception |                                                                           |
|---------------------|---------------------------------------------------------------------------|
| EVIDENCE REQUIRED   | Baseline authority, observed action, trigger, work unit, evidence pointer |
| OWNER               | Escalation owner                                                          |
| CORRECTIVE ACTION   | Stop, downgrade, reauthorize, confirm, or constrain                       |
| CLOSURE STATUS      | Open / remediated / closed                                                |
| REVIEWER            | Authority reviewer                                                        |
| REOPEN CRITERIA     | Repeat drift, scope expansion, failed confirmation                        |
| BOUNDARY NOTE       | Not legal breach finding                                                  |

**Responsibility gap**

|                          |                                                                     |
|--------------------------|---------------------------------------------------------------------|
| <b>EVIDENCE REQUIRED</b> | Missing role owner, affected lifecycle stage, agent role, work unit |
| <b>OWNER</b>             | Governance owner                                                    |
| <b>CORRECTIVE ACTION</b> | Assign owner, update role map, recheck workflow                     |
| <b>CLOSURE STATUS</b>    | Open / assigned / closed                                            |
| <b>REVIEWER</b>          | Governance reviewer                                                 |
| <b>REOPEN CRITERIA</b>   | Owner missing in later run                                          |
| <b>BOUNDARY NOTE</b>     | Not legal liability assignment                                      |

**Tool-action side effect**

|                          |                                                                   |
|--------------------------|-------------------------------------------------------------------|
| <b>EVIDENCE REQUIRED</b> | Tool call, affected system/data, reversibility, rollback evidence |
| <b>OWNER</b>             | Tool-action owner                                                 |
| <b>CORRECTIVE ACTION</b> | Rollback, correction, disable tool, add approval                  |
| <b>CLOSURE STATUS</b>    | Open / corrected / closed                                         |
| <b>REVIEWER</b>          | System owner                                                      |
| <b>REOPEN CRITERIA</b>   | New side effect or rollback failure                               |
| <b>BOUNDARY NOTE</b>     | Not liability conclusion                                          |

**Outcome dispute**

|                          |                                                    |
|--------------------------|----------------------------------------------------|
| <b>EVIDENCE REQUIRED</b> | Output, criteria, dispute reason, reviewer notes   |
| <b>OWNER</b>             | Review owner                                       |
| <b>CORRECTIVE ACTION</b> | Re-review, correct, reject, accept with limitation |
| <b>CLOSURE STATUS</b>    | Open / disputed / remediated / accepted            |
| <b>REVIEWER</b>          | Outcome reviewer                                   |
| <b>REOPEN CRITERIA</b>   | New evidence or recurring dispute                  |
| <b>BOUNDARY NOTE</b>     | Not settlement proof                               |

**Privacy event**

|                          |                                                            |
|--------------------------|------------------------------------------------------------|
| <b>EVIDENCE REQUIRED</b> | Data category, disclosure path, redaction/retention record |
| <b>OWNER</b>             | Privacy owner                                              |
| <b>CORRECTIVE ACTION</b> | Redact, restrict, delete, hold, notify internal process    |
| <b>CLOSURE STATUS</b>    | Open / contained / closed                                  |
| <b>REVIEWER</b>          | Privacy reviewer                                           |
| <b>REOPEN CRITERIA</b>   | Unauthorized exposure or rights request                    |
| <b>BOUNDARY NOTE</b>     | Not legal advice                                           |

**Handoff failure**

|                          |                                                                          |
|--------------------------|--------------------------------------------------------------------------|
| <b>EVIDENCE REQUIRED</b> | Source/target agent, transferred scope, retained scope, rejection reason |
| <b>OWNER</b>             | Process owner                                                            |
| <b>CORRECTIVE ACTION</b> | Return, reassign, clarify constraints                                    |
| <b>CLOSURE STATUS</b>    | Open / reassigned / closed                                               |
| <b>REVIEWER</b>          | Process reviewer                                                         |
| <b>REOPEN CRITERIA</b>   | Repeated failed handoff                                                  |
| <b>BOUNDARY NOTE</b>     | Not legal transfer                                                       |

**Cross-project reuse issue**

|                          |                                                               |
|--------------------------|---------------------------------------------------------------|
| <b>EVIDENCE REQUIRED</b> | Source context, target context, reuse decision, missing reset |
| <b>OWNER</b>             | Reuse owner                                                   |
| <b>CORRECTIVE ACTION</b> | Revalidate, reset, withdraw, update evidence                  |
| <b>CLOSURE STATUS</b>    | Open / revalidated / closed                                   |
| <b>REVIEWER</b>          | Reuse/governance reviewer                                     |
| <b>REOPEN CRITERIA</b>   | Reuse in new context                                          |
| <b>BOUNDARY NOTE</b>     | Not legal reuse clearance                                     |

**Substitution evidence gap**

|                          |                                                                        |
|--------------------------|------------------------------------------------------------------------|
| <b>EVIDENCE REQUIRED</b> | Prior/new component, evidence break, regression or revalidation status |
| <b>OWNER</b>             | Change owner                                                           |
| <b>CORRECTIVE ACTION</b> | Reauthorize, retest, revalidate evidence chain                         |
| <b>CLOSURE STATUS</b>    | Open / revalidated / closed                                            |
| <b>REVIEWER</b>          | Change reviewer                                                        |
| <b>REOPEN CRITERIA</b>   | Component change or failed evidence continuity                         |
| <b>BOUNDARY NOTE</b>     | Not conformance certification                                          |

**Remediation defect**

|                          |                                                     |
|--------------------------|-----------------------------------------------------|
| <b>EVIDENCE REQUIRED</b> | Prior remediation, failed recheck, affected outcome |
| <b>OWNER</b>             | Remediation owner                                   |
| <b>CORRECTIVE ACTION</b> | Reopen, assign corrective action, recheck           |
| <b>CLOSURE STATUS</b>    | Reopened / corrected / closed                       |
| <b>REVIEWER</b>          | Closure reviewer                                    |
| <b>REOPEN CRITERIA</b>   | Failed recheck or recurring exception               |
| <b>BOUNDARY NOTE</b>     | Not regulator closure                               |

**Closure Evidence Fields**

- Event ID.
- Work unit ID.
- Affected MRO.
- Affected outcome.
- Owner.
- Evidence pointer.
- Corrective action.
- Recheck evidence.
- Closure reviewer.
- Closure status.
- Reopen criteria.
- Privacy treatment.
- Boundary statement.

---

## Guide Relationship

Guide 1 should translate closure fields into workflow, ticketing, runtime, and evidence-export architecture. Guide 2 should translate closure fields into incident/dispute/remediation governance routines.

---

# Appendix F — Boundary Language

**Purpose:** Provide safe language for auditability, assurance, validation, MRO, AARM, MPLP, privacy, and professional-use boundaries. **Intended reader:** Author, editors, reviewers, publication QA, Guide 1/2 teams, source QA.

## Source Grounding

Assurance and attestation boundary: AUD-03, BOUND-03. Conformity/certification boundary: BOUND-01, BOUND-02. GAIC dependencies: Validation Lab boundary, MPLP boundary, companion-paper boundary, R4A/R4B/R4C boundary QA.

## Allowed / Forbidden / Replacement Language

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

| Paper status      |                                                                                                              |
|-------------------|--------------------------------------------------------------------------------------------------------------|
| ALLOWED WORDING   | "publication-candidate preparation source"; "internal candidate package"; "source-grounded candidate source" |
| FORBIDDEN WORDING | "published"; "final"; "sealed"; "live public release"                                                        |
| SAFE REPLACEMENT  | "R4 publication-candidate preparation package, not public/final/sealed/live"                                 |

| Auditability      |                                                               |
|-------------------|---------------------------------------------------------------|
| ALLOWED WORDING   | "proposed auditability framework"; "lifecycle evidence guide" |
| FORBIDDEN WORDING | "audit standard"; "official audit methodology"                |
| SAFE REPLACEMENT  | "auditability readiness framework"                            |

Evidence

|                   |                                                 |
|-------------------|-------------------------------------------------|
| ALLOWED WORDING   | "evidence chain supports reconstruction"        |
| FORBIDDEN WORDING | "evidence chain proves compliance"              |
| SAFE REPLACEMENT  | "evidence chain supports responsibility review" |

Agentic Audit Object

|                   |                                       |
|-------------------|---------------------------------------|
| ALLOWED WORDING   | "proposed object model in this paper" |
| FORBIDDEN WORDING | "required schema"; "legal category"   |
| SAFE REPLACEMENT  | "conceptual review object"            |

MRO

|                   |                                                  |
|-------------------|--------------------------------------------------|
| ALLOWED WORDING   | "GAIC-derived governance object"                 |
| FORBIDDEN WORDING | "legal requirement"; "regulator-mandated object" |
| SAFE REPLACEMENT  | "proposed lifecycle responsibility object"       |

AARM

|                   |                                                            |
|-------------------|------------------------------------------------------------|
| ALLOWED WORDING   | "proposed readiness model"                                 |
| FORBIDDEN WORDING | "certification level"; "maturity score"; "assurance grade" |
| SAFE REPLACEMENT  | "readiness vocabulary"                                     |

Assurance

|                   |                                                                |
|-------------------|----------------------------------------------------------------|
| ALLOWED WORDING   | "may support assurance planning subject to professional scope" |
| FORBIDDEN WORDING | "assures"; "guarantees assurance"; "issues assurance opinion"  |
| SAFE REPLACEMENT  | "may support readiness discussion"                             |

**Certification**

|                          |                                                         |
|--------------------------|---------------------------------------------------------|
| <b>ALLOWED WORDING</b>   | "non-certifying review"; "not certification"            |
| <b>FORBIDDEN WORDING</b> | "certified"; "certifies systems"; "certification-ready" |
| <b>SAFE REPLACEMENT</b>  | "reviewable under defined scope"                        |

**Legal/privacy**

|                          |                                                            |
|--------------------------|------------------------------------------------------------|
| <b>ALLOWED WORDING</b>   | "privacy/legal context"; "requires legal review"           |
| <b>FORBIDDEN WORDING</b> | "GDPR compliant"; "legal compliance proof"; "legal advice" |
| <b>SAFE REPLACEMENT</b>  | "privacy-aware evidence architecture"                      |

**Big Four**

|                          |                                                                   |
|--------------------------|-------------------------------------------------------------------|
| <b>ALLOWED WORDING</b>   | "reviewed sources create market context"                          |
| <b>FORBIDDEN WORDING</b> | "Big Four endorse this paper"; "Big Four need this paper"         |
| <b>SAFE REPLACEMENT</b>  | "Big Four sources discuss related AI assurance/governance topics" |

**Audit/professional bodies**

|                          |                                                   |
|--------------------------|---------------------------------------------------|
| <b>ALLOWED WORDING</b>   | "supports terminology or boundary"                |
| <b>FORBIDDEN WORDING</b> | "endorses this paper"; "adopts AARM"              |
| <b>SAFE REPLACEMENT</b>  | "provides external professional-language context" |

**Validation Lab**

|                          |                                                                      |
|--------------------------|----------------------------------------------------------------------|
| <b>ALLOWED WORDING</b>   | "non-certifying evidence adjudication example"                       |
| <b>FORBIDDEN WORDING</b> | "certification authority"; "proves compliance"; "regulator-approved" |
| <b>SAFE REPLACEMENT</b>  | "one optional example of evidence review boundary"                   |

**MPLP**

|                          |                                                      |
|--------------------------|------------------------------------------------------|
| <b>ALLOWED WORDING</b>   | "one optional protocol path"                         |
| <b>FORBIDDEN WORDING</b> | "required"; "industry standard"; "only correct path" |
| <b>SAFE REPLACEMENT</b>  | "one possible implementation path"                   |

**Cognitive OS**

|                          |                                                                                |
|--------------------------|--------------------------------------------------------------------------------|
| <b>ALLOWED WORDING</b>   | "possible runtime path for evidence capture in related product-stack planning" |
| <b>FORBIDDEN WORDING</b> | "required for auditability"                                                    |
| <b>SAFE REPLACEMENT</b>  | "one possible implementation environment, outside this paper requirement"      |

**SoloCrew**

|                          |                                                                      |
|--------------------------|----------------------------------------------------------------------|
| <b>ALLOWED WORDING</b>   | "product proof path for workflow patterns in planning notes"         |
| <b>FORBIDDEN WORDING</b> | "proves enterprise readiness"                                        |
| <b>SAFE REPLACEMENT</b>  | "illustrative product proof path, not an enterprise-readiness claim" |

**SEO/GEO**

|                          |                                                            |
|--------------------------|------------------------------------------------------------|
| <b>ALLOWED WORDING</b>   | "future planning note"                                     |
| <b>FORBIDDEN WORDING</b> | "uplift achieved"; "answer-engine recognized"              |
| <b>SAFE REPLACEMENT</b>  | "future publication/visibility work after artifact exists" |

**Procurement**

|                          |                                             |
|--------------------------|---------------------------------------------|
| <b>ALLOWED WORDING</b>   | "readiness conversation input"              |
| <b>FORBIDDEN WORDING</b> | "recommended vendor"; "procurement ranking" |
| <b>SAFE REPLACEMENT</b>  | "non-procurement framework"                 |

## Later insurability white paper

**ALLOWED WORDING** "future insurability paper"

**FORBIDDEN WORDING** "insurance guarantee"; "underwriting standard"

**SAFE REPLACEMENT** "future source research required"

## Big Four Boundary

Use:

- "The reviewed Big Four sources discuss AI assurance, trusted AI, governance, risk, controls, audit transformation, and agentic AI adoption."
- "These sources create market context for this paper."

Do not use:

- "Big Four firms endorse this paper."
- "The Big Four need this paper."
- "this paper replaces Big Four methodology."

## Audit / Professional Body Boundary

Use professional sources for terminology and boundary. Do not imply that PCAOB, IAASB, AICPA, ISACA, IIA, ISO, NIST, OECD, EU institutions, Singapore agencies, ICO, or any other organization endorses or adopts this paper.

## Legal / Privacy Boundary

Use privacy and legal sources as context. State that jurisdiction-specific interpretation belongs to qualified legal/privacy professionals. Do not say any evidence object in this paper proves GDPR, EU AI Act, UK GDPR, or other legal compliance.

## Validation Lab Boundary

Validation Lab may be described only as one non-certifying evidence adjudication example. It does not certify systems, issue compliance certificates, provide legal opinions, issue assurance opinions, act as regulator, act as conformity assessment body, prove enterprise readiness, or guarantee insurability.

## MPLP Boundary

MPLP may be discussed only as one optional protocol path. This paper does not require MPLP, does not claim MPLP is industry-standard, and does not claim MPLP proves compliance, auditability, assurance, or enterprise readiness.

---

## AARM Boundary

AARM is a proposed readiness model. It does not score vendors, certify systems, issue assurance, prove legal compliance, bind regulators, or provide procurement guidance. L5 "Assurance-Ready" means evidence architecture may support qualified assurance planning or review within a defined scope, subject to professional scope and judgment. It does not mean assurance has been issued.

---

---

# Source Note and Citation Register

R4 uses source-ID citation notes rather than public footnotes. Source IDs appear in chapter source notes and table source-support columns. The candidate package keeps the full source register and chapter-level citation map as internal package records.

---

## R4 Citation Style Decision

- Internal style for R4: source IDs in text and tables, plus a package-level source register.
  - Public style for a later wave: decide whether to convert source IDs into endnotes, footnotes, or an appendix source register.
  - Quote guidance: paraphrase by default; avoid direct quotation unless a later citation QA wave approves a short, necessary excerpt.
  - Boundary guidance: source presence does not imply endorsement, standard-setting, certification, assurance, legal advice, procurement approval, or regulator approval.
- 

## Source Role Summary

- BF sources: market context only.
- AUD sources: audit/assurance terminology and boundary context.
- AI sources: governance context, not legal compliance proof.
- EVID sources: observability, provenance, logs, and evidence-chain distinction.
- PRIV sources: minimization, retention, selective disclosure, and privacy/evidence tension.
- BOUND sources: validation, verification, attestation, certification, and conformity-assessment boundaries.
- GAIC-SOURCE: MROs, RCCS-M/ALCS context, Validation Lab boundary, and companion-paper sequencing.

---

# Artifact Integrity Note

R4 generated an internal candidate source package and checksum file. Candidate HTML/PDF/DOCX artifacts were intentionally deferred because the existing renderer scripts are GAIC public-white-paper specific and direct reuse would risk touching public/sealed GAIC artifact assumptions. Artifact generation is documented in the internal generation plan and the R4 artifact-generation decision report.

---

---

## Source Note and Citation Register

The source register and citation map below are part of the internal candidate package. Source IDs are used as internal citation handles for review and future public citation styling.

# AIAAWP R4 Source Register

**Status:** R4 publication-candidate preparation source register **Boundary:** Sources listed here are used for context, terminology, grounding, or boundary discipline. Inclusion does not imply endorsement of this paper, GAIC, MPLP, AARM, Validation Lab, or Jearon Wong.

---

## Source Use Rules

- Big Four sources are market context only.
- Audit/professional sources support audit evidence, assurance, controls, internal audit, and boundary language.
- AI governance/public guidance sources support governance context, risk management, accountability, documentation, monitoring, and human oversight.
- Evidence/provenance/observability/logging sources support the distinction between logs/traces and responsibility-linked audit evidence chains.
- Privacy sources support minimization, retention, selective disclosure, and evidence/privacy tension.
- Boundary sources support distinctions among validation, assurance, attestation, certification, conformity assessment, and audit opinion.
- GAIC source truth supports MROs, RCCS-M/ALCS context, Evidence-Based Validation Pattern, Validation Lab boundary, and companion-paper boundary.

# Register

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

|                              |                                                      |
|------------------------------|------------------------------------------------------|
| <b>BF-01</b>                 |                                                      |
| <b>TITLE</b>                 | AI agents: Scaling faster than expected              |
| <b>ORGANIZATION</b>          | Deloitte                                             |
| <b>SOURCE CLASS</b>          | Big Four market context                              |
| <b>USE ROLE</b>              | Agentic AI adoption and enterprise readiness context |
| <b>CHAPTERS USED</b>         | 0, 2, 14, 15                                         |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase only                                      |
| <b>BOUNDARY NOTE</b>         | No endorsement claim.                                |

|                              |                                             |
|------------------------------|---------------------------------------------|
| <b>BF-02</b>                 |                                             |
| <b>TITLE</b>                 | Assurance for AI                            |
| <b>ORGANIZATION</b>          | PwC                                         |
| <b>SOURCE CLASS</b>          | Big Four market context                     |
| <b>USE ROLE</b>              | AI assurance market context                 |
| <b>CHAPTERS USED</b>         | 0, 1, 14, 15                                |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase only                             |
| <b>BOUNDARY NOTE</b>         | Context only; not validation of this paper. |

|                              |                                                                                      |
|------------------------------|--------------------------------------------------------------------------------------|
| <b>BF-03</b>                 |                                                                                      |
| <b>TITLE</b>                 | EY Assurance releases new technology capabilities strengthening confidence and trust |
| <b>ORGANIZATION</b>          | EY                                                                                   |
| <b>SOURCE CLASS</b>          | Big Four market context                                                              |
| <b>USE ROLE</b>              | Audit transformation and assurance technology context                                |
| <b>CHAPTERS USED</b>         | 1, 15                                                                                |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase only                                                                      |
| <b>BOUNDARY NOTE</b>         | No methodology replacement claim.                                                    |

**BF-04**

|                              |                                                             |
|------------------------------|-------------------------------------------------------------|
| <b>TITLE</b>                 | KPMG Trusted AI                                             |
| <b>ORGANIZATION</b>          | KPMG                                                        |
| <b>SOURCE CLASS</b>          | Big Four market context                                     |
| <b>USE ROLE</b>              | Trusted AI, governance, privacy, and accountability context |
| <b>CHAPTERS USED</b>         | 2, 10, 14                                                   |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase only                                             |
| <b>BOUNDARY NOTE</b>         | Context only.                                               |

**BF-05**

|                              |                                                                                                          |
|------------------------------|----------------------------------------------------------------------------------------------------------|
| <b>TITLE</b>                 | AI Agents Move Beyond Experimentation as Leaders Prepare for Competitive Transformation Within 24 Months |
| <b>ORGANIZATION</b>          | KPMG                                                                                                     |
| <b>SOURCE CLASS</b>          | Big Four market context                                                                                  |
| <b>USE ROLE</b>              | Agentic AI adoption and risk context                                                                     |
| <b>CHAPTERS USED</b>         | 0, 2, 14, 15                                                                                             |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase only                                                                                          |
| <b>BOUNDARY NOTE</b>         | No adoption or endorsement claim.                                                                        |

**AUD-01**

|                              |                                                                    |
|------------------------------|--------------------------------------------------------------------|
| <b>TITLE</b>                 | AS 1105: Audit Evidence                                            |
| <b>ORGANIZATION</b>          | PCAOB                                                              |
| <b>SOURCE CLASS</b>          | Audit/professional body                                            |
| <b>USE ROLE</b>              | Audit evidence terminology and sufficiency/appropriateness framing |
| <b>CHAPTERS USED</b>         | 0, 3, 4, 5, 6, 7, 13, 16                                           |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase; short definition quote only if later approved          |
| <b>BOUNDARY NOTE</b>         | This paper is not a PCAOB audit standard.                          |

**AUD-02**

|                         |                                                        |
|-------------------------|--------------------------------------------------------|
| <b>TITLE</b>            | ISA 500 (Revised), Audit Evidence                      |
| <b>ORGANIZATION</b>     | IAASB                                                  |
| <b>SOURCE CLASS</b>     | Audit/professional body                                |
| <b>USE ROLE</b>         | Audit evidence and technology-enabled evidence context |
| <b>CHAPTERS USED</b>    | 3, 4, 7, 13                                            |
| <b>QUOTE/PARAPHRASE</b> | Paraphrase only                                        |
| <b>NOTE</b>             |                                                        |
| <b>BOUNDARY NOTE</b>    | This paper is not ISA methodology.                     |

**AUD-03**

|                         |                                                  |
|-------------------------|--------------------------------------------------|
| <b>TITLE</b>            | ISAE 3000 (Revised)                              |
| <b>ORGANIZATION</b>     | IAASB                                            |
| <b>SOURCE CLASS</b>     | Audit/professional body                          |
| <b>USE ROLE</b>         | Assurance engagement boundary                    |
| <b>CHAPTERS USED</b>    | 1, 12, 13, 15, 16                                |
| <b>QUOTE/PARAPHRASE</b> | Paraphrase only                                  |
| <b>NOTE</b>             |                                                  |
| <b>BOUNDARY NOTE</b>    | This paper does not issue assurance conclusions. |

**AUD-04**

|                         |                                                                      |
|-------------------------|----------------------------------------------------------------------|
| <b>TITLE</b>            | Trust Services Criteria                                              |
| <b>ORGANIZATION</b>     | AICPA                                                                |
| <b>SOURCE CLASS</b>     | Audit/professional body                                              |
| <b>USE ROLE</b>         | Control vocabulary, monitoring, privacy/security, control activities |
| <b>CHAPTERS USED</b>    | 5, 7, 8, 10, 11, 14                                                  |
| <b>QUOTE/PARAPHRASE</b> | Paraphrase only                                                      |
| <b>NOTE</b>             |                                                                      |
| <b>BOUNDARY NOTE</b>    | This paper is not SOC or attestation guidance.                       |

**AUD-05**

|                              |                                                                                                   |
|------------------------------|---------------------------------------------------------------------------------------------------|
| <b>TITLE</b>                 | Audit and Assurance Guidance for the NIST Cybersecurity Framework 2.0 and Artificial Intelligence |
| <b>ORGANIZATION</b>          | ISACA                                                                                             |
| <b>SOURCE CLASS</b>          | Audit/professional body                                                                           |
| <b>USE ROLE</b>              | AI audit and assurance practice context                                                           |
| <b>CHAPTERS USED</b>         | 6, 7, 8, 11, 13, 15                                                                               |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase only                                                                                   |
| <b>BOUNDARY NOTE</b>         | No ISACA endorsement.                                                                             |

**AUD-06**

|                              |                                                                 |
|------------------------------|-----------------------------------------------------------------|
| <b>TITLE</b>                 | Global Technology Audit Guide: Auditing Artificial Intelligence |
| <b>ORGANIZATION</b>          | IIA                                                             |
| <b>SOURCE CLASS</b>          | Audit/professional body                                         |
| <b>USE ROLE</b>              | Internal audit, AI governance, controls, human oversight        |
| <b>CHAPTERS USED</b>         | 2, 7, 8, 11, 13, 14, 15                                         |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase only                                                 |
| <b>BOUNDARY NOTE</b>         | Internal audit language only.                                   |

**AUD-07**

|                              |                             |
|------------------------------|-----------------------------|
| <b>TITLE</b>                 | The IIA's Three Lines Model |
| <b>ORGANIZATION</b>          | IIA                         |
| <b>SOURCE CLASS</b>          | Audit/professional body     |
| <b>USE ROLE</b>              | Governance role separation  |
| <b>CHAPTERS USED</b>         | 1, 14, 15                   |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase only             |
| <b>BOUNDARY NOTE</b>         | Not an enterprise mandate.  |

**AI-01**

|                         |                                                            |
|-------------------------|------------------------------------------------------------|
| <b>TITLE</b>            | Artificial Intelligence Risk Management Framework 1.0      |
| <b>ORGANIZATION</b>     | NIST                                                       |
| <b>SOURCE CLASS</b>     | Regulator/public guidance                                  |
| <b>USE ROLE</b>         | AI risk management, governance, monitoring, accountability |
| <b>CHAPTERS USED</b>    | 0, 2, 3, 5, 6, 13, 14, 16                                  |
| <b>QUOTE/PARAPHRASE</b> | Paraphrase only                                            |
| <b>NOTE</b>             |                                                            |
| <b>BOUNDARY NOTE</b>    | Guidance context, not legal proof.                         |

**AI-02**

|                         |                                                          |
|-------------------------|----------------------------------------------------------|
| <b>TITLE</b>            | NIST AI 600-1 Generative AI Profile                      |
| <b>ORGANIZATION</b>     | NIST                                                     |
| <b>SOURCE CLASS</b>     | Regulator/public guidance                                |
| <b>USE ROLE</b>         | Generative AI risk, incidents, monitoring, documentation |
| <b>CHAPTERS USED</b>    | 2, 7, 11, 14                                             |
| <b>QUOTE/PARAPHRASE</b> | Paraphrase only                                          |
| <b>NOTE</b>             |                                                          |
| <b>BOUNDARY NOTE</b>    | Guidance context only.                                   |

**AI-03**

|                         |                                                        |
|-------------------------|--------------------------------------------------------|
| <b>TITLE</b>            | OECD AI Principles                                     |
| <b>ORGANIZATION</b>     | OECD                                                   |
| <b>SOURCE CLASS</b>     | Public guidance                                        |
| <b>USE ROLE</b>         | Responsible AI accountability and transparency context |
| <b>CHAPTERS USED</b>    | 1, 2, 14                                               |
| <b>QUOTE/PARAPHRASE</b> | Paraphrase only                                        |
| <b>NOTE</b>             |                                                        |
| <b>BOUNDARY NOTE</b>    | Policy context only.                                   |

**AI-04**

|                              |                                         |
|------------------------------|-----------------------------------------|
| <b>TITLE</b>                 | ISO/IEC 42001:2023 AI management system |
| <b>ORGANIZATION</b>          | ISO                                     |
| <b>SOURCE CLASS</b>          | Standards body                          |
| <b>USE ROLE</b>              | AI management-system context            |
| <b>CHAPTERS USED</b>         | 2, 13, 14                               |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase public summary only          |
| <b>BOUNDARY NOTE</b>         | No compliance claim.                    |

**AI-05**

|                              |                                       |
|------------------------------|---------------------------------------|
| <b>TITLE</b>                 | ISO/IEC 23894:2023 AI risk management |
| <b>ORGANIZATION</b>          | ISO                                   |
| <b>SOURCE CLASS</b>          | Standards body                        |
| <b>USE ROLE</b>              | AI risk-management context            |
| <b>CHAPTERS USED</b>         | 2, 13, 14                             |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase public summary only        |
| <b>BOUNDARY NOTE</b>         | No standard compliance claim.         |

**AI-06**

|                              |                                                                |
|------------------------------|----------------------------------------------------------------|
| <b>TITLE</b>                 | Regulation (EU) 2024/1689 Artificial Intelligence Act          |
| <b>ORGANIZATION</b>          | European Union                                                 |
| <b>SOURCE CLASS</b>          | Regulator/public guidance                                      |
| <b>USE ROLE</b>              | AI governance, logging, human oversight, documentation context |
| <b>CHAPTERS USED</b>         | 1, 2, 4, 7, 14                                                 |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase only                                                |
| <b>BOUNDARY NOTE</b>         | Legal context only; no legal advice.                           |

**AI-07**

|                              |                                                 |
|------------------------------|-------------------------------------------------|
| <b>TITLE</b>                 | Model AI Governance Framework for Generative AI |
| <b>ORGANIZATION</b>          | AI Verify Foundation / IMDA                     |
| <b>SOURCE CLASS</b>          | Public guidance                                 |
| <b>USE ROLE</b>              | Generative AI governance and testing context    |
| <b>CHAPTERS USED</b>         | 1, 2, 7, 14                                     |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase only                                 |
| <b>BOUNDARY NOTE</b>         | No endorsement.                                 |

**AI-08**

|                              |                                                             |
|------------------------------|-------------------------------------------------------------|
| <b>TITLE</b>                 | Model AI Governance Framework for Agentic AI                |
| <b>ORGANIZATION</b>          | Singapore MDDI / AI Verify Foundation                       |
| <b>SOURCE CLASS</b>          | Public guidance                                             |
| <b>USE ROLE</b>              | Agentic AI governance, autonomy, accountability, safeguards |
| <b>CHAPTERS USED</b>         | 0, 2, 3, 5, 6, 14                                           |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase only                                             |
| <b>BOUNDARY NOTE</b>         | Context only.                                               |

**AI-09**

|                              |                                    |
|------------------------------|------------------------------------|
| <b>TITLE</b>                 | Guidance on AI and Data Protection |
| <b>ORGANIZATION</b>          | UK ICO                             |
| <b>SOURCE CLASS</b>          | Regulator/public guidance          |
| <b>USE ROLE</b>              | AI/privacy context                 |
| <b>CHAPTERS USED</b>         | 9, 10, 14                          |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase only                    |
| <b>BOUNDARY NOTE</b>         | No legal advice.                   |

**EVID-01**

|                              |                                                                                    |
|------------------------------|------------------------------------------------------------------------------------|
| <b>TITLE</b>                 | PROV Overview                                                                      |
| <b>ORGANIZATION</b>          | W3C                                                                                |
| <b>SOURCE CLASS</b>          | Technical documentation                                                            |
| <b>USE ROLE</b>              | Provenance vocabulary for entities, activities, agents, derivation, responsibility |
| <b>CHAPTERS USED</b>         | 3, 4, 5, 8, 9, 13                                                                  |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase only                                                                    |
| <b>BOUNDARY NOTE</b>         | Technical vocabulary, not audit standard.                                          |

**EVID-02**

|                              |                                              |
|------------------------------|----------------------------------------------|
| <b>TITLE</b>                 | OpenTelemetry Observability Primer           |
| <b>ORGANIZATION</b>          | OpenTelemetry                                |
| <b>SOURCE CLASS</b>          | Technical documentation                      |
| <b>USE ROLE</b>              | Logs, metrics, traces, observability context |
| <b>CHAPTERS USED</b>         | 0, 4, 8, 9, 14                               |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase only                              |
| <b>BOUNDARY NOTE</b>         | Do not criticize vendors.                    |

**EVID-03**

|                              |                                                          |
|------------------------------|----------------------------------------------------------|
| <b>TITLE</b>                 | SP 800-92 Guide to Computer Security Log Management      |
| <b>ORGANIZATION</b>          | NIST                                                     |
| <b>SOURCE CLASS</b>          | Public guidance                                          |
| <b>USE ROLE</b>              | Log management context                                   |
| <b>CHAPTERS USED</b>         | 4, 9, 10                                                 |
| <b>QUOTE/PARAPHRASE NOTE</b> | Paraphrase only                                          |
| <b>BOUNDARY NOTE</b>         | Older official logging source; not agentic auditability. |

**EVID-04**

|                         |                                                             |
|-------------------------|-------------------------------------------------------------|
| <b>TITLE</b>            | Cybersecurity Framework 2.0                                 |
| <b>ORGANIZATION</b>     | NIST                                                        |
| <b>SOURCE CLASS</b>     | Public guidance                                             |
| <b>USE ROLE</b>         | Governance, detection, response, recovery, incident context |
| <b>CHAPTERS USED</b>    | 4, 11, 14                                                   |
| <b>QUOTE/PARAPHRASE</b> | Paraphrase only                                             |
| <b>NOTE</b>             |                                                             |
| <b>BOUNDARY NOTE</b>    | Cybersecurity context only.                                 |

**PRIV-01**

|                         |                                                          |
|-------------------------|----------------------------------------------------------|
| <b>TITLE</b>            | GDPR Article 5                                           |
| <b>ORGANIZATION</b>     | European Union                                           |
| <b>SOURCE CLASS</b>     | Regulator/public guidance                                |
| <b>USE ROLE</b>         | Minimization, storage limitation, accountability context |
| <b>CHAPTERS USED</b>    | 9, 10                                                    |
| <b>QUOTE/PARAPHRASE</b> | Paraphrase; short quote only if later approved           |
| <b>NOTE</b>             |                                                          |
| <b>BOUNDARY NOTE</b>    | Legal context, not advice.                               |

**PRIV-02**

|                         |                                                            |
|-------------------------|------------------------------------------------------------|
| <b>TITLE</b>            | EDPB Guidelines 4/2019 on Article 25                       |
| <b>ORGANIZATION</b>     | EDPB                                                       |
| <b>SOURCE CLASS</b>     | Regulator/public guidance                                  |
| <b>USE ROLE</b>         | Data protection by design/default and minimization context |
| <b>CHAPTERS USED</b>    | 9, 10                                                      |
| <b>QUOTE/PARAPHRASE</b> | Paraphrase only                                            |
| <b>NOTE</b>             |                                                            |
| <b>BOUNDARY NOTE</b>    | No GDPR proof.                                             |

**PRIV-03**

|                         |                                                |
|-------------------------|------------------------------------------------|
| <b>TITLE</b>            | NIST Privacy Framework 1.0                     |
| <b>ORGANIZATION</b>     | NIST                                           |
| <b>SOURCE CLASS</b>     | Public guidance                                |
| <b>USE ROLE</b>         | Privacy risk management and governance context |
| <b>CHAPTERS USED</b>    | 10, 14                                         |
| <b>QUOTE/PARAPHRASE</b> | Paraphrase only                                |
| <b>NOTE</b>             |                                                |
| <b>BOUNDARY NOTE</b>    | Guidance only.                                 |

**PRIV-04**

|                         |                           |
|-------------------------|---------------------------|
| <b>TITLE</b>            | Data minimisation         |
| <b>ORGANIZATION</b>     | UK ICO                    |
| <b>SOURCE CLASS</b>     | Regulator/public guidance |
| <b>USE ROLE</b>         | Data minimization context |
| <b>CHAPTERS USED</b>    | 10                        |
| <b>QUOTE/PARAPHRASE</b> | Paraphrase only           |
| <b>NOTE</b>             |                           |
| <b>BOUNDARY NOTE</b>    | No legal advice.          |

**PRIV-05**

|                         |                                          |
|-------------------------|------------------------------------------|
| <b>TITLE</b>            | Storage limitation                       |
| <b>ORGANIZATION</b>     | UK ICO                                   |
| <b>SOURCE CLASS</b>     | Regulator/public guidance                |
| <b>USE ROLE</b>         | Retention and storage limitation context |
| <b>CHAPTERS USED</b>    | 10                                       |
| <b>QUOTE/PARAPHRASE</b> | Paraphrase only                          |
| <b>NOTE</b>             |                                          |
| <b>BOUNDARY NOTE</b>    | No retention mandate.                    |

**BOUND-01**

|                         |                                                     |
|-------------------------|-----------------------------------------------------|
| <b>TITLE</b>            | CASCO - Committee on conformity assessment          |
| <b>ORGANIZATION</b>     | ISO                                                 |
| <b>SOURCE CLASS</b>     | Standards body                                      |
| <b>USE ROLE</b>         | Conformity-assessment boundary                      |
| <b>CHAPTERS USED</b>    | 12, 13                                              |
| <b>QUOTE/PARAPHRASE</b> | Paraphrase public page only                         |
| <b>NOTE</b>             |                                                     |
| <b>BOUNDARY NOTE</b>    | Validation Lab is not a conformity assessment body. |

**BOUND-02**

|                         |                                       |
|-------------------------|---------------------------------------|
| <b>TITLE</b>            | ISO/IEC 17029:2019 summary            |
| <b>ORGANIZATION</b>     | ISO                                   |
| <b>SOURCE CLASS</b>     | Standards body                        |
| <b>USE ROLE</b>         | Validation/verification body boundary |
| <b>CHAPTERS USED</b>    | 12, 13                                |
| <b>QUOTE/PARAPHRASE</b> | Paraphrase public summary only        |
| <b>NOTE</b>             |                                       |
| <b>BOUNDARY NOTE</b>    | No conformance claim.                 |

**BOUND-03**

|                         |                                              |
|-------------------------|----------------------------------------------|
| <b>TITLE</b>            | Attestation standards                        |
| <b>ORGANIZATION</b>     | AICPA                                        |
| <b>SOURCE CLASS</b>     | Audit/professional body                      |
| <b>USE ROLE</b>         | Attestation and practitioner-report boundary |
| <b>CHAPTERS USED</b>    | 12, 13, 15                                   |
| <b>QUOTE/PARAPHRASE</b> | Paraphrase only                              |
| <b>NOTE</b>             |                                              |
| <b>BOUNDARY NOTE</b>    | This paper is not an attestation engagement. |

**GAIC-SOURCE**

|                              |                                                                                                                 |
|------------------------------|-----------------------------------------------------------------------------------------------------------------|
| <b>TITLE</b>                 | GAIC v0.3.2-FRC-R3 source truth                                                                                 |
| <b>ORGANIZATION</b>          | Jearon Wong / GAIC source master                                                                                |
| <b>SOURCE CLASS</b>          | GAIC-derived source truth                                                                                       |
| <b>USE ROLE</b>              | MROs, RCCS-M/ALCS context, Evidence-Based Validation Pattern, Validation Lab boundary, companion-paper boundary |
| <b>CHAPTERS USED</b>         | 1, 3, 5, 6, 9, 10, 11, 12, 13, 16; appendices                                                                   |
| <b>QUOTE/PARAPHRASE NOTE</b> | Internal source reference                                                                                       |
| <b>BOUNDARY NOTE</b>         | GAIC-derived constructs must be labeled as author framework, not external standard.                             |

---

# R4 Source Register Result

**PASS - all sources used in the R4 candidate source are drawn from the R0 inventory or GAIC source truth.**

---

# R4 Citation Style

R4 uses source IDs in text, tables, and source notes with this register as the package-level lookup. Public endnote or footnote formatting remains a later publication-candidate QA decision.

---

---

# Package Citation Map

# AIAAWP R4 Citation Map

**Status:** R4 publication-candidate preparation citation map **Boundary:** This map supports R4 candidate source review. It does not create final citation formatting, quote source text, or claim endorsement by any source organization.

# Chapter Citation Map

Table rendered as semantic row cards to preserve GAIC-style readability across HTML, PDF, and DOCX review surfaces.

## 0 Executive Summary

|                 |                                                                              |
|-----------------|------------------------------------------------------------------------------|
| SOURCE IDS USED | BF-01, BF-02, BF-05, AUD-01, AI-01, EVID-02, GAIC-SOURCE                     |
| SOURCE ROLE     | Market context; audit evidence; governance; observability; GAIC source truth |
| USE PLACEMENT   | Main argument and source note                                                |
| CITATION RISK   | Big Four endorsement overclaim                                               |
| R4 REVIEW NOTE  | Keep Big Four as context only.                                               |

## 1 Scope, Audience, Boundary

|                 |                                                                         |
|-----------------|-------------------------------------------------------------------------|
| SOURCE IDS USED | AUD-03, BOUND-01, BOUND-02, BOUND-03, AUD-07, BF-02, BF-03, GAIC-SOURCE |
| SOURCE ROLE     | Boundary, audience, assurance/attestation context                       |
| USE PLACEMENT   | Boundary section and table notes                                        |
| CITATION RISK   | Treating boundary sources as endorsement                                |
| R4 REVIEW NOTE  | Repeat non-standard and non-certifying boundary.                        |

## 2 Difference from Model Governance

|                 |                                                                |
|-----------------|----------------------------------------------------------------|
| SOURCE IDS USED | BF-01, BF-04, BF-05, AI-01, AI-02, AI-04, AI-05, AI-08, AUD-06 |
| SOURCE ROLE     | Market and AI governance context                               |
| USE PLACEMENT   | Main text and Table 1                                          |
| CITATION RISK   | Dismissing model governance                                    |
| R4 REVIEW NOTE  | Maintain "necessary but insufficient" framing.                 |

3 Audit Object Shift

|                 |                                                     |
|-----------------|-----------------------------------------------------|
| SOURCE IDS USED | AUD-01, AUD-02, AI-01, AI-08, EVID-01, GAIC-SOURCE  |
| SOURCE ROLE     | Audit evidence, governance, provenance, GAIC MROs   |
| USE PLACEMENT   | Main text                                           |
| CITATION RISK   | Treating Agentic Audit Object as externally sourced |
| R4 REVIEW NOTE  | Label as author synthesis in this paper.            |

4 Logs vs Evidence Chains

|                 |                                                                     |
|-----------------|---------------------------------------------------------------------|
| SOURCE IDS USED | AUD-01, AUD-02, EVID-01, EVID-02, EVID-03, EVID-04, AI-06           |
| SOURCE ROLE     | Audit evidence, observability, logging, incident/governance context |
| USE PLACEMENT   | Main text and Table 2                                               |
| CITATION RISK   | Saying logs are useless                                             |
| R4 REVIEW NOTE  | State logs are useful evidence ingredients.                         |

5 Agentic Audit Object Model

|                 |                                                          |
|-----------------|----------------------------------------------------------|
| SOURCE IDS USED | AUD-01, AUD-04, AI-01, EVID-01, GAIC-SOURCE              |
| SOURCE ROLE     | Evidence, controls, governance, provenance, GAIC objects |
| USE PLACEMENT   | Main text and object table                               |
| CITATION RISK   | Mandatory schema overclaim                               |
| R4 REVIEW NOTE  | Keep model as proposed architecture.                     |

6 MRO Mapping

|                 |                                                  |
|-----------------|--------------------------------------------------|
| SOURCE IDS USED | GAIC-SOURCE, AUD-01, AUD-05, AI-01               |
| SOURCE ROLE     | GAIC MRO source truth; evidence/control language |
| USE PLACEMENT   | Main text and Table 3                            |
| CITATION RISK   | MRO legal mandate overclaim                      |
| R4 REVIEW NOTE  | Keep MROs as GAIC-derived governance objects.    |

7 Evidence Request List

|                 |                                               |
|-----------------|-----------------------------------------------|
| SOURCE IDS USED | AUD-01, AUD-02, AUD-04, AUD-05, AUD-06, AI-06 |
| SOURCE ROLE     | Audit evidence, controls, AI audit practice   |
| USE PLACEMENT   | Main text and Table 4                         |
| CITATION RISK   | Evidence request as formal audit procedure    |
| R4 REVIEW NOTE  | Mark as readiness/evidence architecture.      |

8 Lifecycle Walkthrough

|                 |                                                       |
|-----------------|-------------------------------------------------------|
| SOURCE IDS USED | AUD-04, AUD-05, AUD-06, EVID-01, EVID-02, GAIC-SOURCE |
| SOURCE ROLE     | Controls, audit practice, provenance, traces          |
| USE PLACEMENT   | Main text and Table 5                                 |
| CITATION RISK   | Formal audit procedure implication                    |
| R4 REVIEW NOTE  | Keep walkthrough illustrative.                        |

9 Evidence Partitioning

|                 |                                                                          |
|-----------------|--------------------------------------------------------------------------|
| SOURCE IDS USED | EVID-01, EVID-02, EVID-03, PRIV-01, PRIV-02, PRIV-03, AI-09, GAIC-SOURCE |
| SOURCE ROLE     | Provenance/logging/privacy/context partitioning                          |
| USE PLACEMENT   | Main text and Table 6                                                    |
| CITATION RISK   | Universal disclosure or retention                                        |
| R4 REVIEW NOTE  | Stress scope, access, and minimization.                                  |

10 Privacy and Minimization

|                 |                                                                 |
|-----------------|-----------------------------------------------------------------|
| SOURCE IDS USED | PRIV-01, PRIV-02, PRIV-03, PRIV-04, PRIV-05, AI-09, GAIC-SOURCE |
| SOURCE ROLE     | Privacy and retention context                                   |
| USE PLACEMENT   | Main text and Table 7                                           |
| CITATION RISK   | Legal advice or GDPR proof                                      |
| R4 REVIEW NOTE  | Require legal/privacy review in R3.                             |

11 Exception and Closure

|                 |                                                     |
|-----------------|-----------------------------------------------------|
| SOURCE IDS USED | EVID-04, AI-02, AUD-04, AUD-05, AUD-06, GAIC-SOURCE |
| SOURCE ROLE     | Incident, control, remediation, audit practice      |
| USE PLACEMENT   | Main text and Table 10                              |
| CITATION RISK   | Legal remedy/settlement overclaim                   |
| R4 REVIEW NOTE  | Closure is evidence state only.                     |

12 Third-Party Validation Boundary

|                 |                                                         |
|-----------------|---------------------------------------------------------|
| SOURCE IDS USED | AUD-03, BOUND-01, BOUND-02, BOUND-03, GAIC-SOURCE       |
| SOURCE ROLE     | Assurance, attestation, conformity, validation boundary |
| USE PLACEMENT   | Main text and Table 8                                   |
| CITATION RISK   | Certification or assurance opinion overclaim            |
| R4 REVIEW NOTE  | Review every Validation Lab paragraph.                  |

13 AARM

|                 |                                                                                  |
|-----------------|----------------------------------------------------------------------------------|
| SOURCE IDS USED | AUD-01, AUD-05, AUD-06, AI-01, BOUND-03, GAIC-SOURCE                             |
| SOURCE ROLE     | Evidence, AI audit practice, governance, attestation boundary, AARM source truth |
| USE PLACEMENT   | Main text and Table 9                                                            |
| CITATION RISK   | Maturity/certification overclaim                                                 |
| R4 REVIEW NOTE  | Keep "what it does not prove" column.                                            |

14 Enterprise Readiness

|                 |                                                                       |
|-----------------|-----------------------------------------------------------------------|
| SOURCE IDS USED | BF-02, BF-04, AUD-04, AUD-07, AI-01, PRIV-03, GAIC-SOURCE             |
| SOURCE ROLE     | Enterprise governance, controls, three lines, privacy, market context |
| USE PLACEMENT   | Main text                                                             |
| CITATION RISK   | Procurement/readiness guarantee                                       |
| R4 REVIEW NOTE  | Bridge only to Guide 1/Guide 2.                                       |

**15 Audit and Assurance Firm Use**

|                        |                                                                       |
|------------------------|-----------------------------------------------------------------------|
| <b>SOURCE IDS USED</b> | BF-02, BF-03, AUD-03, AUD-05, AUD-06, BOUND-03                        |
| <b>SOURCE ROLE</b>     | Market context, professional practice, assurance/attestation boundary |
| <b>USE PLACEMENT</b>   | Main text                                                             |
| <b>CITATION RISK</b>   | Big Four endorsement or methodology replacement                       |
| <b>R4 REVIEW NOTE</b>  | Use "may use as discussion/readiness structure."                      |

**16 Conclusion**

|                        |                                                                      |
|------------------------|----------------------------------------------------------------------|
| <b>SOURCE IDS USED</b> | GAIC-SOURCE, AUD-01, AUD-03, AI-01                                   |
| <b>SOURCE ROLE</b>     | Object chain, audit evidence, assurance boundary, governance context |
| <b>USE PLACEMENT</b>   | Main text                                                            |
| <b>CITATION RISK</b>   | Insurability overclaim before WP3                                    |
| <b>R4 REVIEW NOTE</b>  | Keep WP3 bridge high level.                                          |

## Appendix Citation Map

| Appendix                | Source IDs used                                       | Source role                                      | Citation risk                          | R4 re-view note                              |
|-------------------------|-------------------------------------------------------|--------------------------------------------------|----------------------------------------|----------------------------------------------|
| A Evidence Request List | AUD-01, AUD-02, AUD-04, AUD-05, AUD-06, GAIC-SOURCE   | Evidence/control vocabulary and MRO mapping      | Formal audit-procedure implication     | Keep request list as readiness architecture. |
| B Walkthrough Template  | AUD-04, AUD-05, AUD-06, EVID-01, EVID-02, GAIC-SOURCE | Controls, provenance, traces, MRO lifecycle      | Audit procedure implication            | Keep illustrative.                           |
| C MRO Mapping           | GAIC-SOURCE, AUD-01, AUD-05, AI-01                    | GAIC MRO source truth and audit/control language | Legal mandate overclaim                | Verify MRO names exactly.                    |
| D Readiness Matrix      | AUD-01, AUD-05, AUD-06, AI-01, BOUND-03, GAIC-SOURCE  | Evidence, audit practice, governance, boundary   | Certification/maturity-score overclaim | Keep no scores.                              |
| E Closure Checklist     | EVID-04, AI-02, AUD-04, AUD-05, AUD-06,               | Incident, control, remediation, MRO closure      | Legal closure overclaim                | Closure means evidence state only.           |

| Appendix            | Source IDs used                                   | Source role                                                                     | Citation risk                                        | R4 re-view note              |
|---------------------|---------------------------------------------------|---------------------------------------------------------------------------------|------------------------------------------------------|------------------------------|
|                     | GAIC-SOURCE                                       |                                                                                 |                                                      |                              |
| F Boundary Language | AUD-03, BOUND-01, BOUND-02, BOUND-03, GAIC-SOURCE | Assurance, conformity, attestation, Validation Lab and companion-paper boundary | Boundary wording becoming legal disclaimer overreach | Keep as editorial guardrail. |

---

# R4 Citation Map Result

**PASS - R4 candidate source uses R0 source IDs and GAIC source truth with explicit source-role and citation-risk notes.**

---

AIAAWP-2026-v0.1-R4-CANDIDATE - public research candidate. Generated 2026-05-18T02:16:35Z. Deployment verification pending. No final/sealed/live claim.